

CAP 781

MACHINE LEARNING

Tanzeela Javid Kaloo (32638)

Assistant Professor

System And Architecture

Lovely Professional University

UNIT – V

Neural Network and Deep Learning

Tanzeela Javid Kaloo (32638)

Assistant Professor

System And Architecture

Lovely Professional University

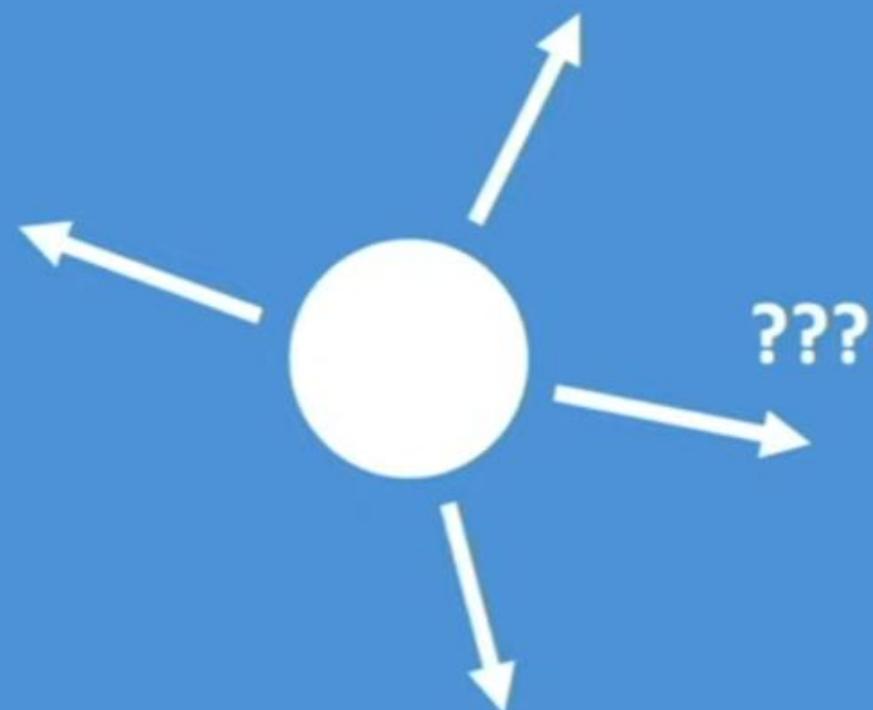
Content

- Introduction to Neural Networks,
- Perceptron and Multilayer Perceptron (MLPs),
- Activation Functions: Sigmoid, ReLU, and others,
- Convolutional Neural Networks (CNNs) for Image Recognition,
- Recurrent Neural Networks (RNNs) for Sequence Data,
- Transfer Learning and Pre-trained Models

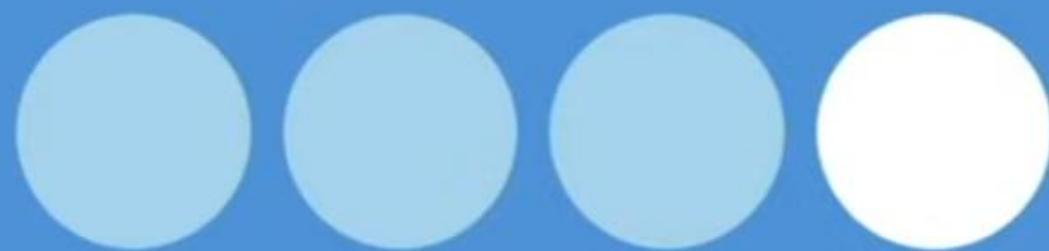
Given an image of a ball,
can you predict where it will go next?



Given an image of a ball,
can you predict where it will go next?



Given an image of a ball,
can you predict where it will go next?



Given an image of a ball,
can you predict where it will go next?



Sequences in the Wild



Audio

Sequences in the Wild



Audio

Sequences in the Wild

6.S191 Introduction to Deep Learning

Text

character:

6 . S | 9 |

word: Introduction to Deep Learning
Text

character:

6 . S | 9 |

word:

Introduct



Learning



6 . S | 9 |



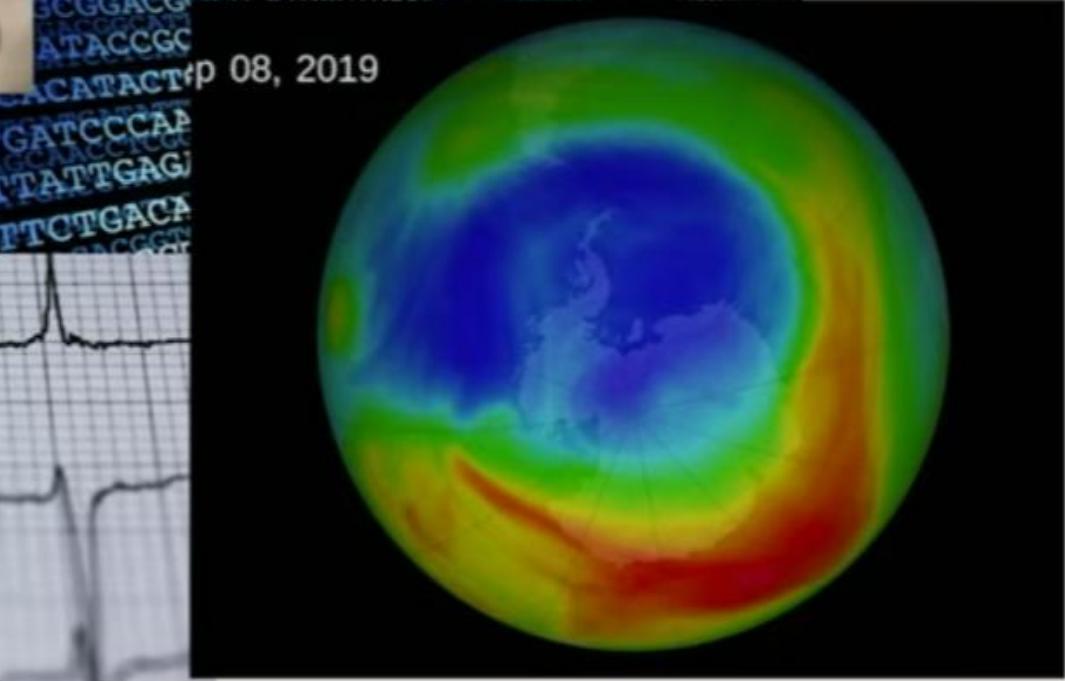
word: Introduct

p Learning



08, 2019

A close-up view of a DNA sequence, showing a series of nucleotide bases (A, T, C, G) arranged in a vertical column. The sequence is partially visible, showing "GGTACCCCTTCACATAC", "TGCCAGACTGATCCC", and "TAAGAGGAATTATTGAG".



Sequence Modeling Applications



One to One
Binary Classification



"Will I pass this class?"
Student → Pass?

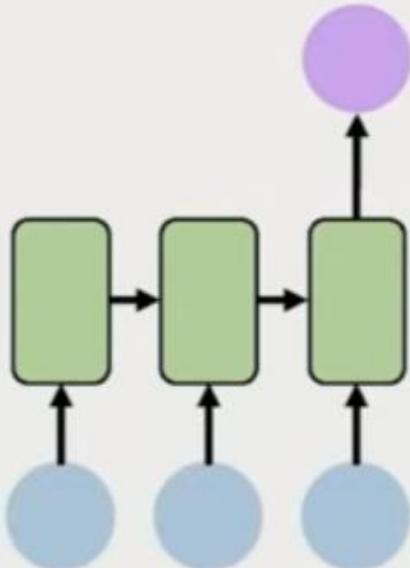
Sequence Modeling Applications



One to One
Binary Classification



"Will I pass this class?"
Student → Pass?



Many to One
Sentiment Classification



Ivar Hagendoorn
@ivarhagendoorn



The @MIT Introduction to #DeepLearning is definitely one of the best courses of its kind currently available online introtodeeplearning.com

12:45 PM - 12 Feb 2018



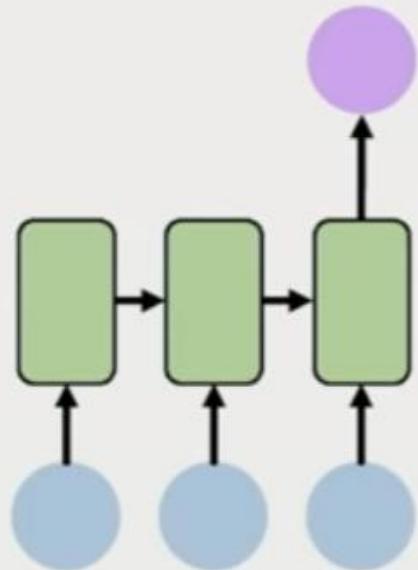
Sequence Modeling Applications



One to One
Binary Classification



"Will I pass this class?"
Student → Pass?



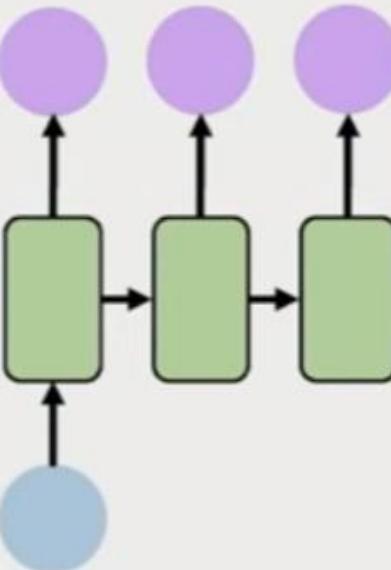
Many to One
Sentiment Classification



Ivar Hagendoorn
@ivarhagendoorn

The @MIT Introduction to #DeepLearning is definitely one of the best courses of its kind currently available online introtodeeplearning.com

12:40 PM - 12 Feb 2018



One to Many
Image Captioning



"A baseball player throws a ball."

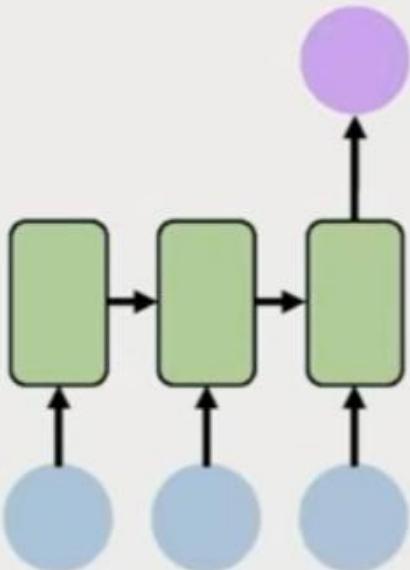
Sequence Modeling Applications



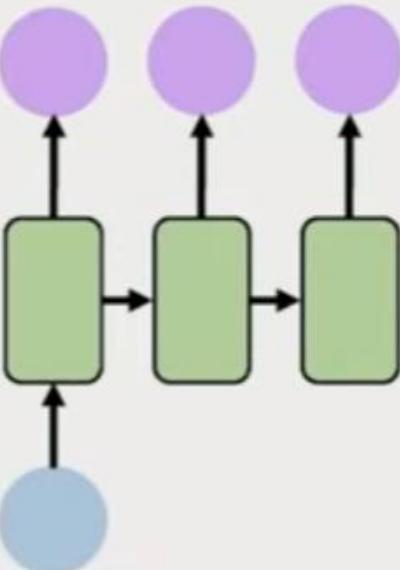
One to One
Binary Classification



"Will I pass this class?"
Student → Pass?



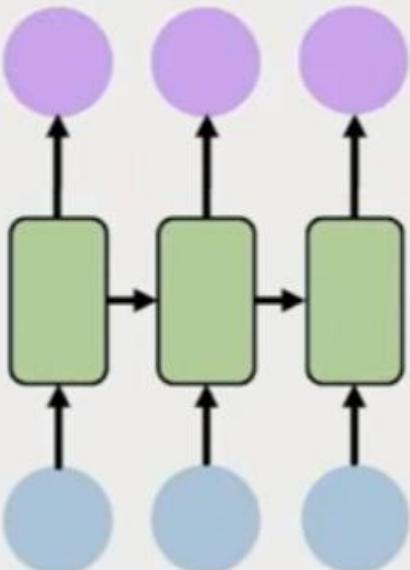
Many to One
Sentiment Classification



One to Many
Image Captioning



"A baseball player throws a ball."

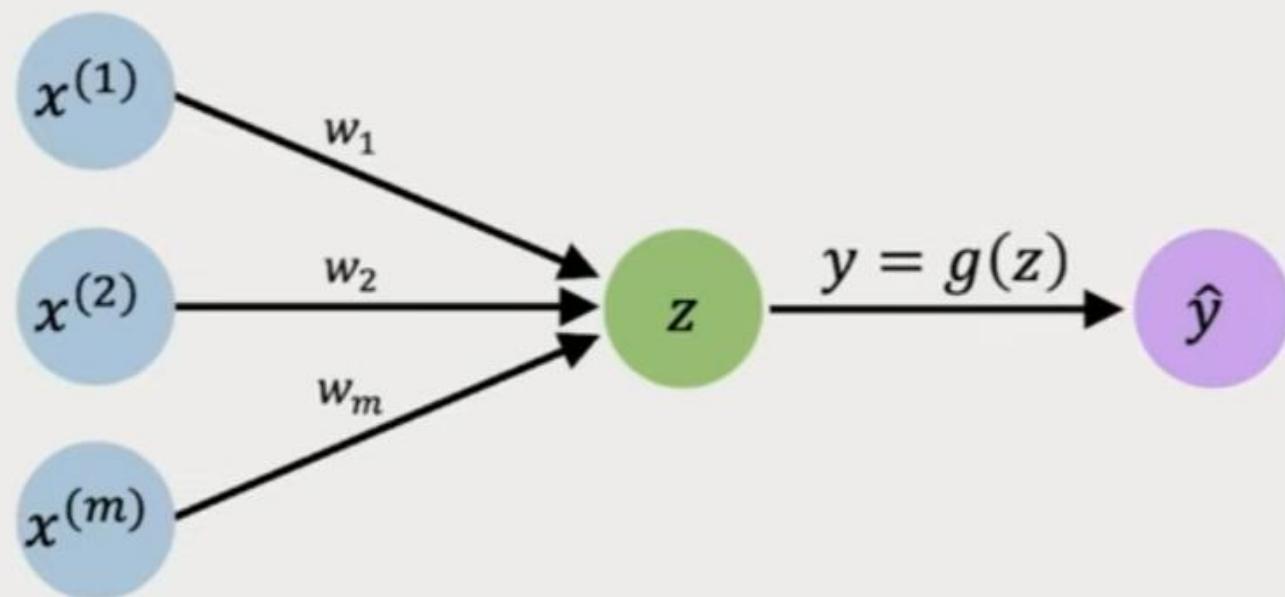


Many to Many
Machine Translation

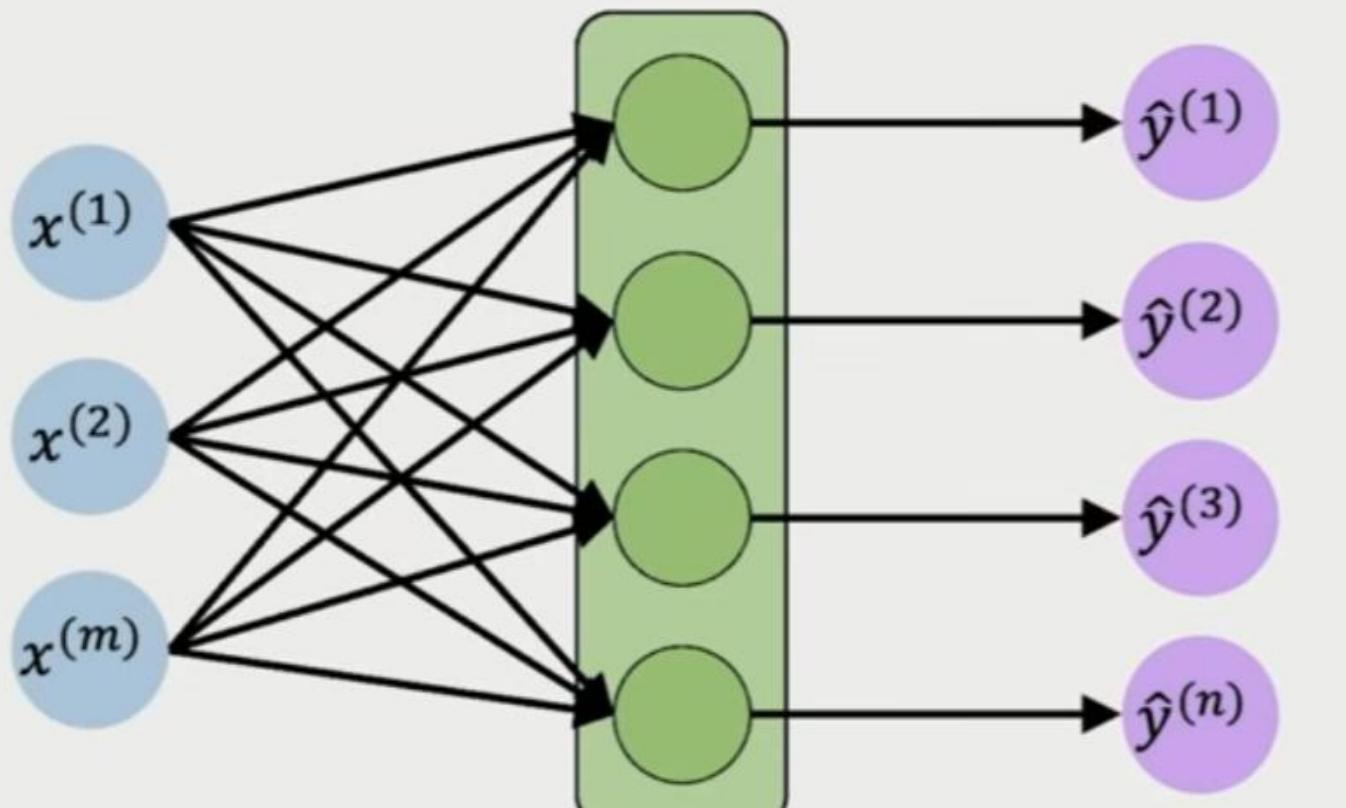


Neurons with Recurrence

The Perceptron Revisited



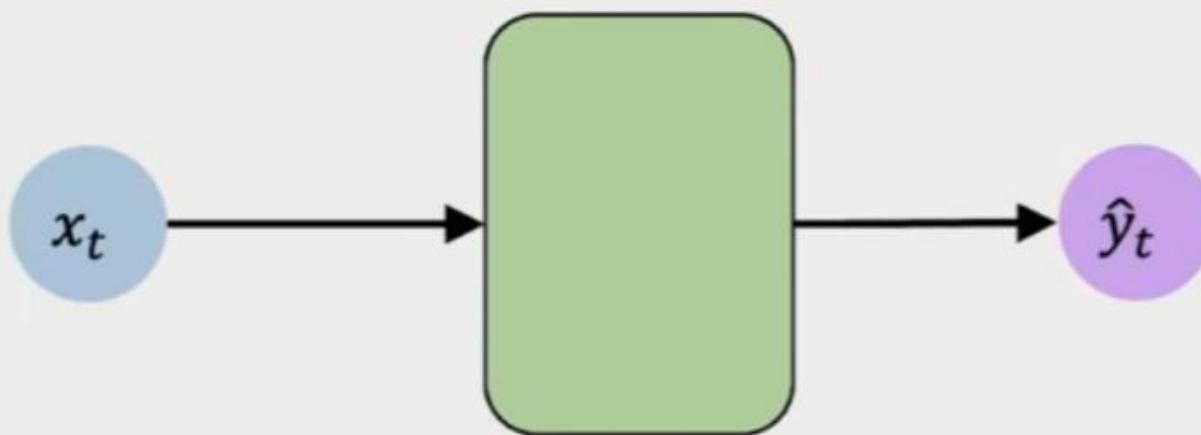
Feed-Forward Networks Revisited



$$\boldsymbol{x} \in \mathbb{R}^m$$

$$\hat{\boldsymbol{y}} \in \mathbb{R}^n$$

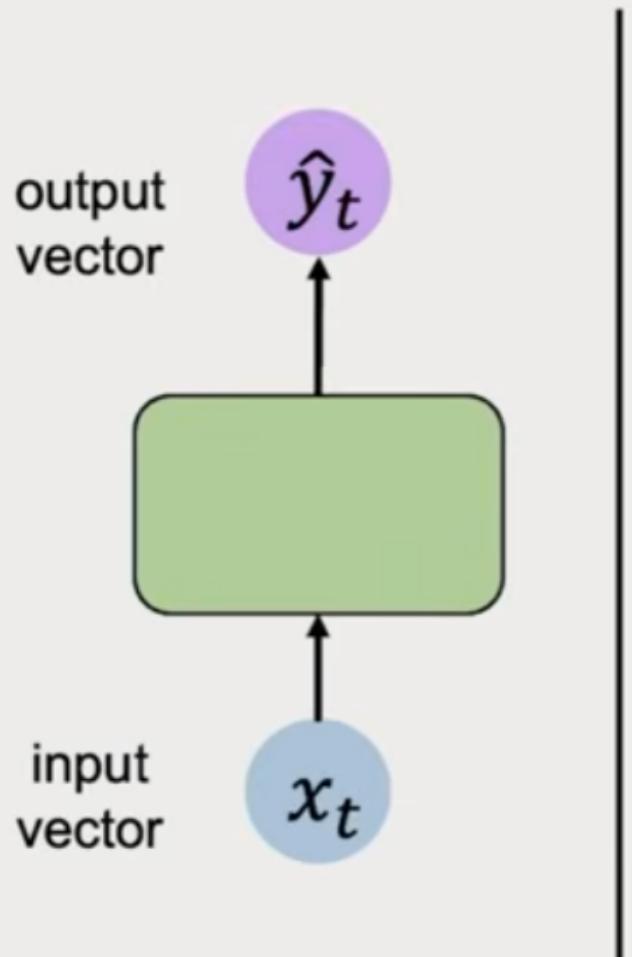
Feed-Forward Networks Revisited



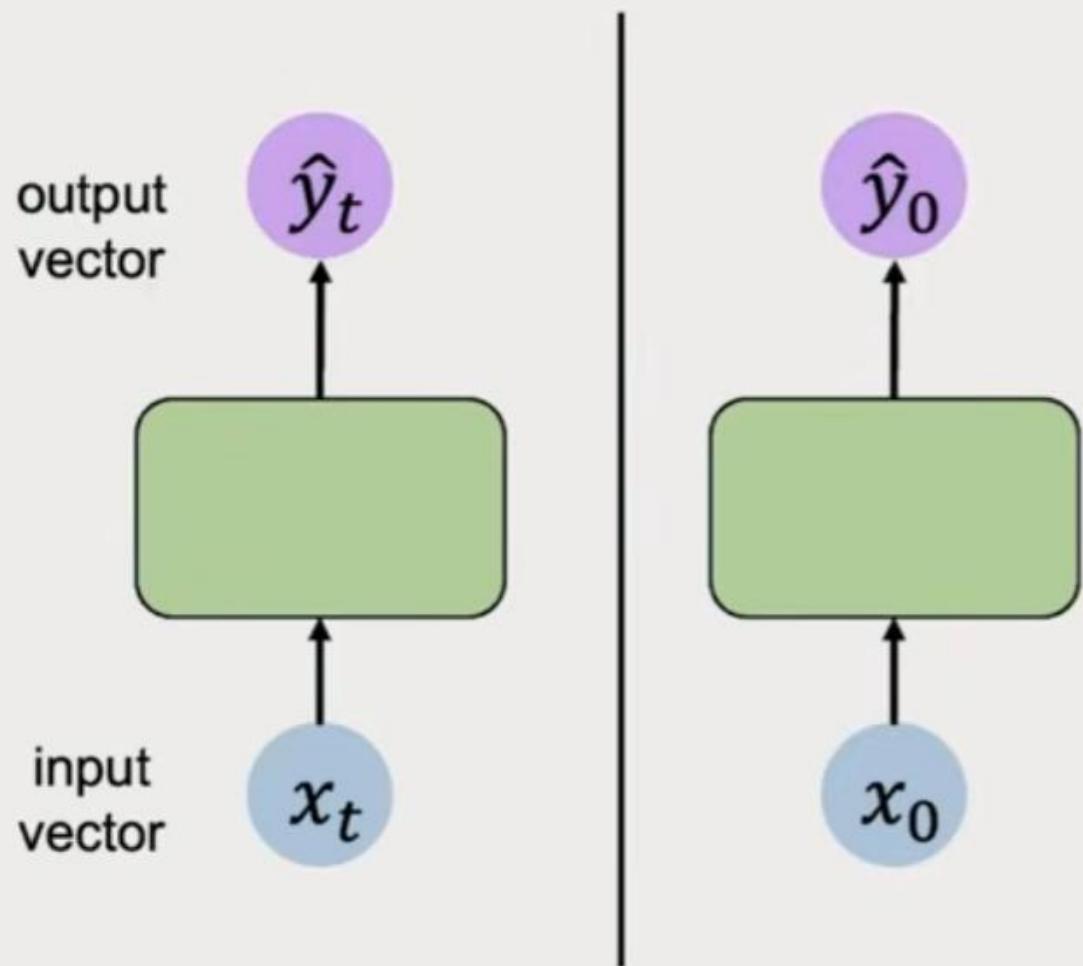
$$x_t \in \mathbb{R}^m$$

$$\hat{y}_t \in \mathbb{R}^n$$

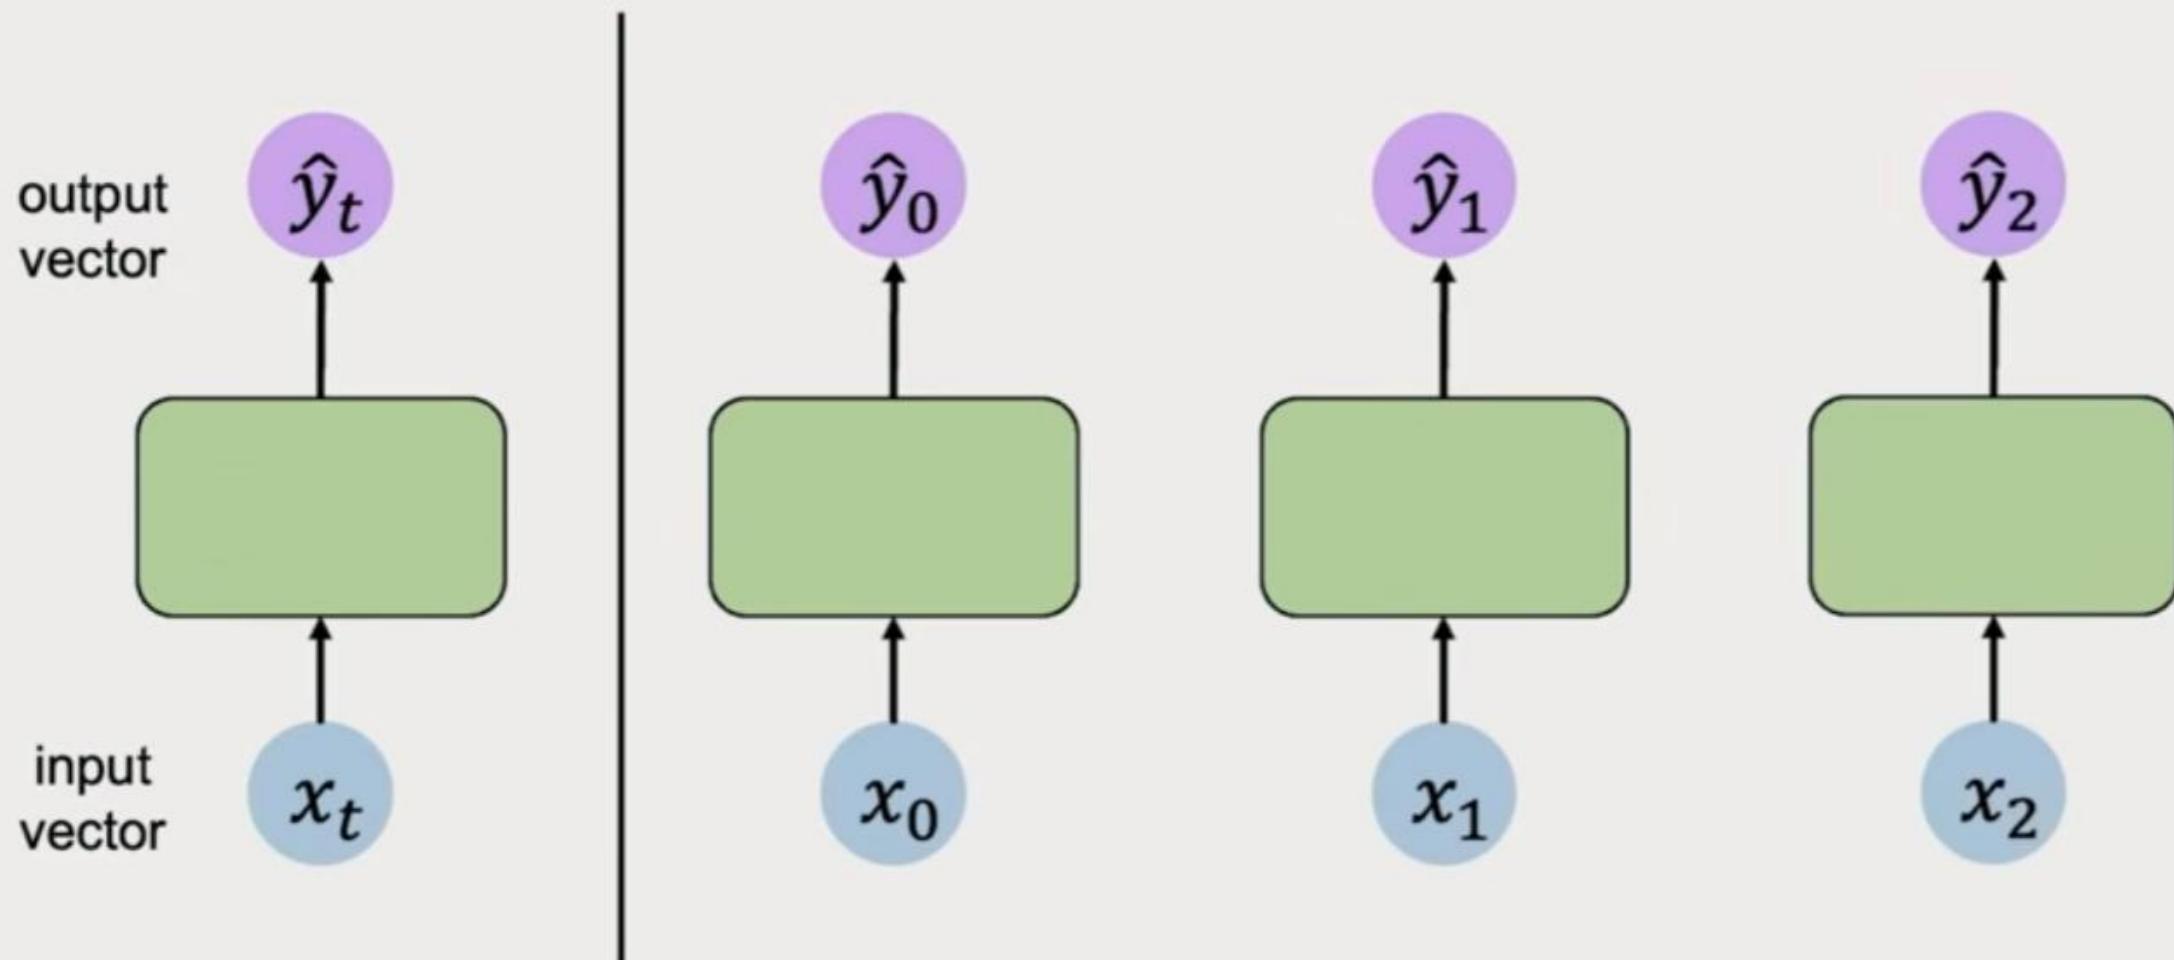
Handling Individual Time Steps



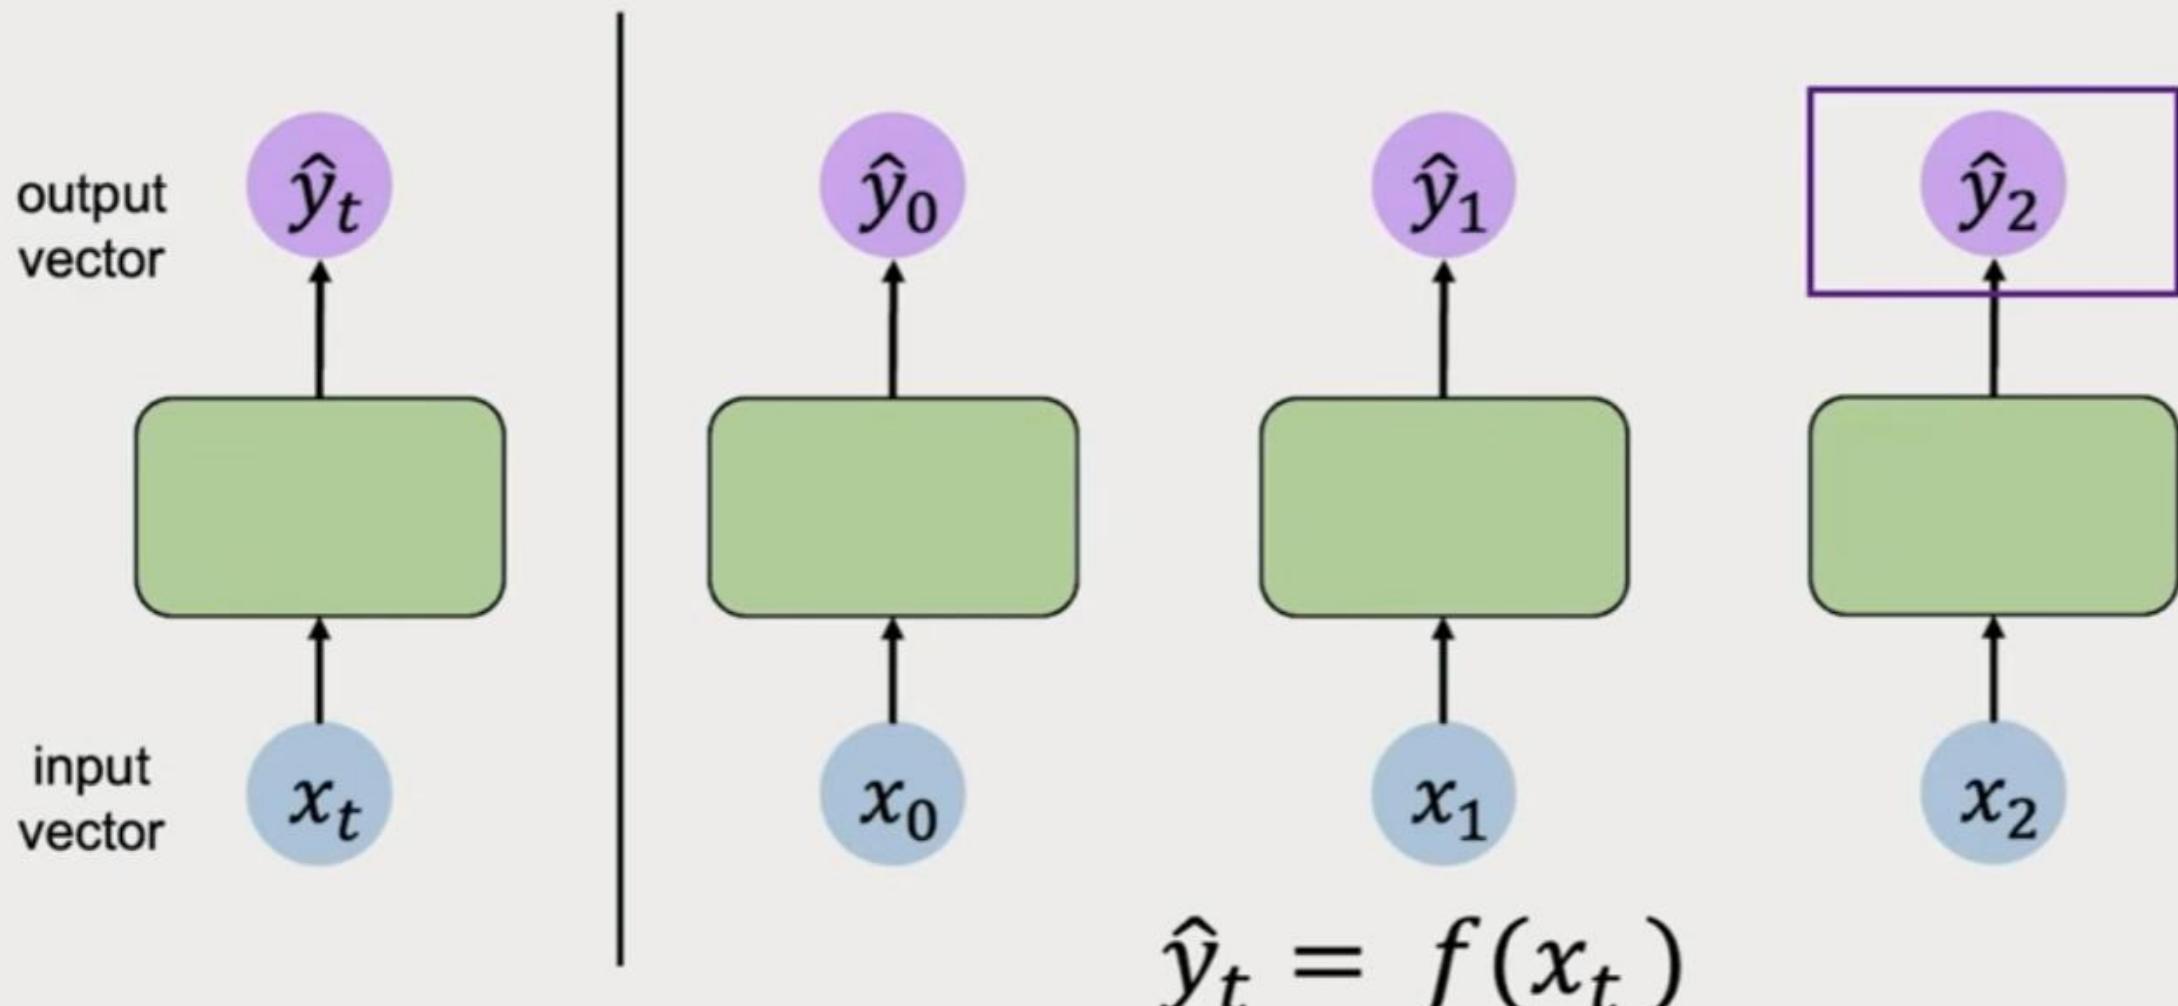
Handling Individual Time Steps



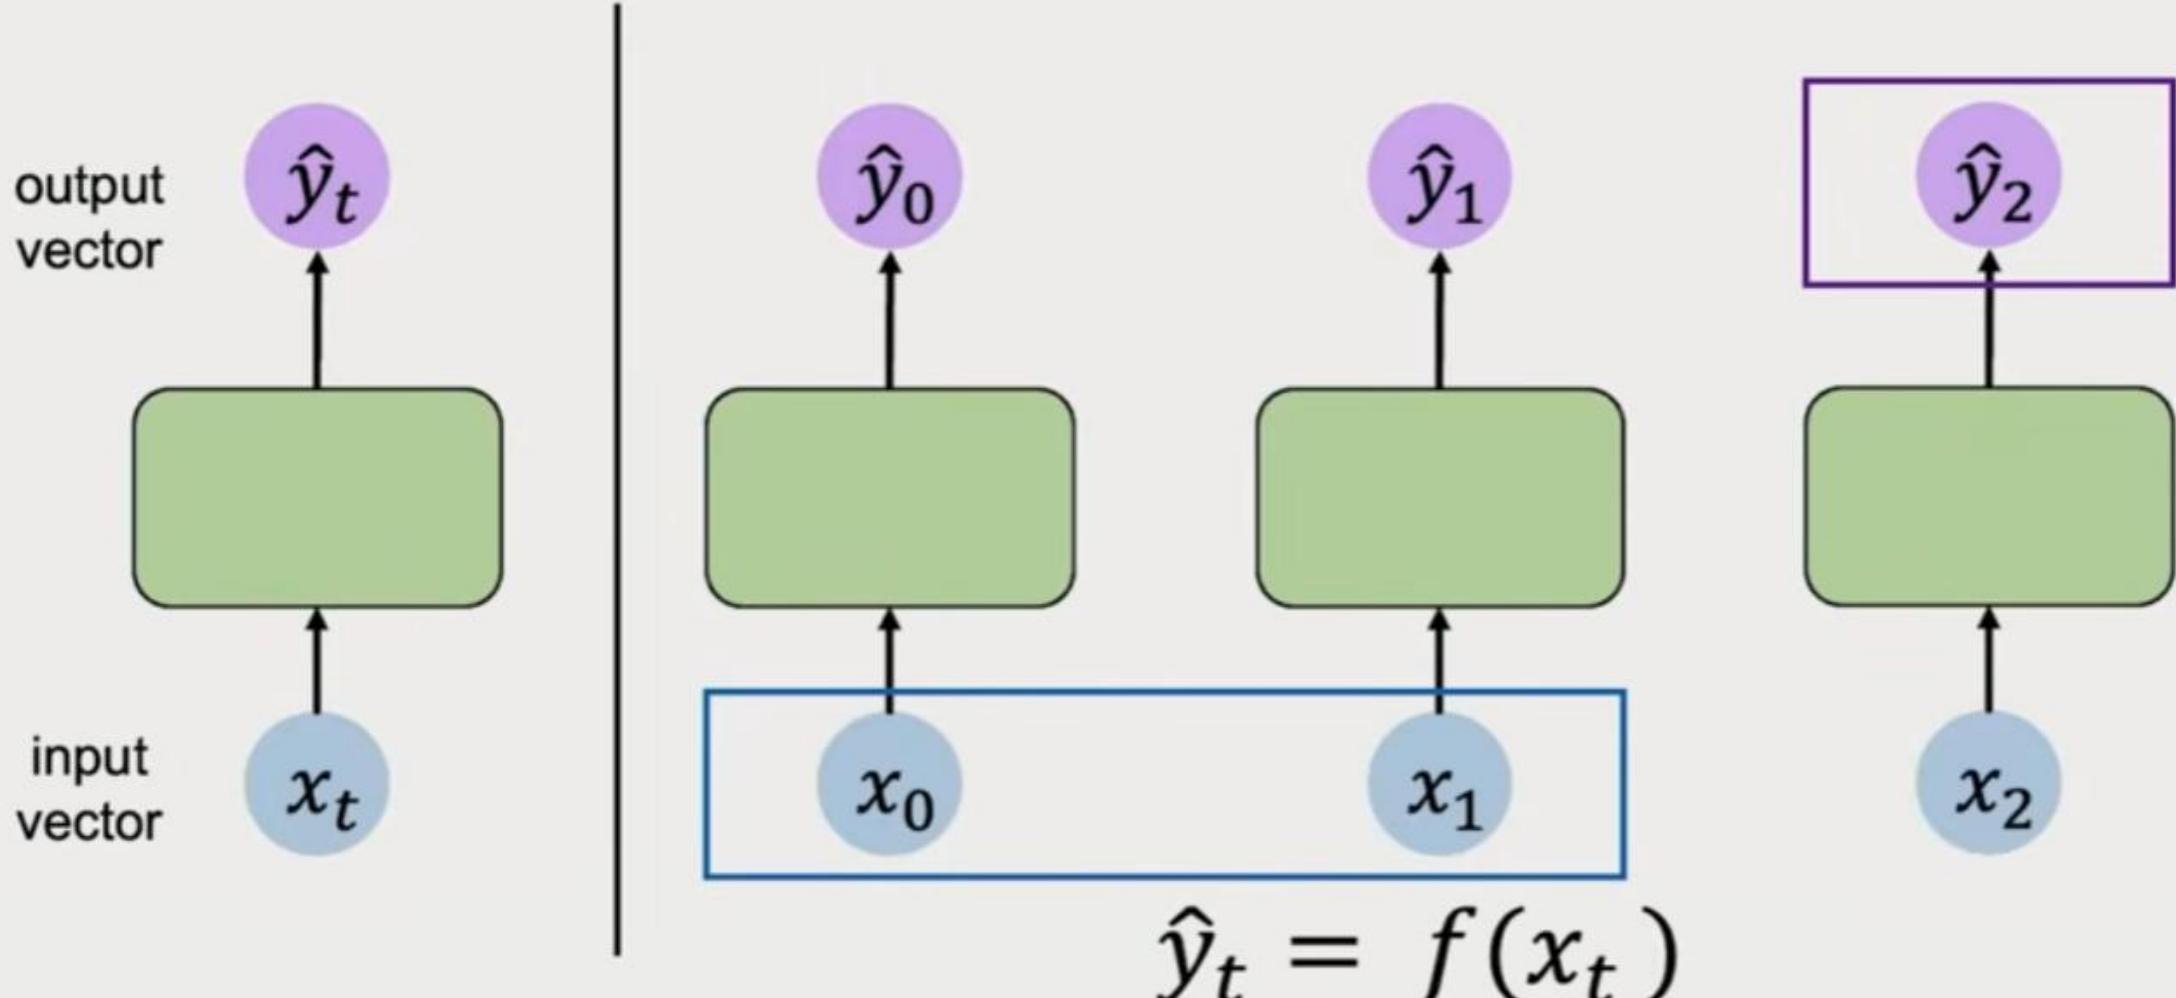
Handling Individual Time Steps



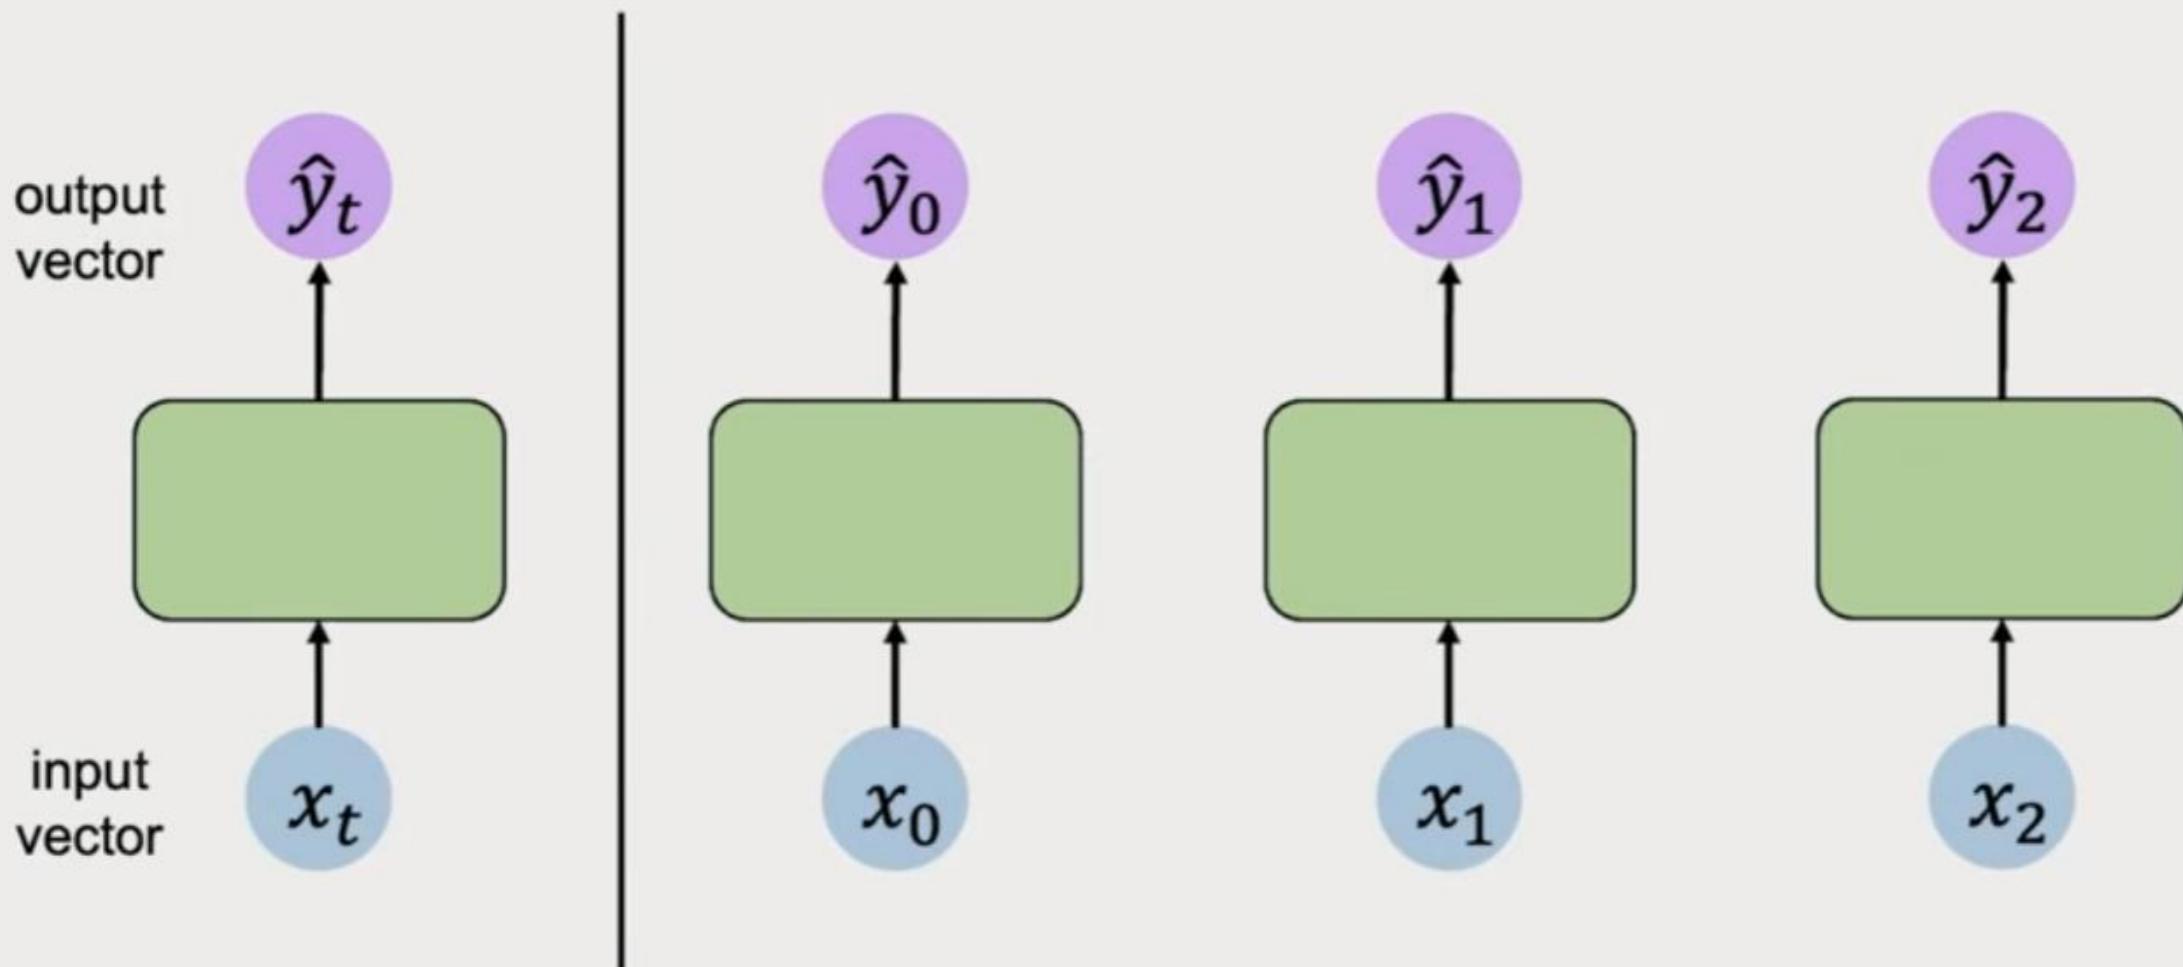
Handling Individual Time Steps



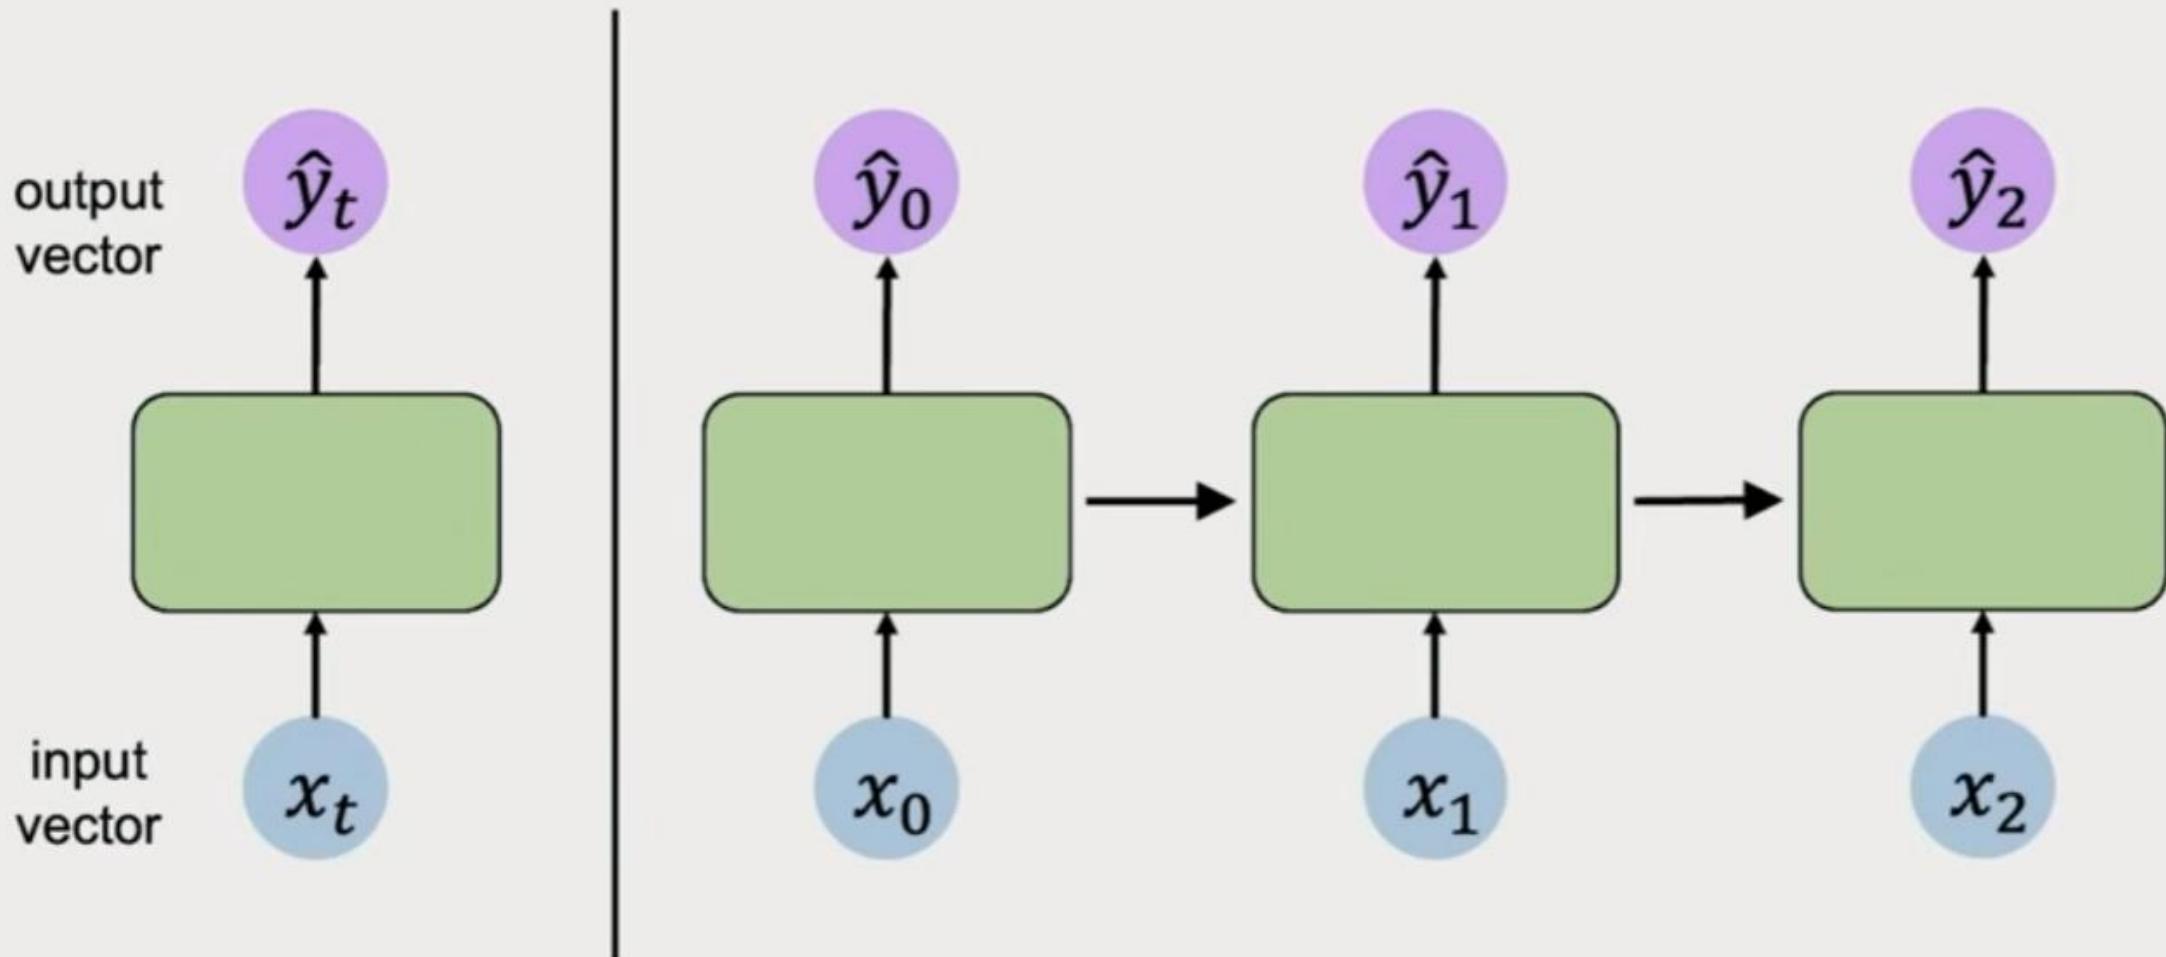
Handling Individual Time Steps



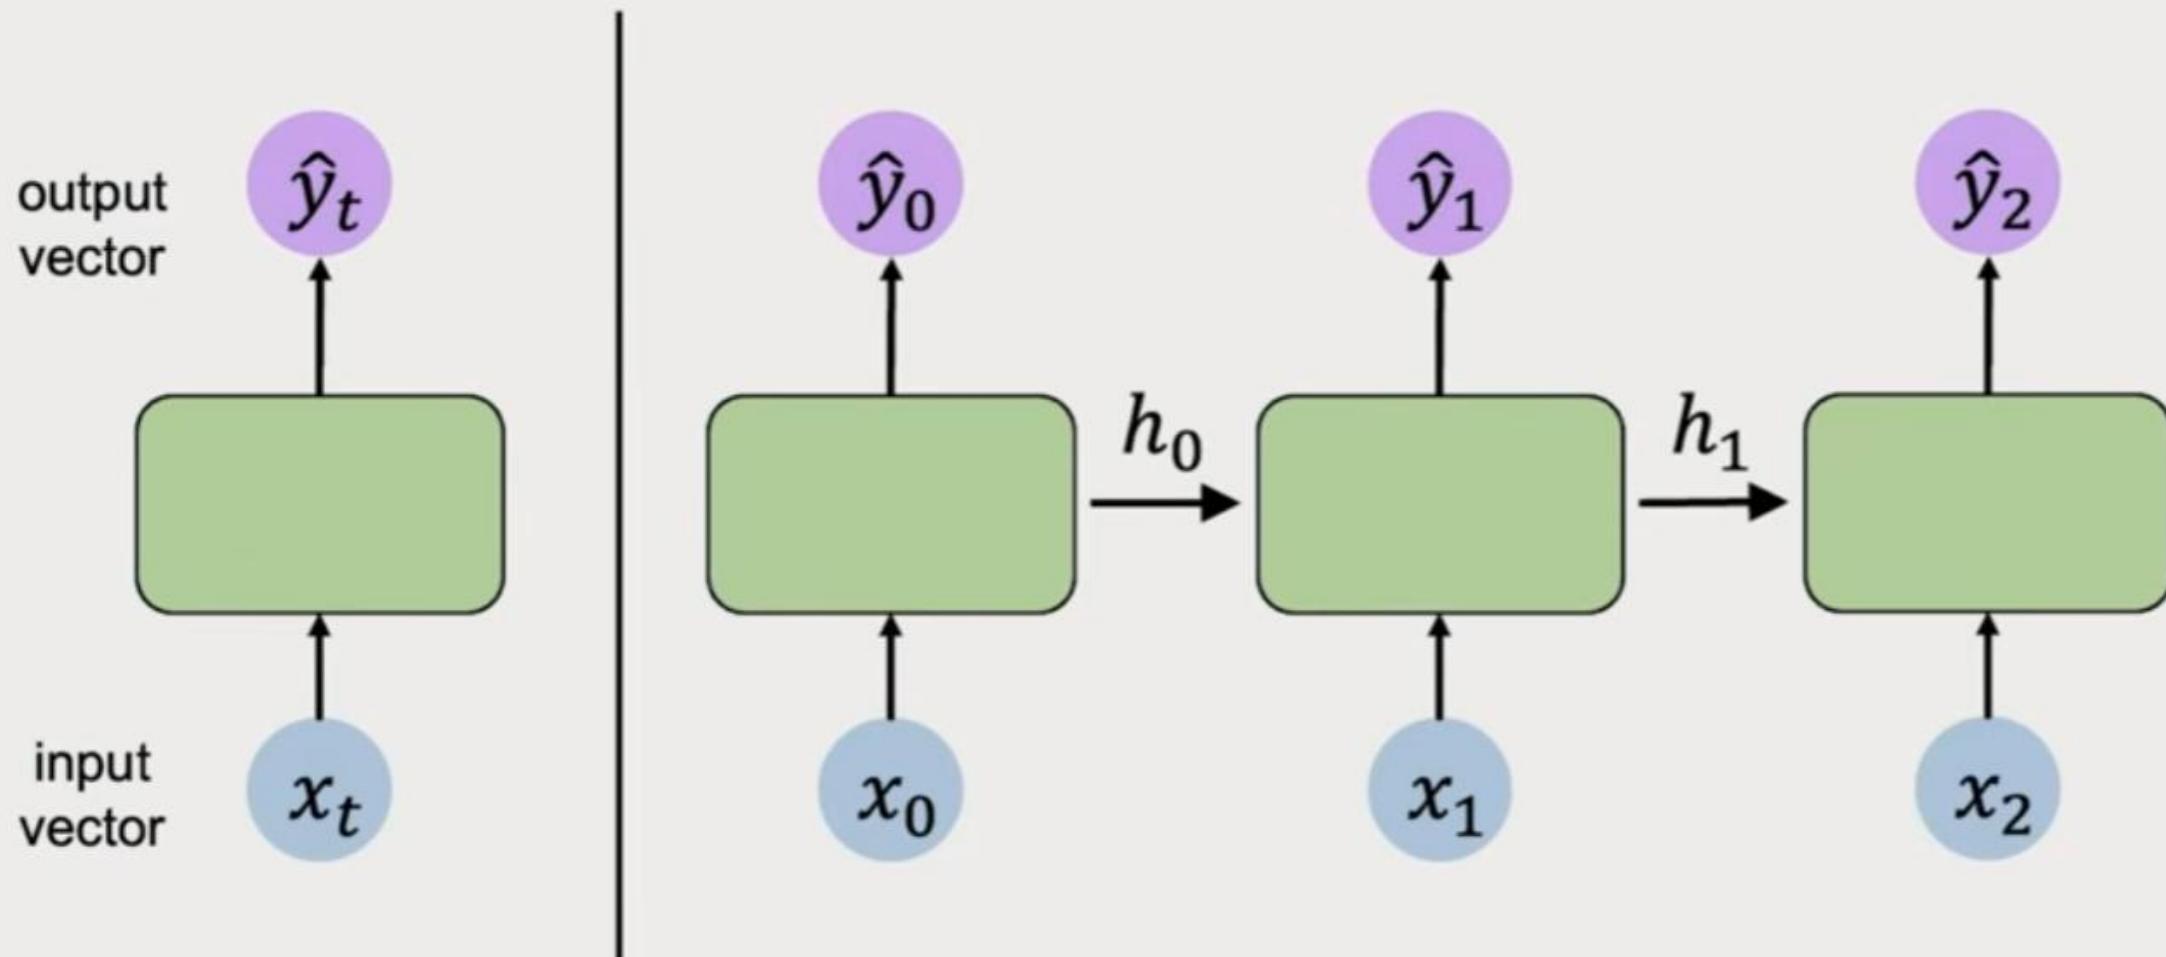
Neurons with Recurrence



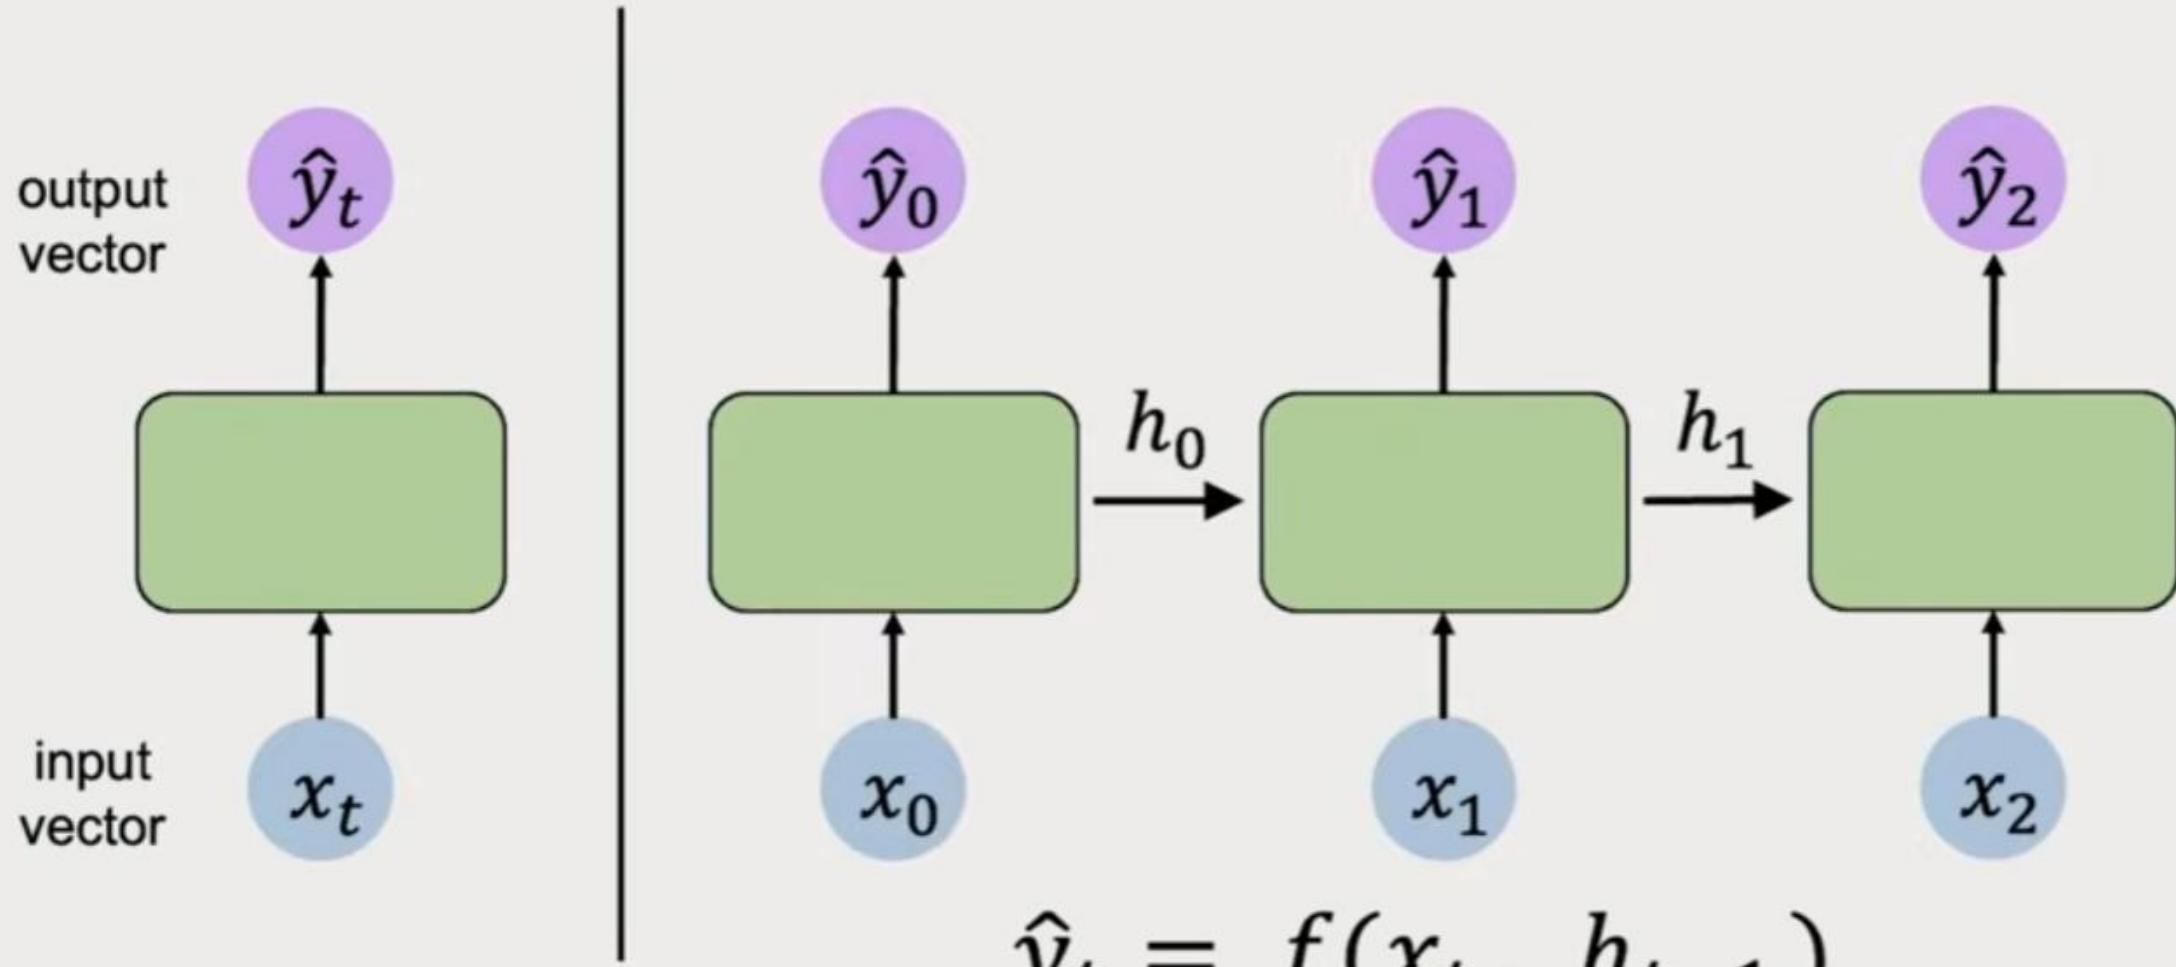
Neurons with Recurrence



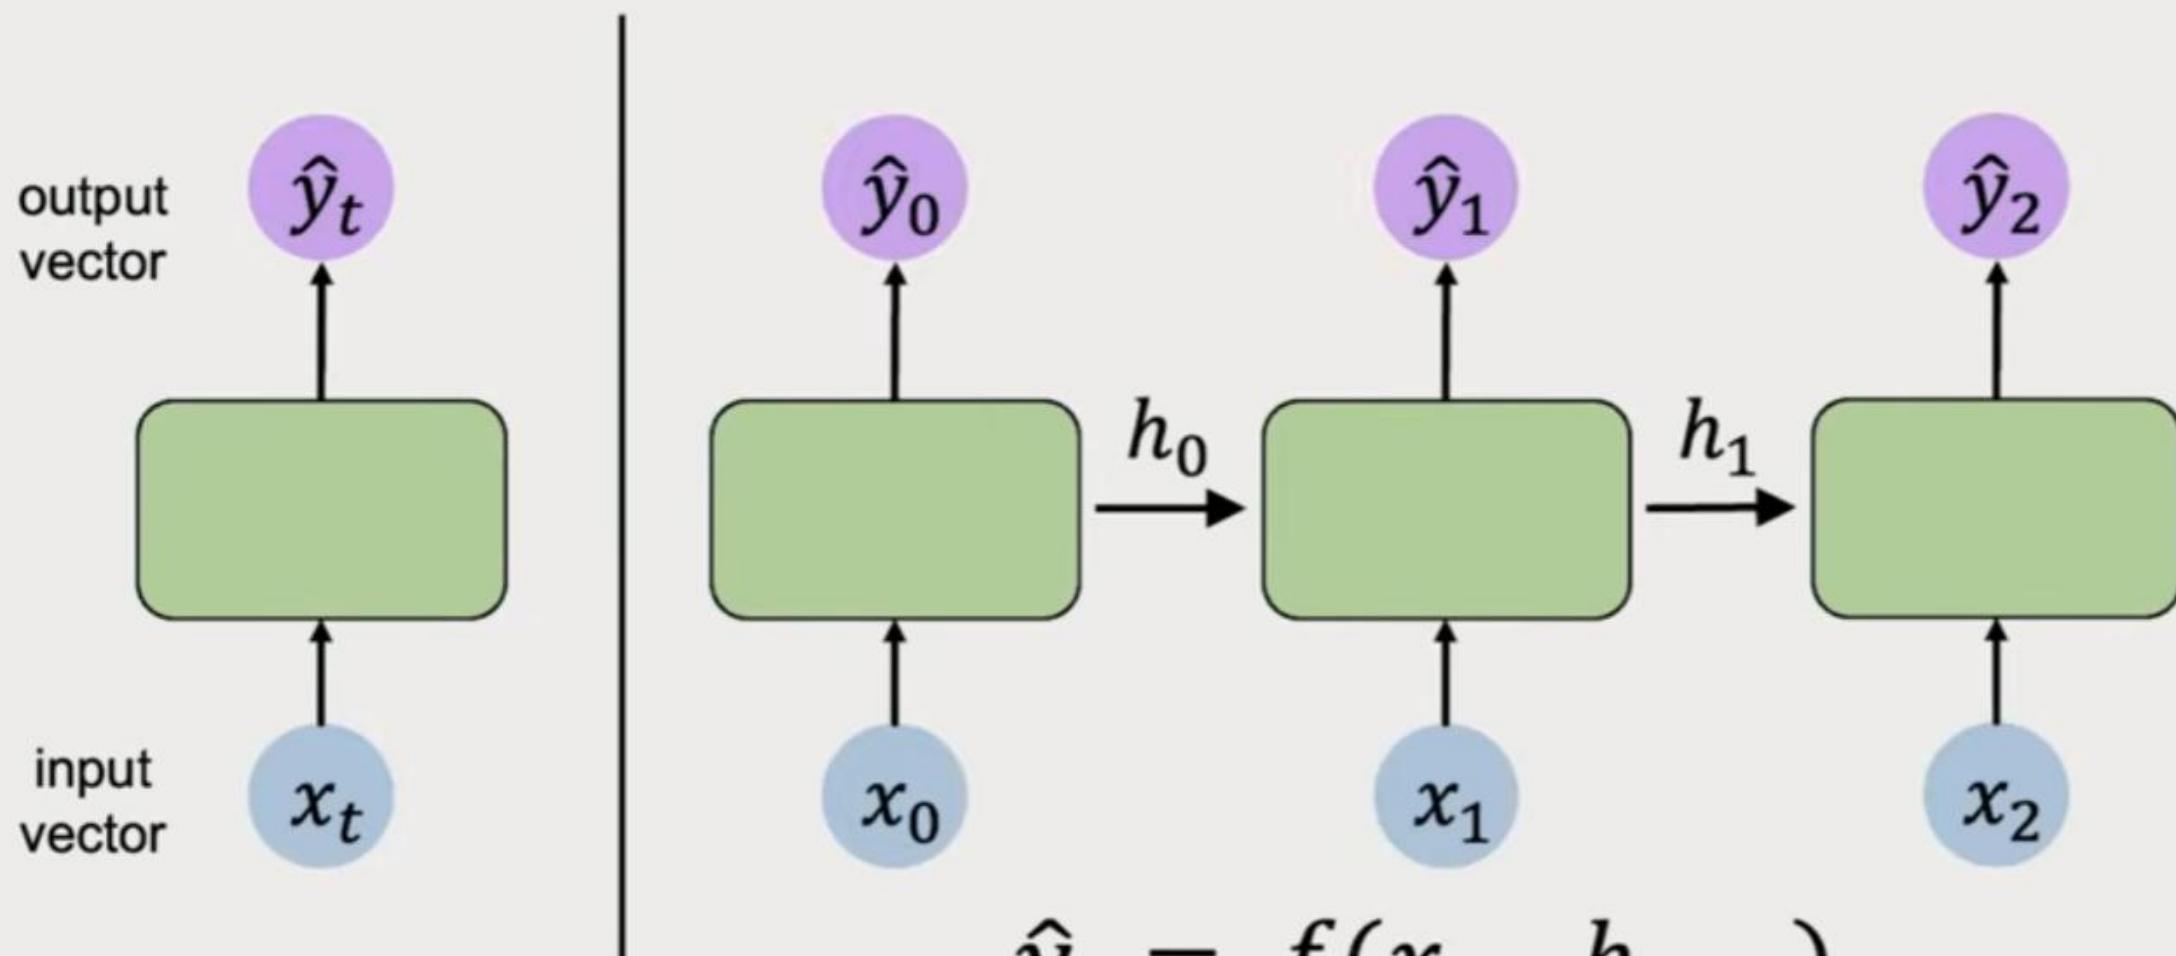
Neurons with Recurrence



Neurons with Recurrence



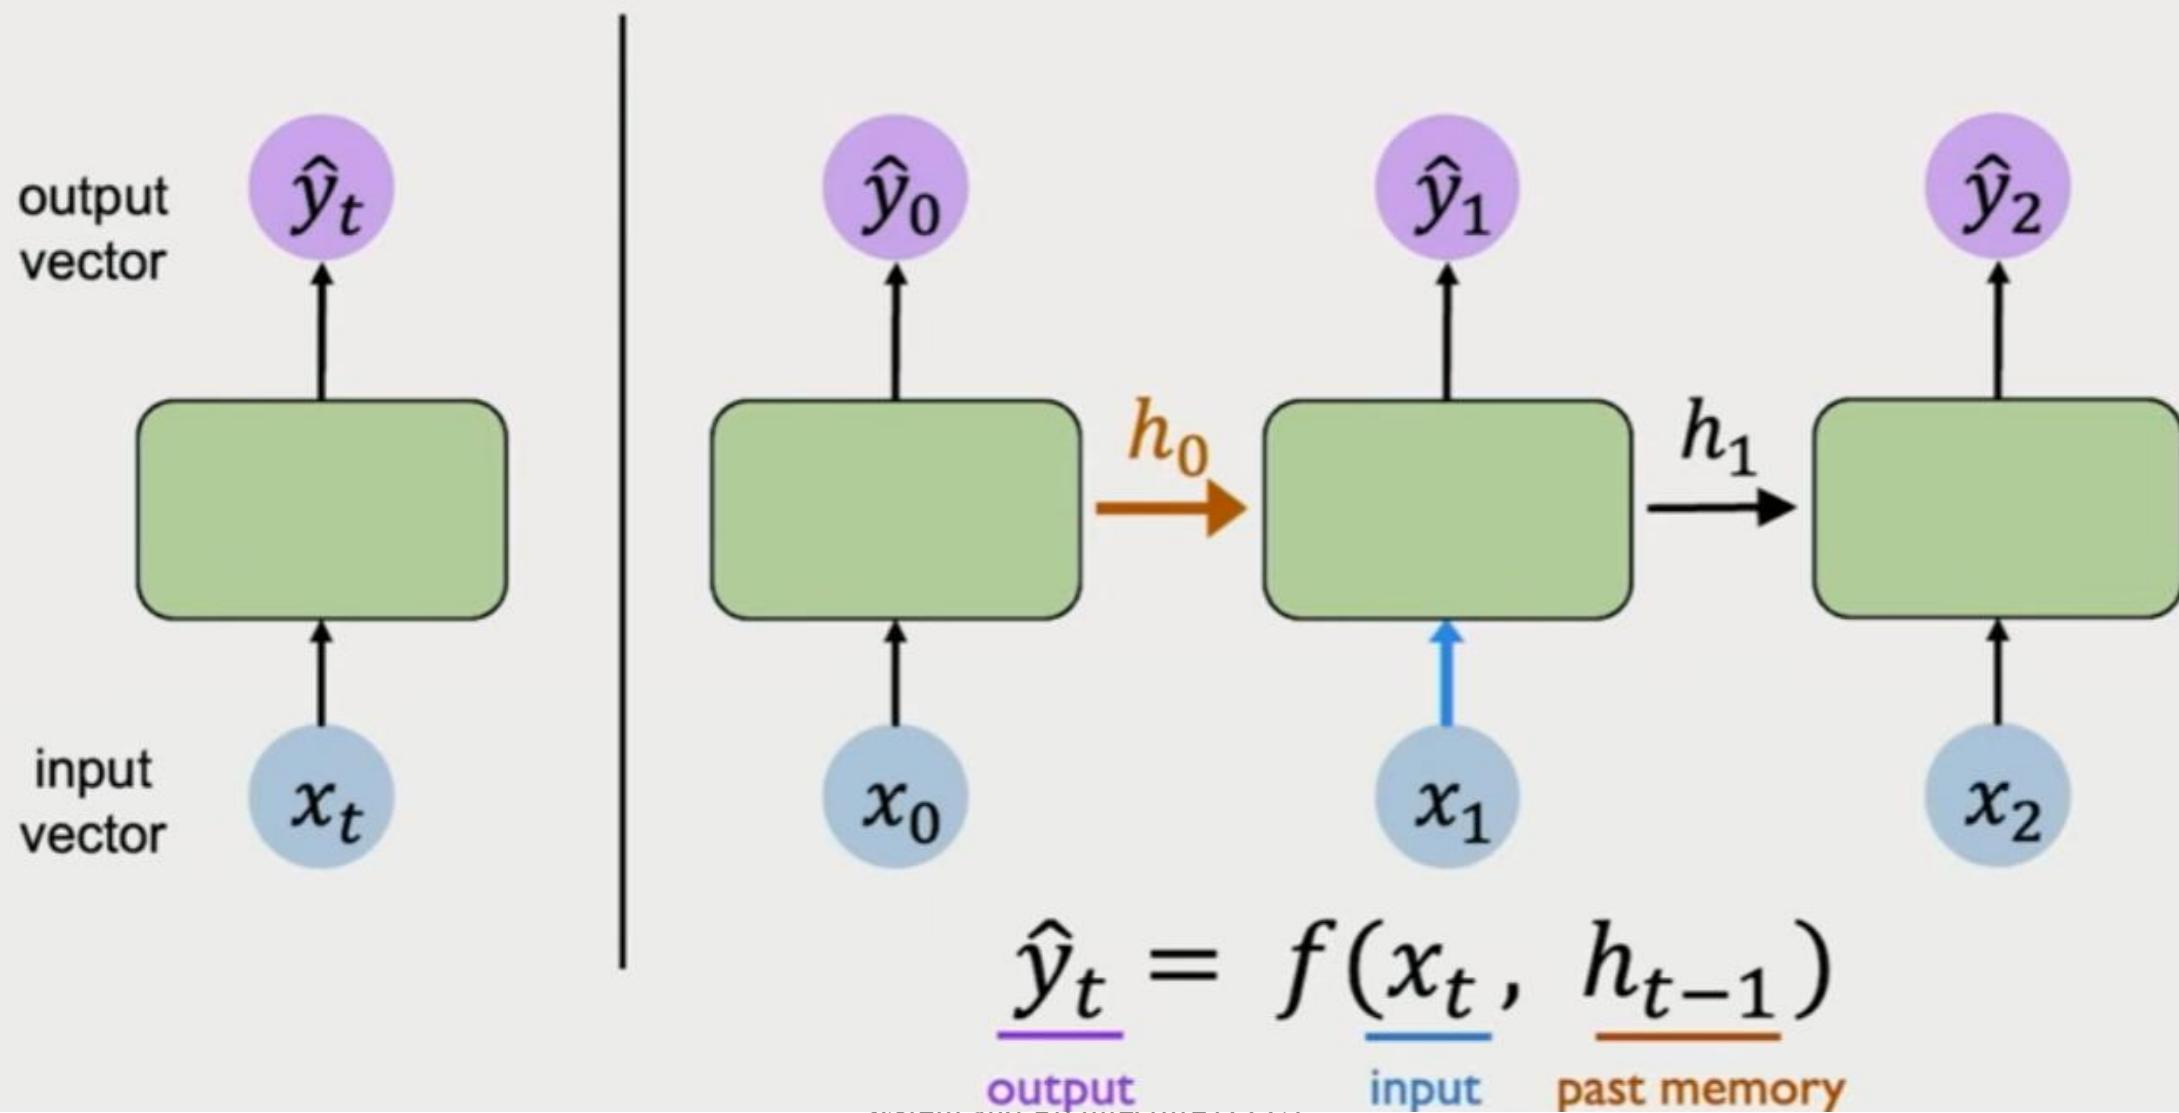
Neurons with Recurrence



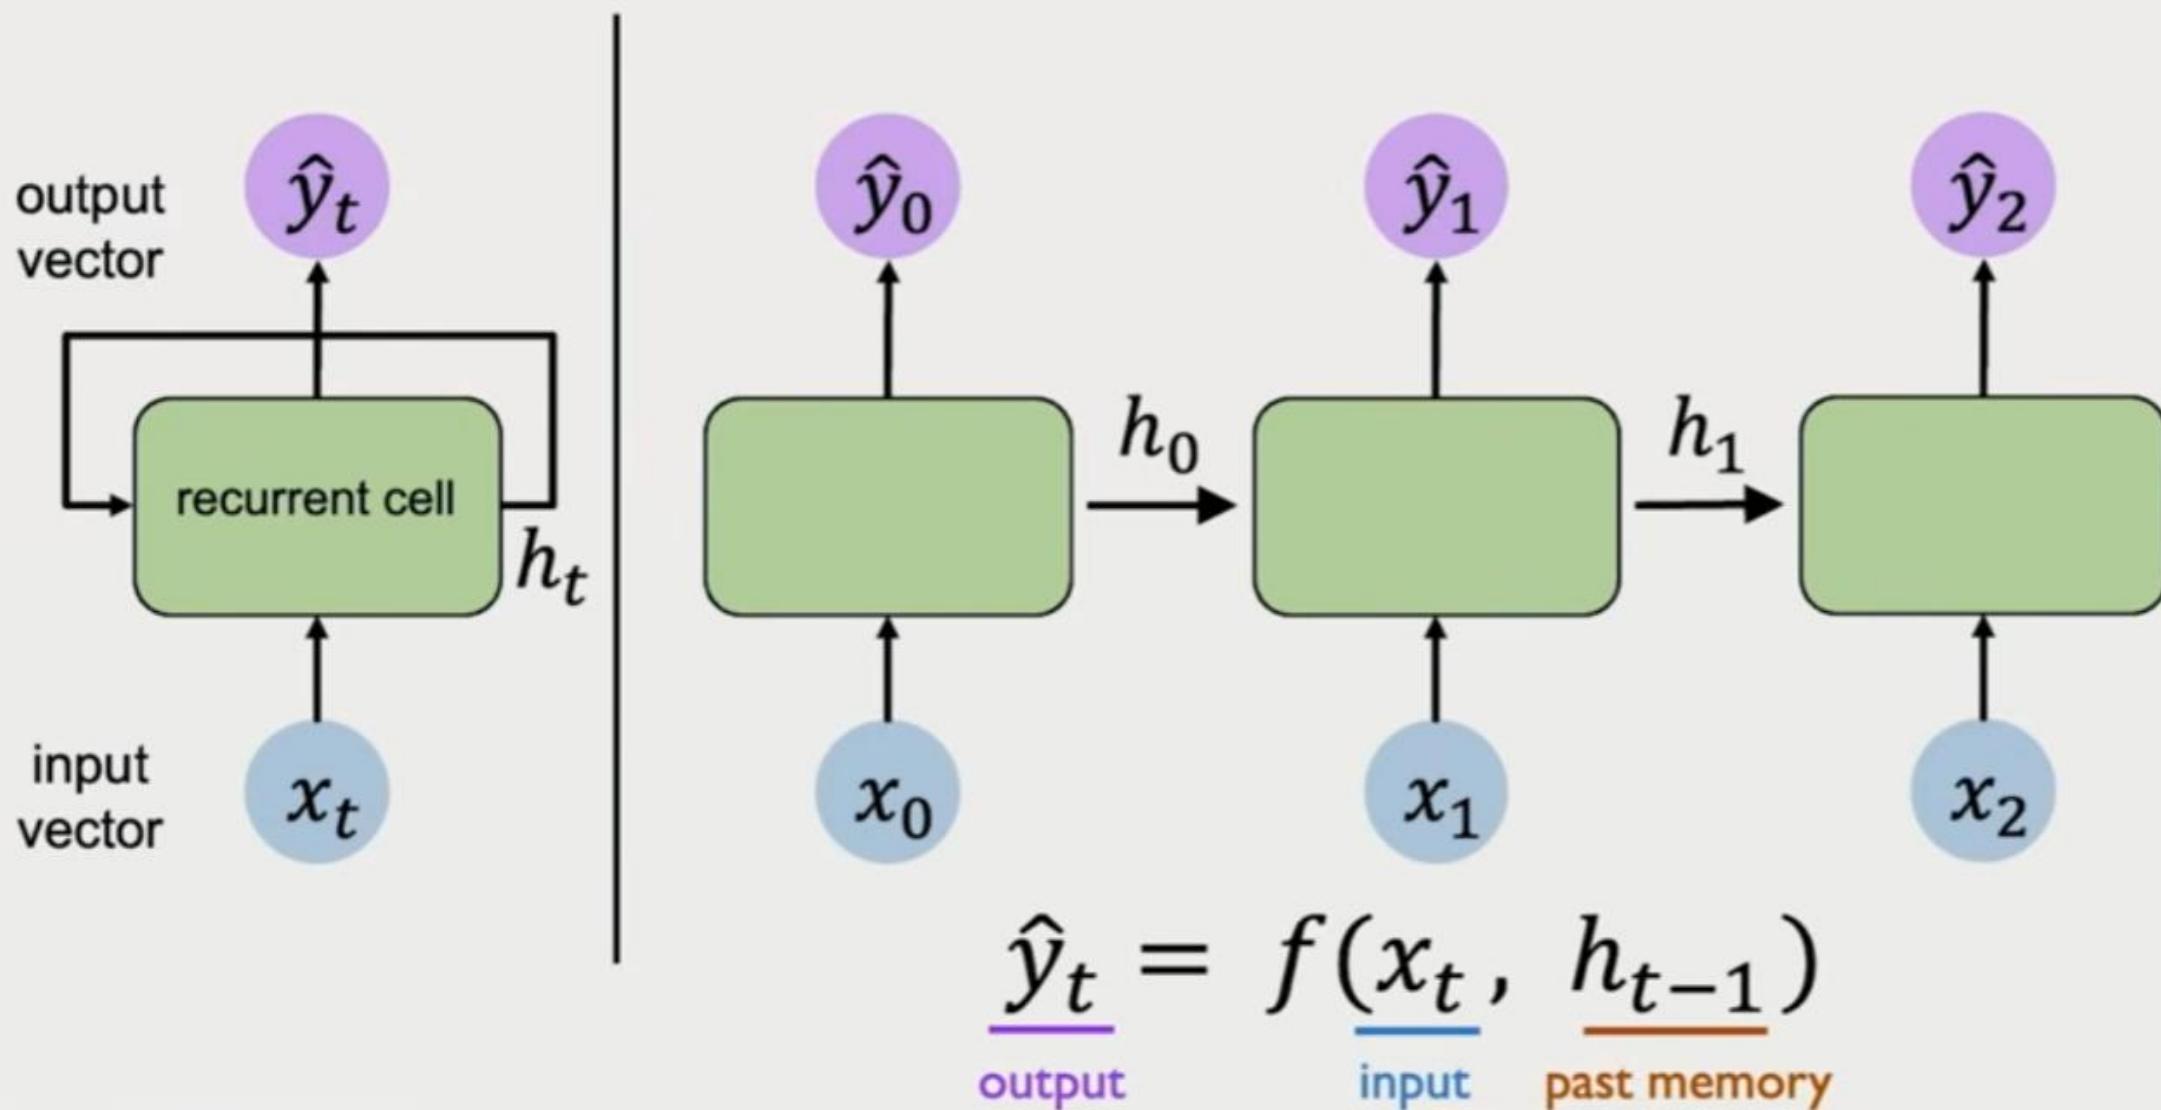
$$\underline{\hat{y}_t} = f(x_t, h_{t-1})$$

output

Neurons with Recurrence

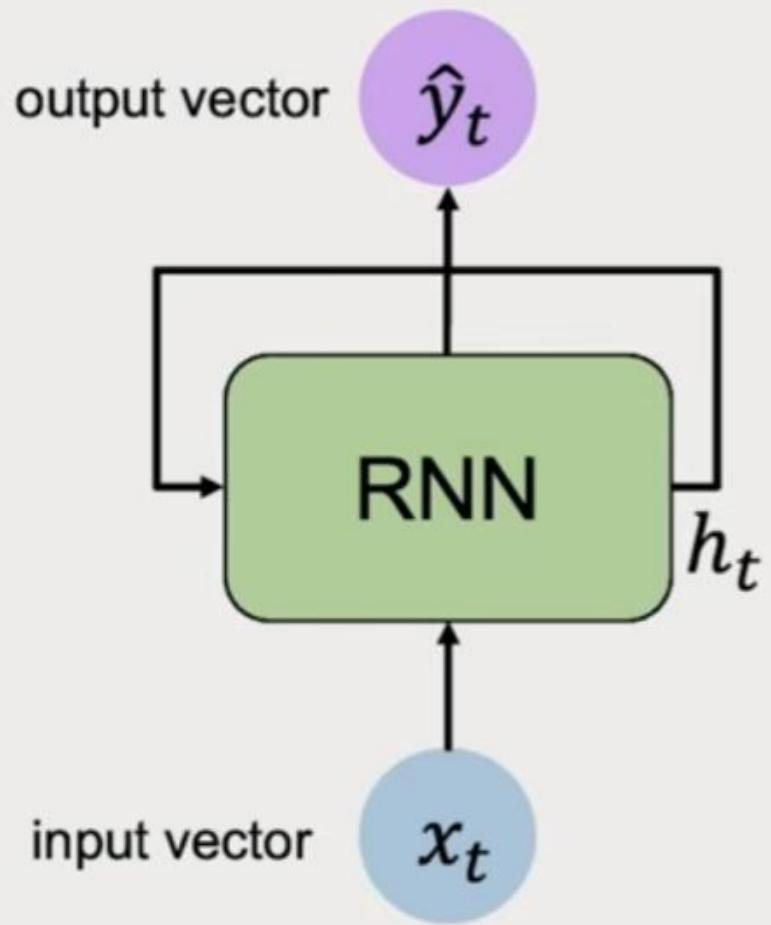


Neurons with Recurrence



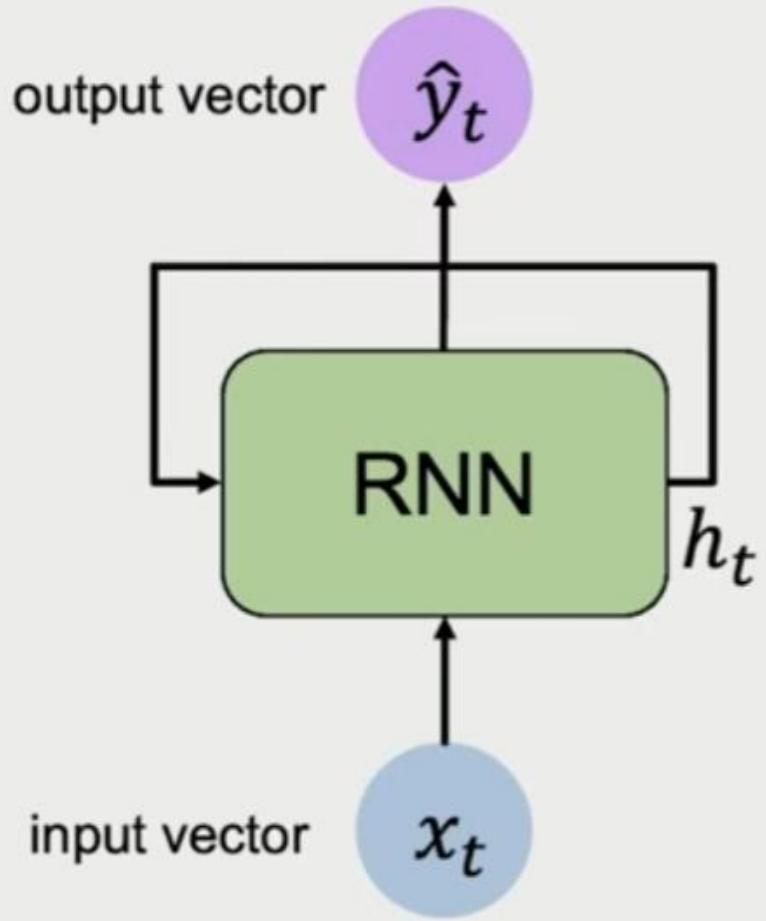
Recurrent Neural Networks (RNNs)

Recurrent Neural Networks (RNNs)



RNNs have a **state**, h_t , that is updated **at each time step** as a sequence is processed

Recurrent Neural Networks (RNNs)

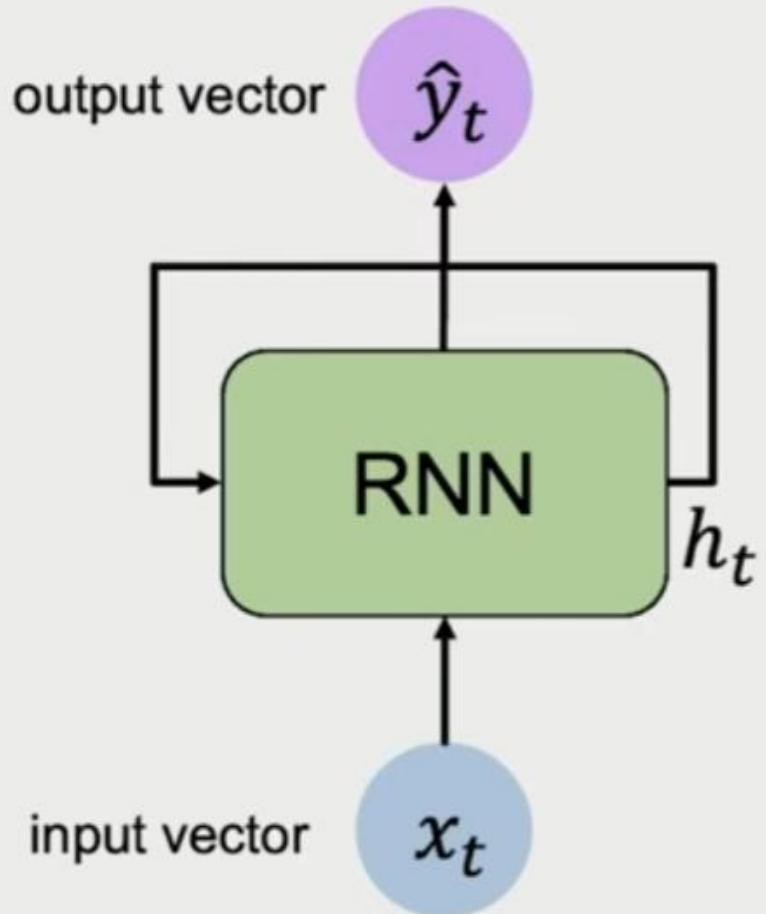


Apply a **recurrence relation** at every time step to process a sequence:

$$h_t = f_W(x_t, h_{t-1})$$

RNNs have a **state**, h_t , that is updated **at each time step** as a sequence is processed

Recurrent Neural Networks (RNNs)

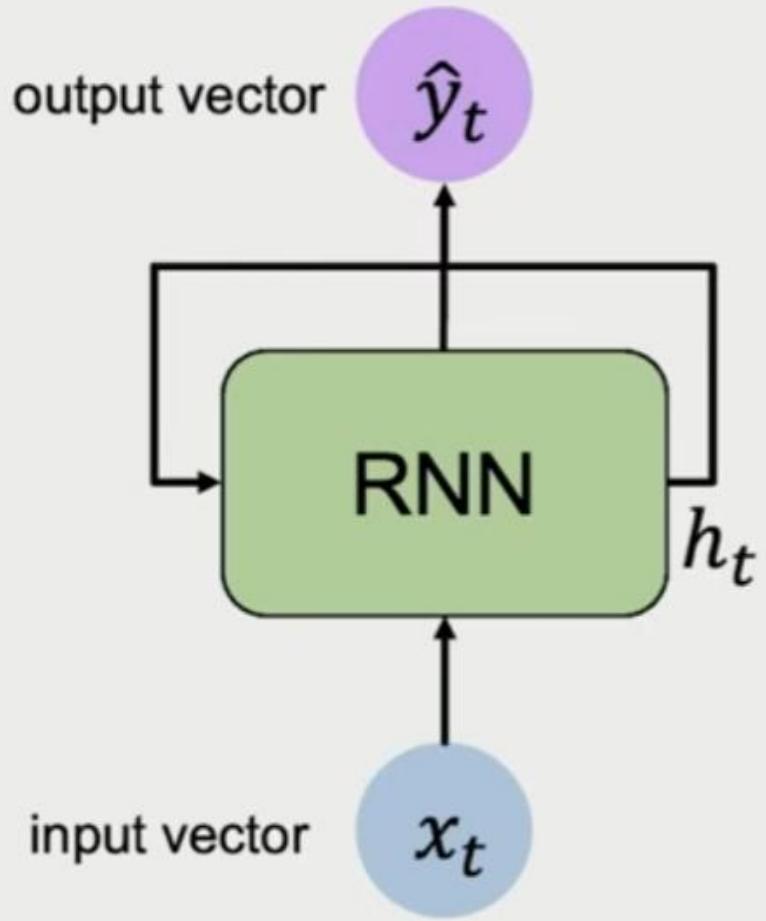


Apply a **recurrence relation** at every time step to process a sequence:

$$h_t = f_W(x_t, h_{t-1})$$

RNNs have a **state**, h_t , that is updated **at each time step** as a sequence is processed

Recurrent Neural Networks (RNNs)

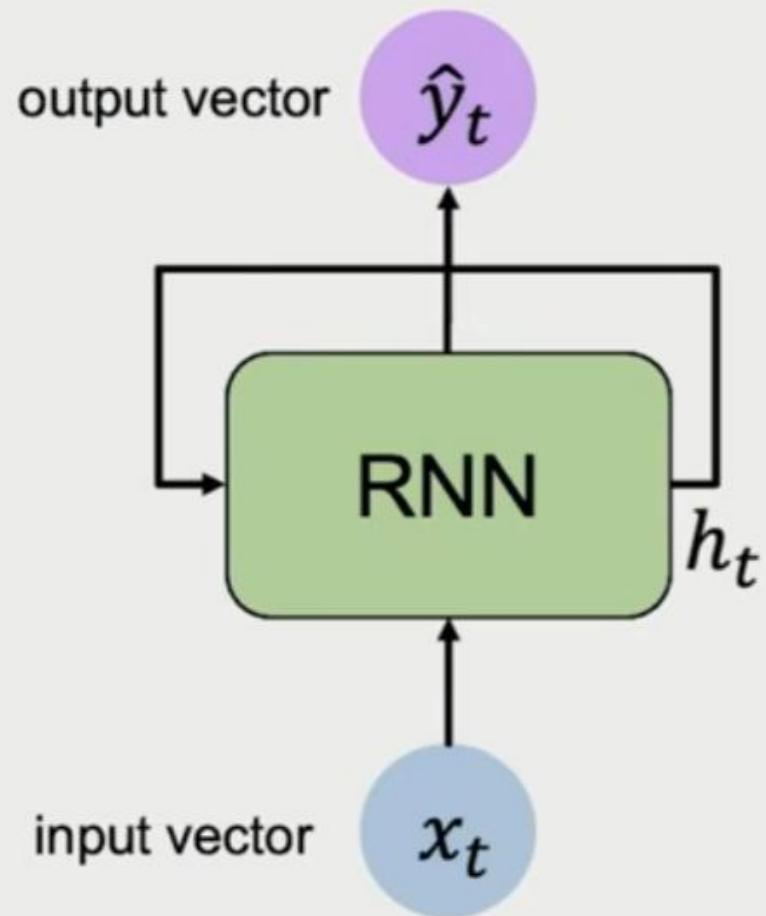


Apply a **recurrence relation** at every time step to process a sequence:

$$h_t = f_W(x_t, h_{t-1})$$

RNNs have a **state**, h_t , that is updated **at each time step** as a sequence is processed

Recurrent Neural Networks (RNNs)



Apply a **recurrence relation** at every time step to process a sequence:

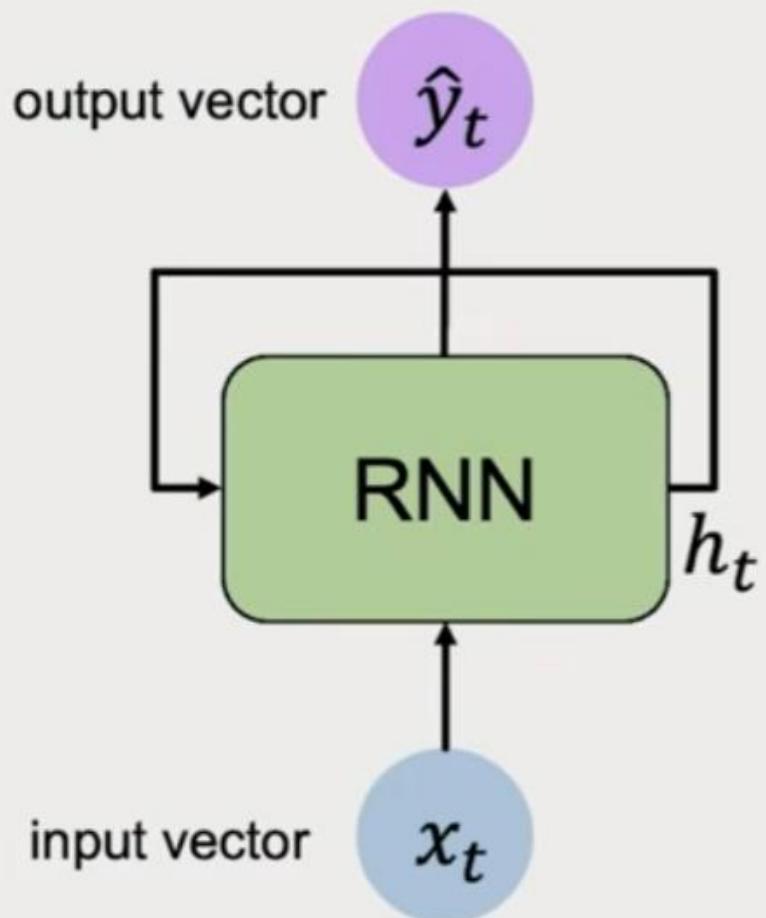
$$h_t = f_W(x_t, h_{t-1})$$

cell state function
 with weights
 w

input old state

RNNs have a **state**, h_t , that is updated **at each time step** as a sequence is processed

Recurrent Neural Networks (RNNs)



Apply a **recurrence relation** at every time step to process a sequence:

$$h_t = f_W(x_t, h_{t-1})$$

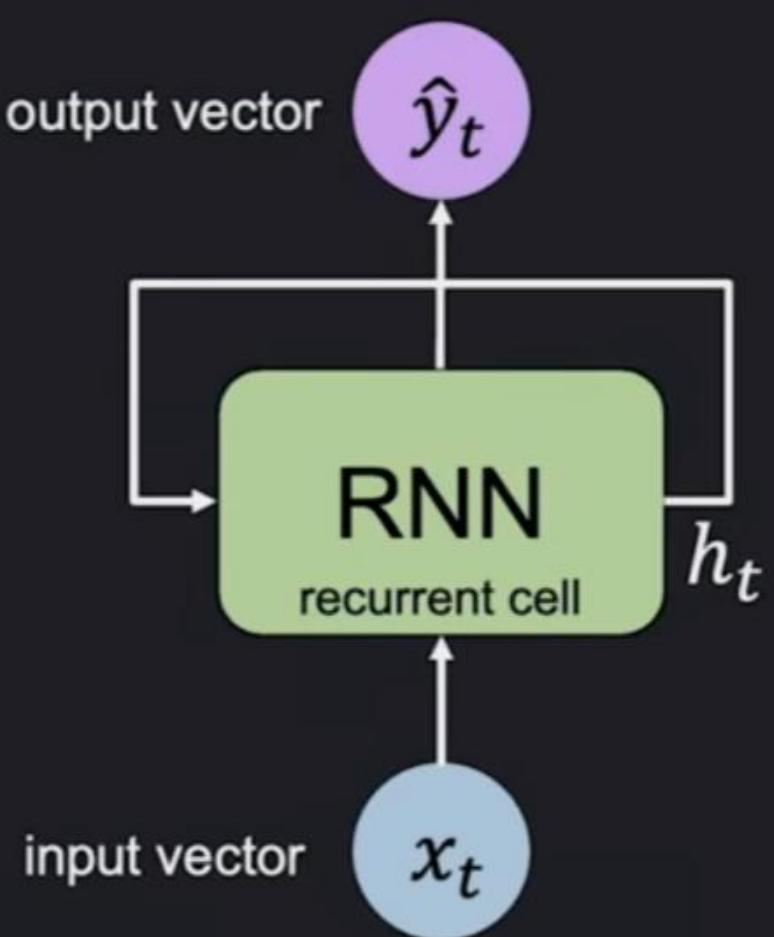
cell state function
 with weights
 W input old state

Note: the same function and set of parameters are used at every time step

RNNs have a **state**, h_t , that is updated **at each time step** as a sequence is processed

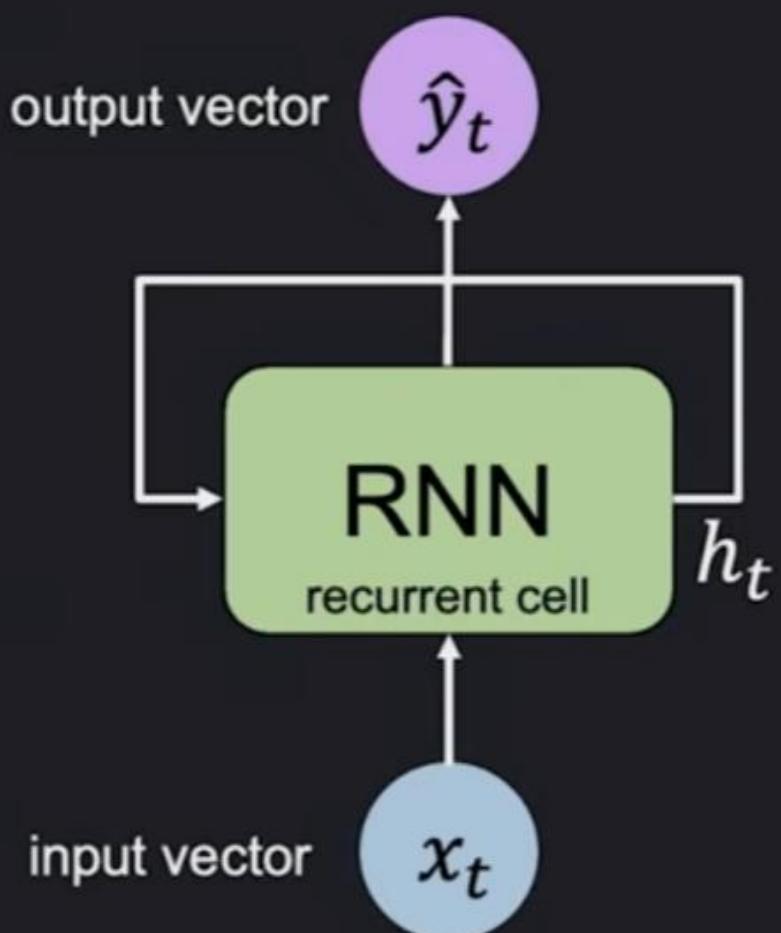
RNN Intuition

```
my_rnn = RNN()  
hidden_state = [0, 0, 0, 0]  
  
sentence = ["I", "love", "recurrent", "neural"]  
  
for word in sentence:  
    prediction, hidden_state = my_rnn(word, hidden_state)  
  
next_word_prediction = prediction  
# >>> "networks!"
```

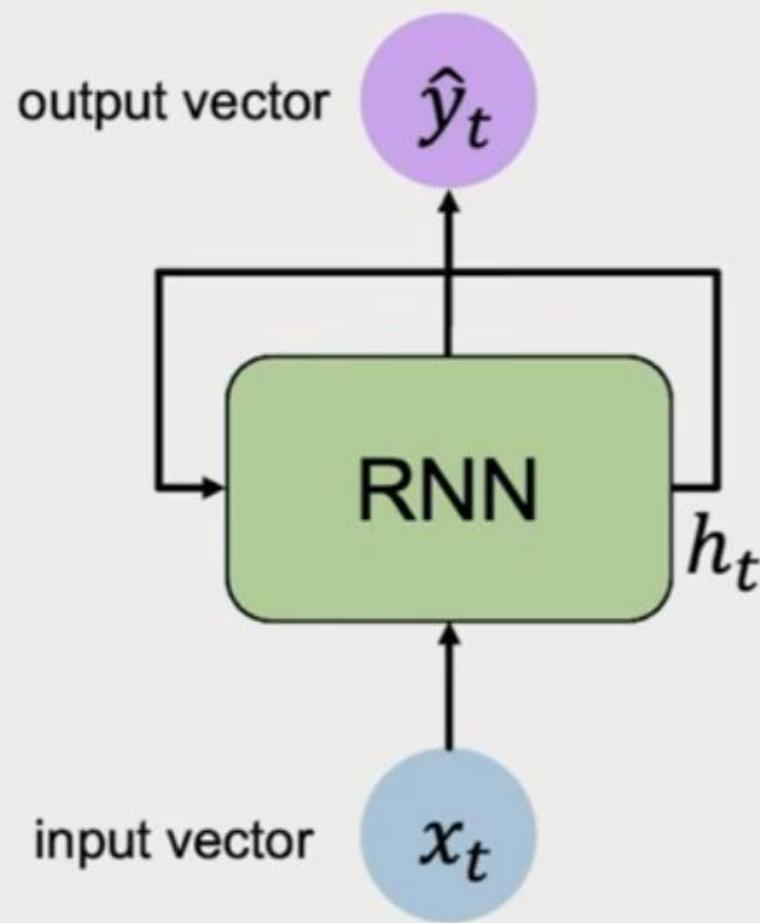


RNN Intuition

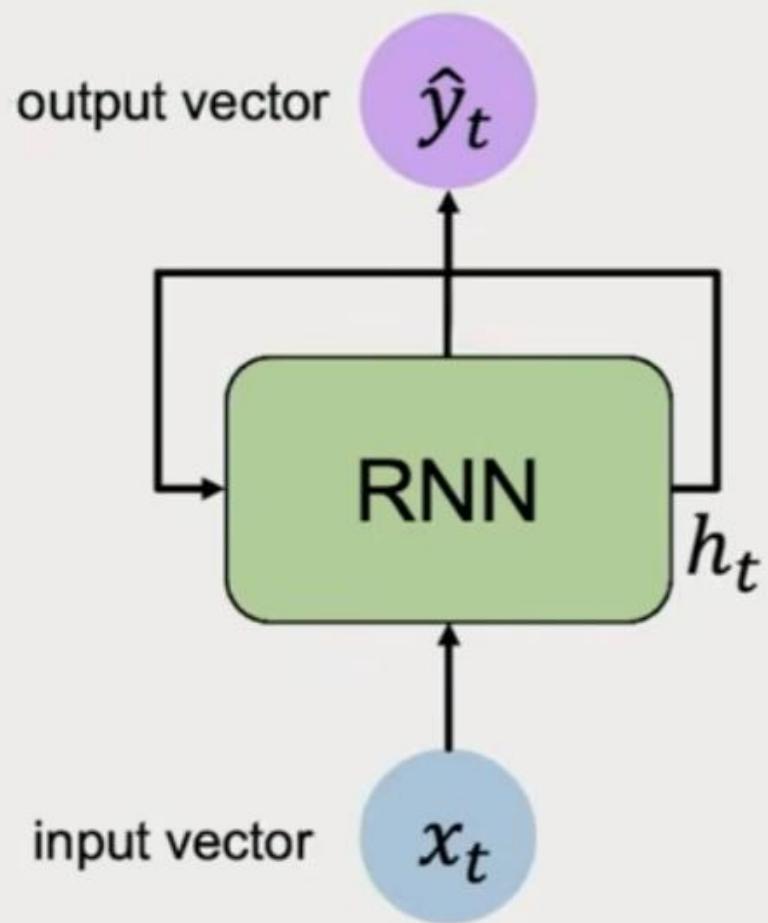
```
my_rnn = RNN()  
hidden_state = [0, 0, 0, 0]  
  
sentence = ["I", "love", "recurrent", "neural"]  
  
for word in sentence:  
    prediction, hidden_state = my_rnn(word, hidden_state)  
  
    next_word_prediction = prediction  
    # >>> "networks!"
```



RNN State Update and Output



RNN State Update and Output



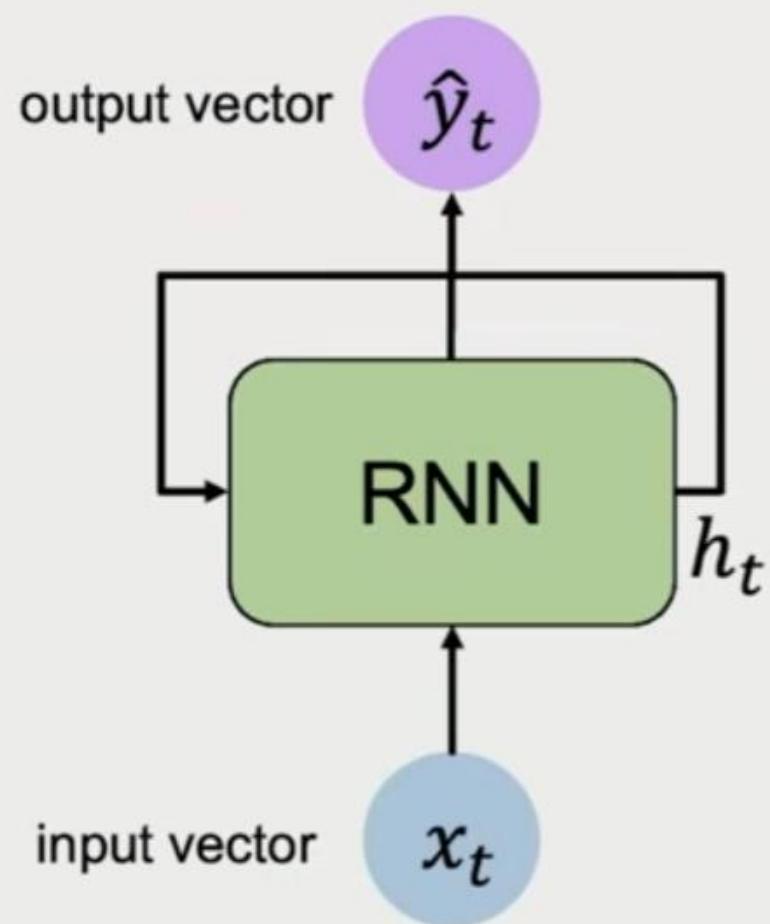
Update Hidden State

$$h_t = \tanh(W_{hh}^T h_{t-1} + W_{xh}^T x_t)$$

Input Vector

$$x_t$$

RNN State Update and Output



Output Vector

$$\hat{y}_t = \mathbf{W}_{hy}^T h_t$$

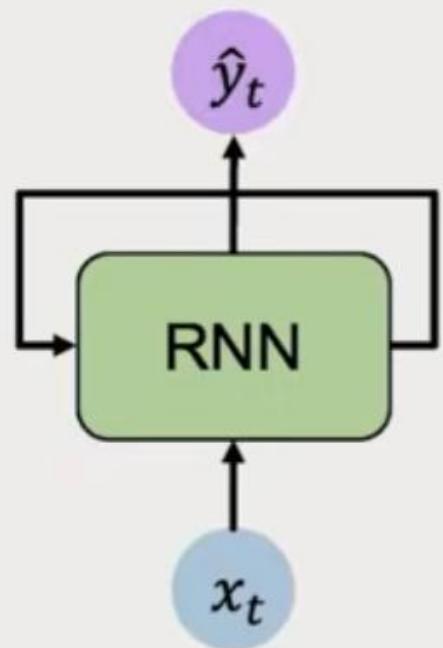
Update Hidden State

$$h_t = \tanh(\mathbf{W}_{hh}^T h_{t-1} + \mathbf{W}_{xh}^T x_t)$$

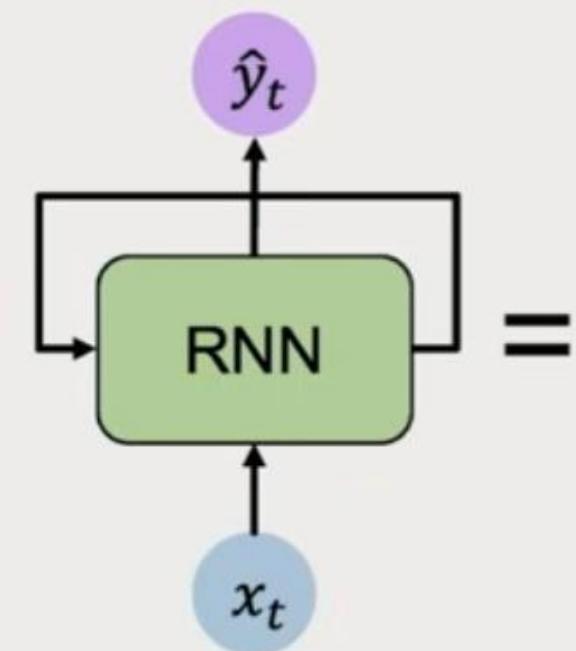
Input Vector

$$x_t$$

RNNs: Computational Graph Across Time

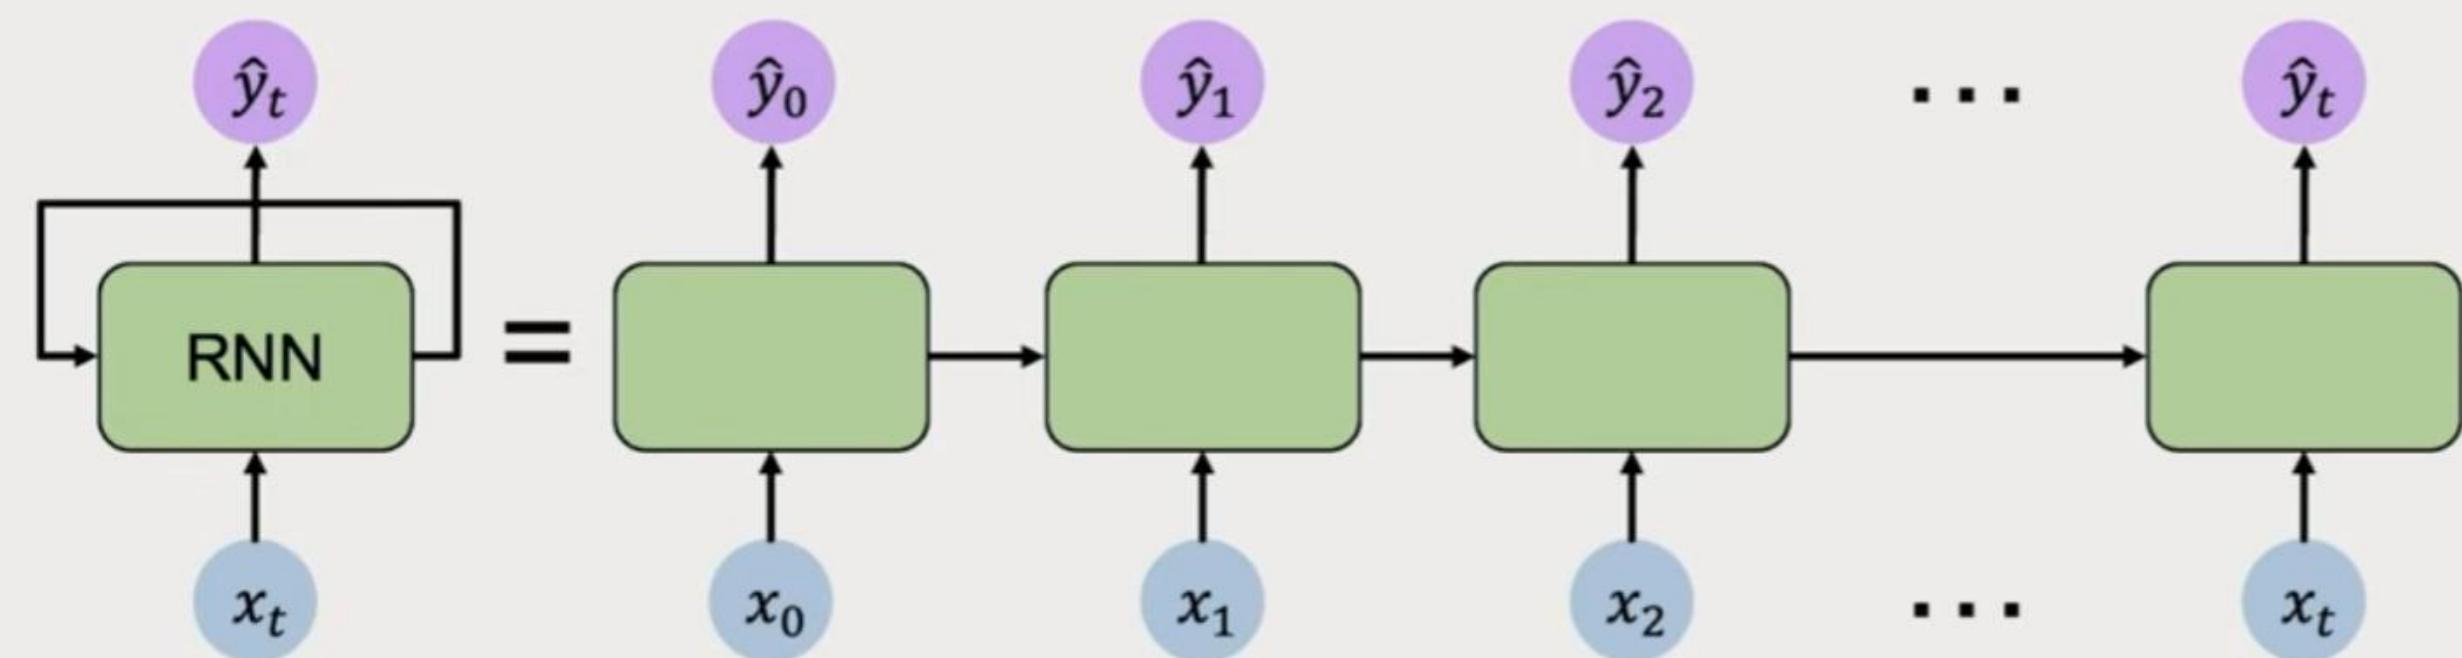


RNNs: Computational Graph Across Time

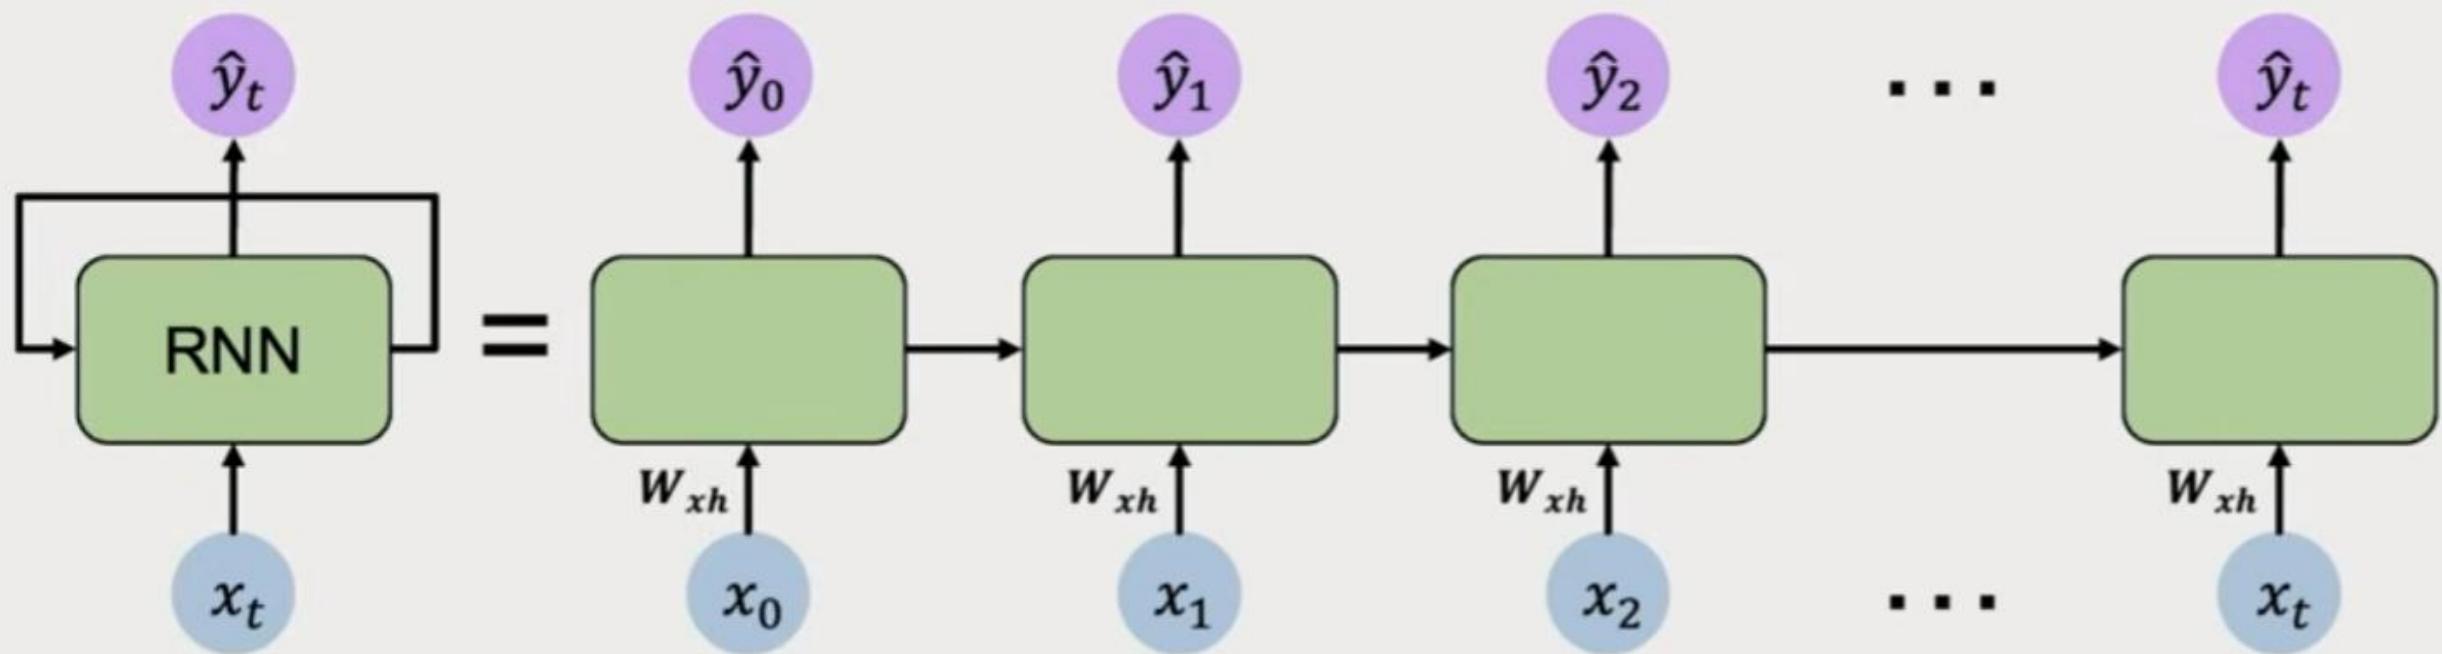


= Represent as computational graph unrolled across time

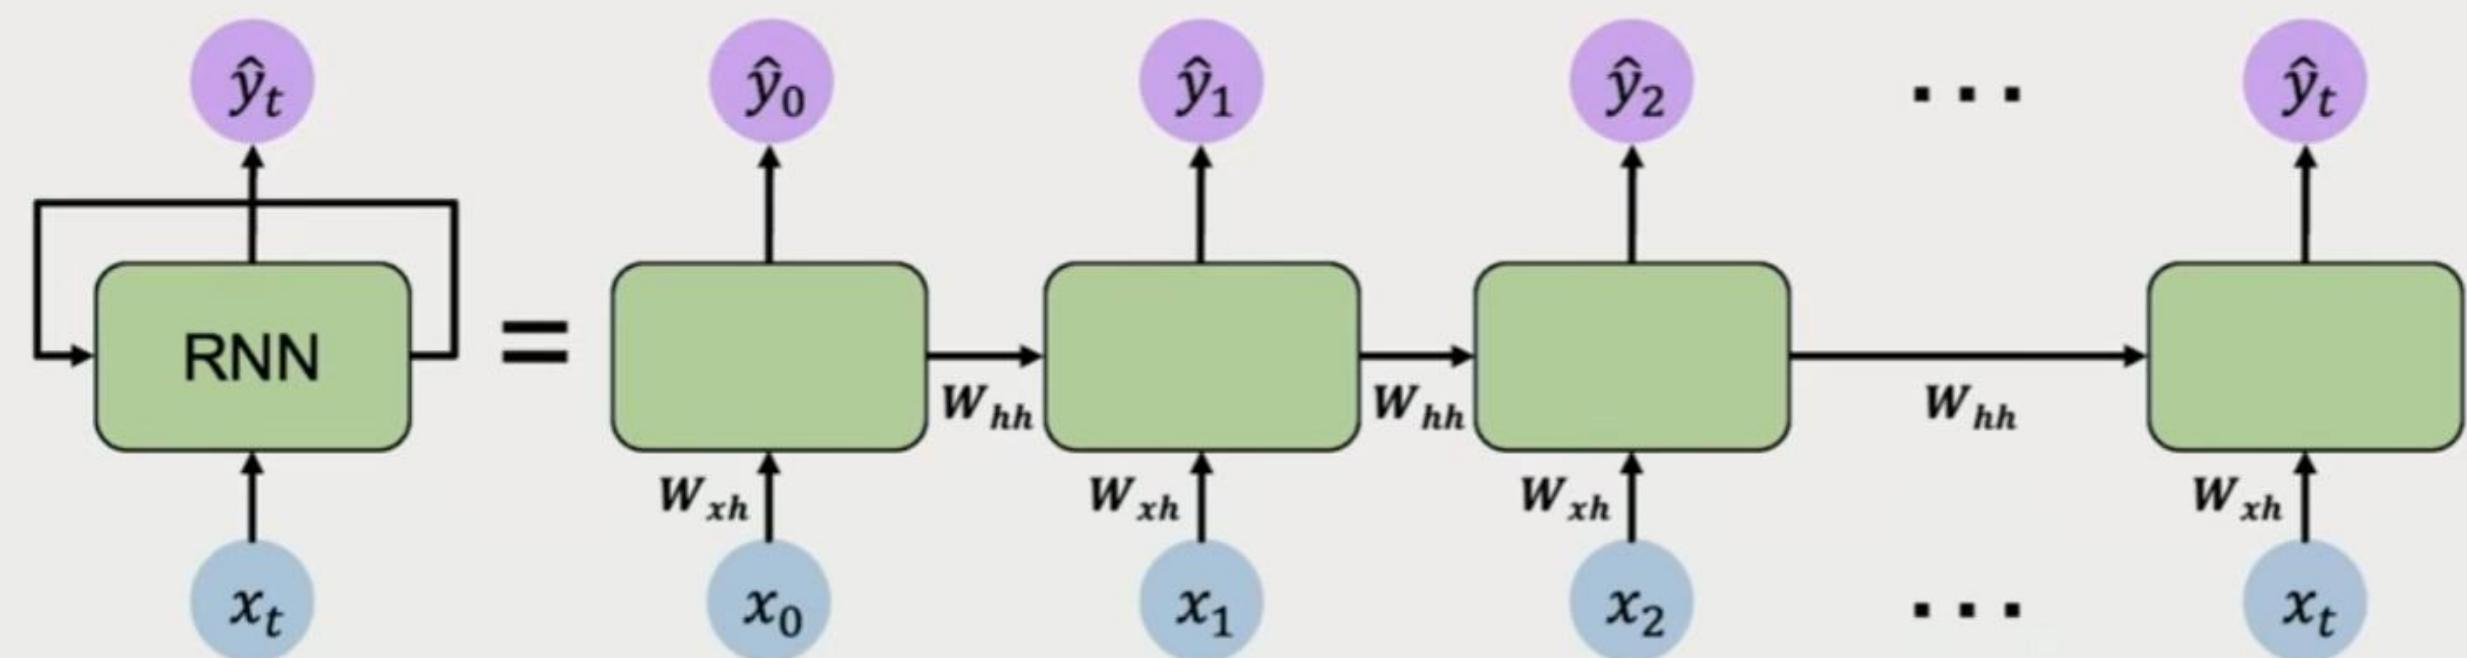
RNNs: Computational Graph Across Time



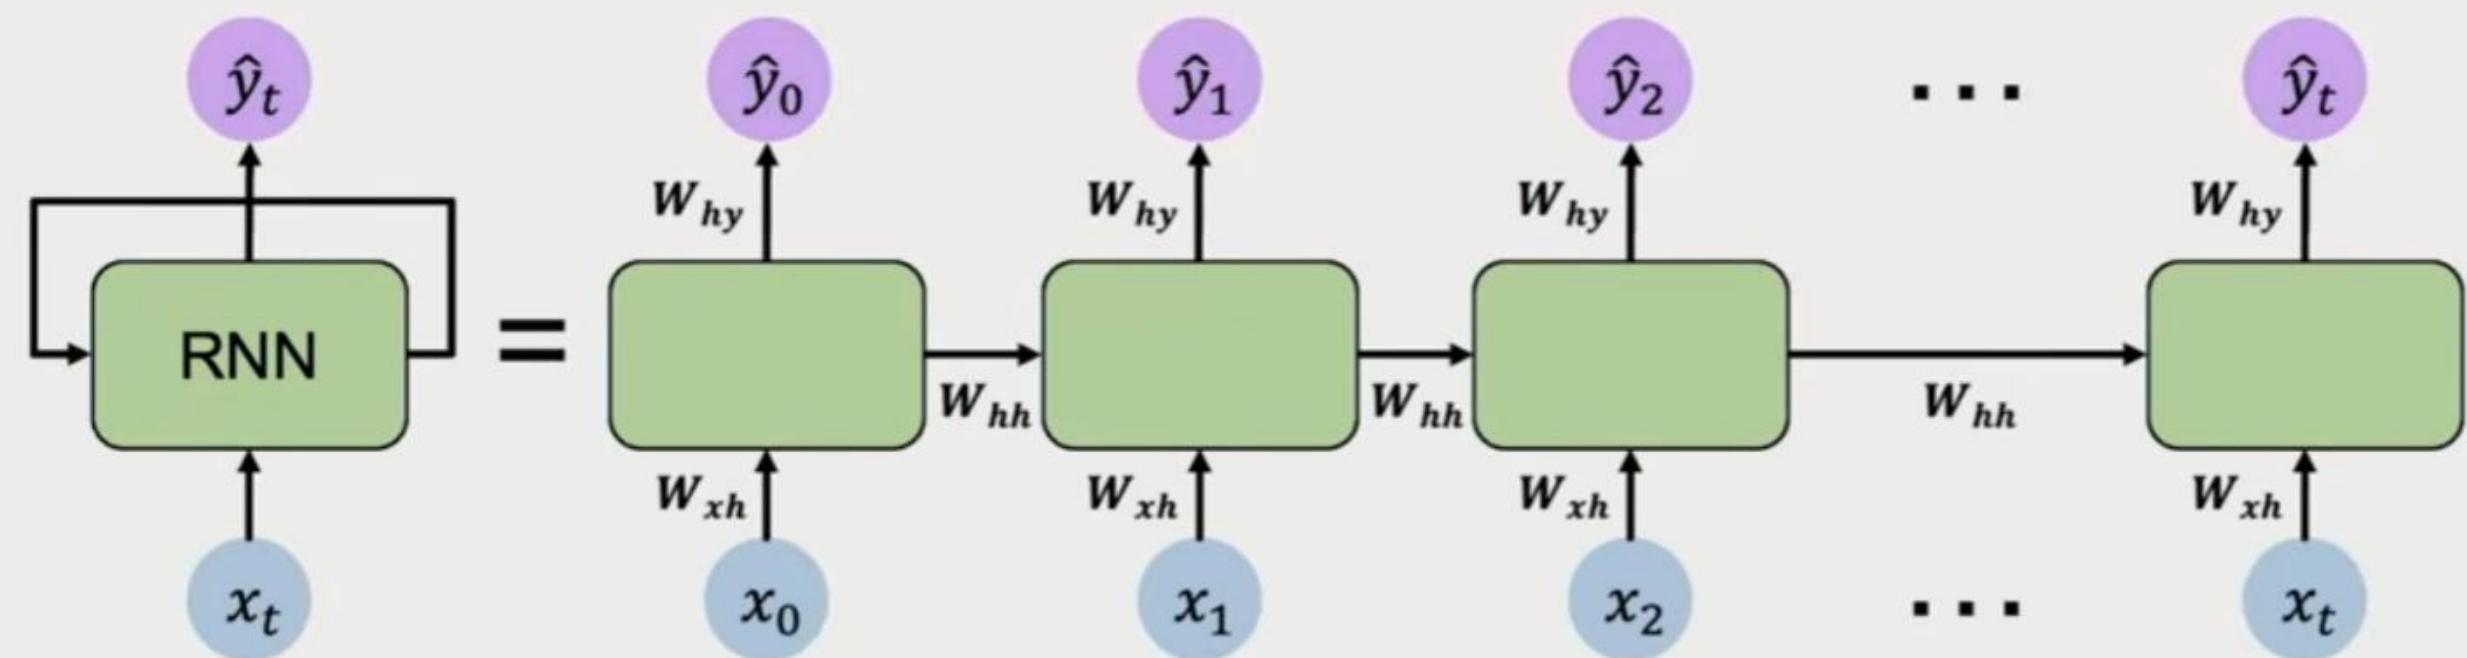
RNNs: Computational Graph Across Time



RNNs: Computational Graph Across Time

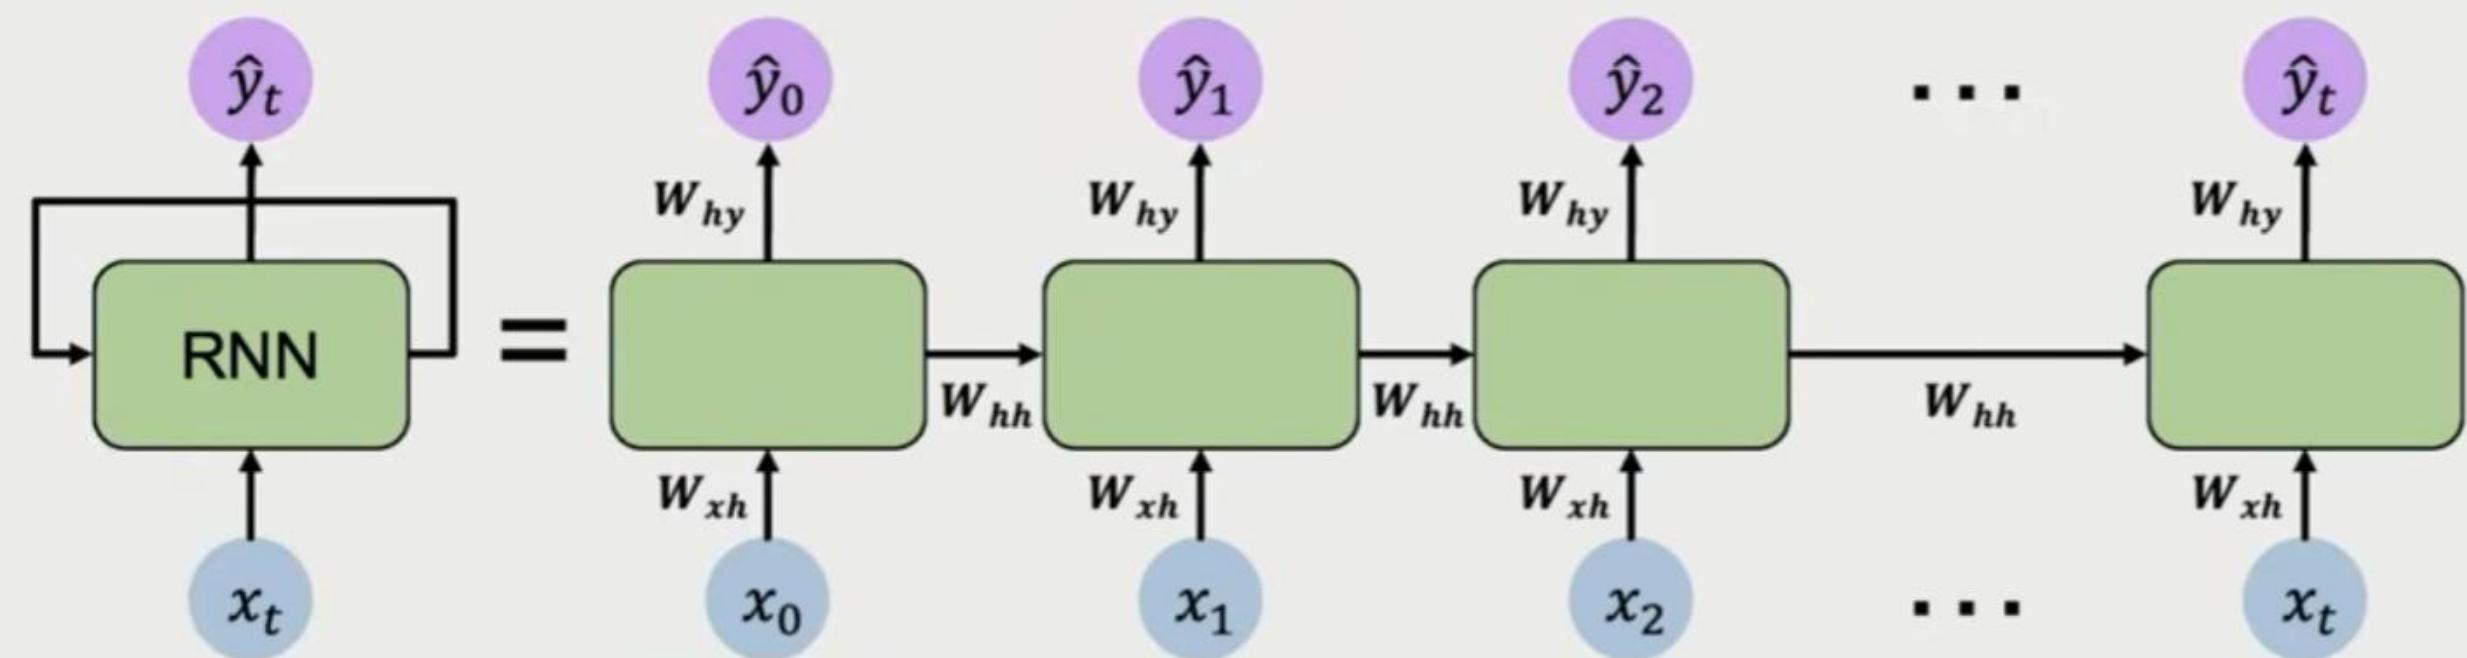


RNNs: Computational Graph Across Time



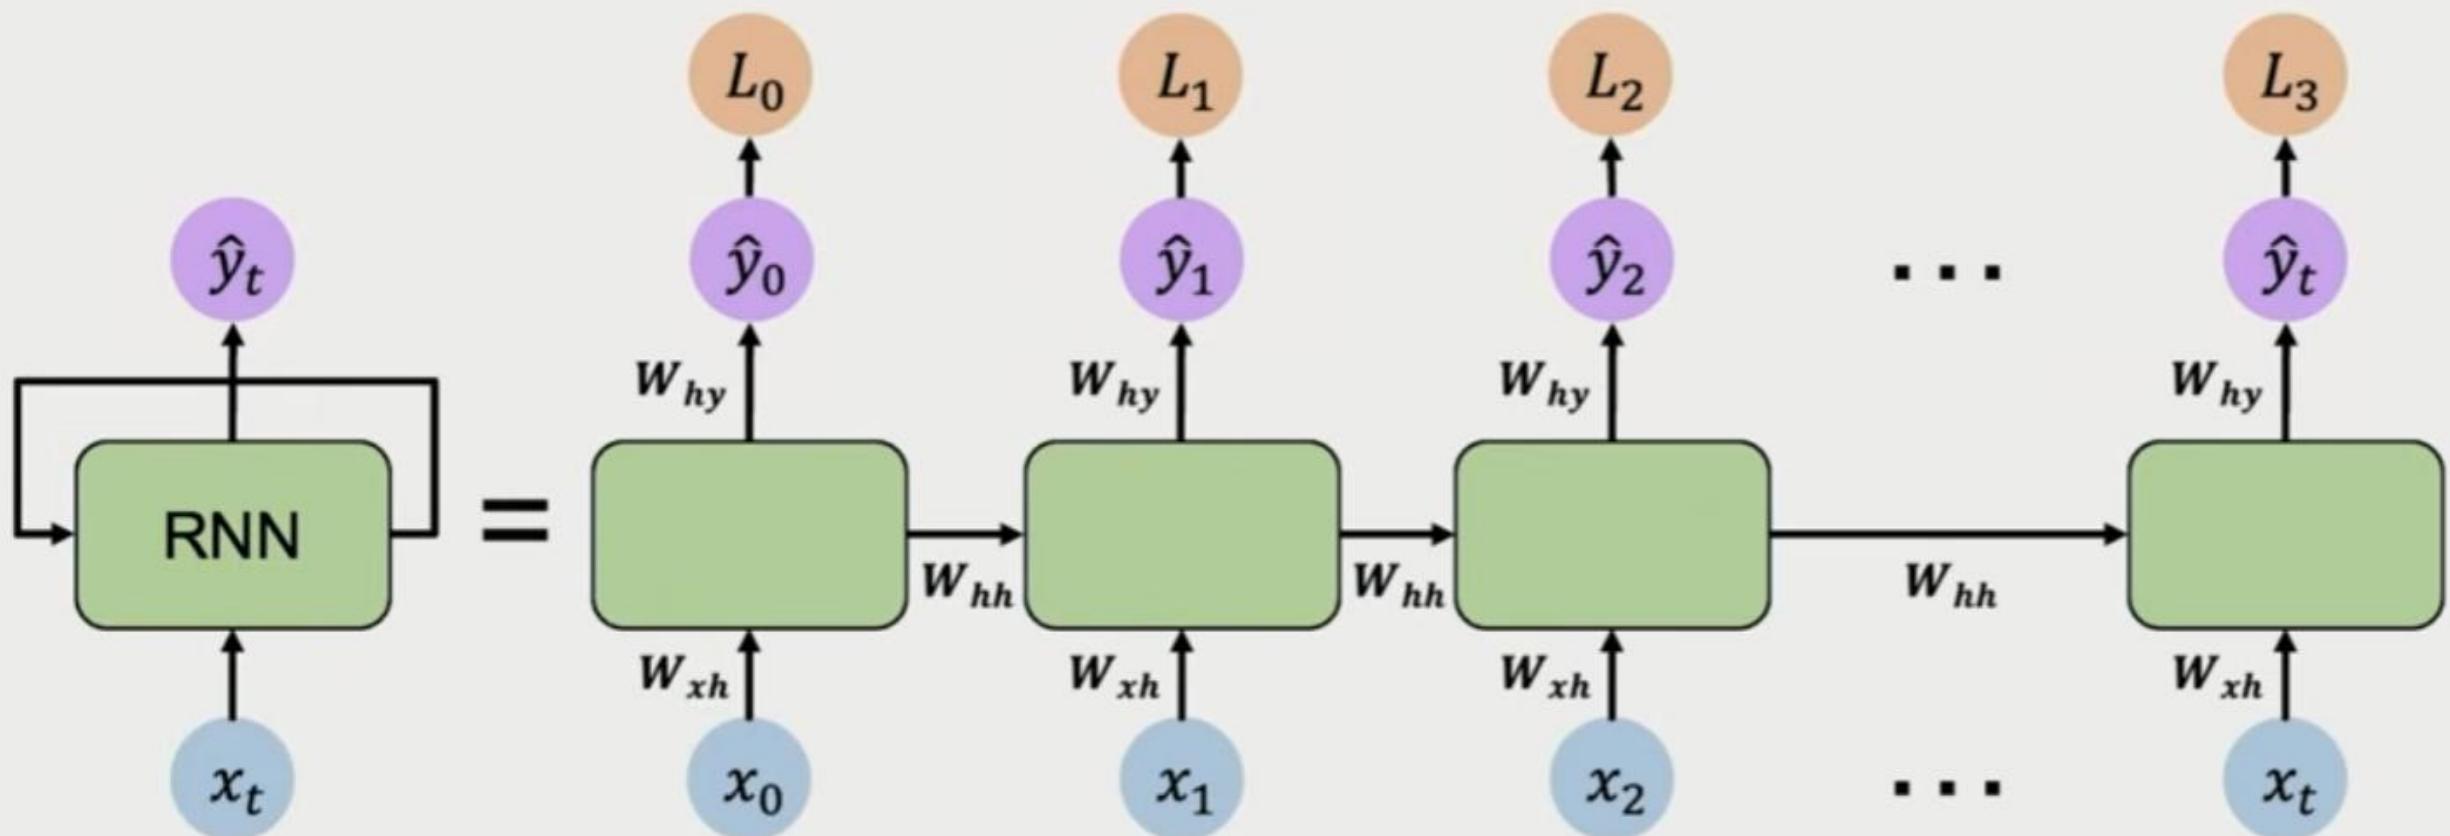
RNNs: Computational Graph Across Time

Re-use the **same weight matrices** at every time step



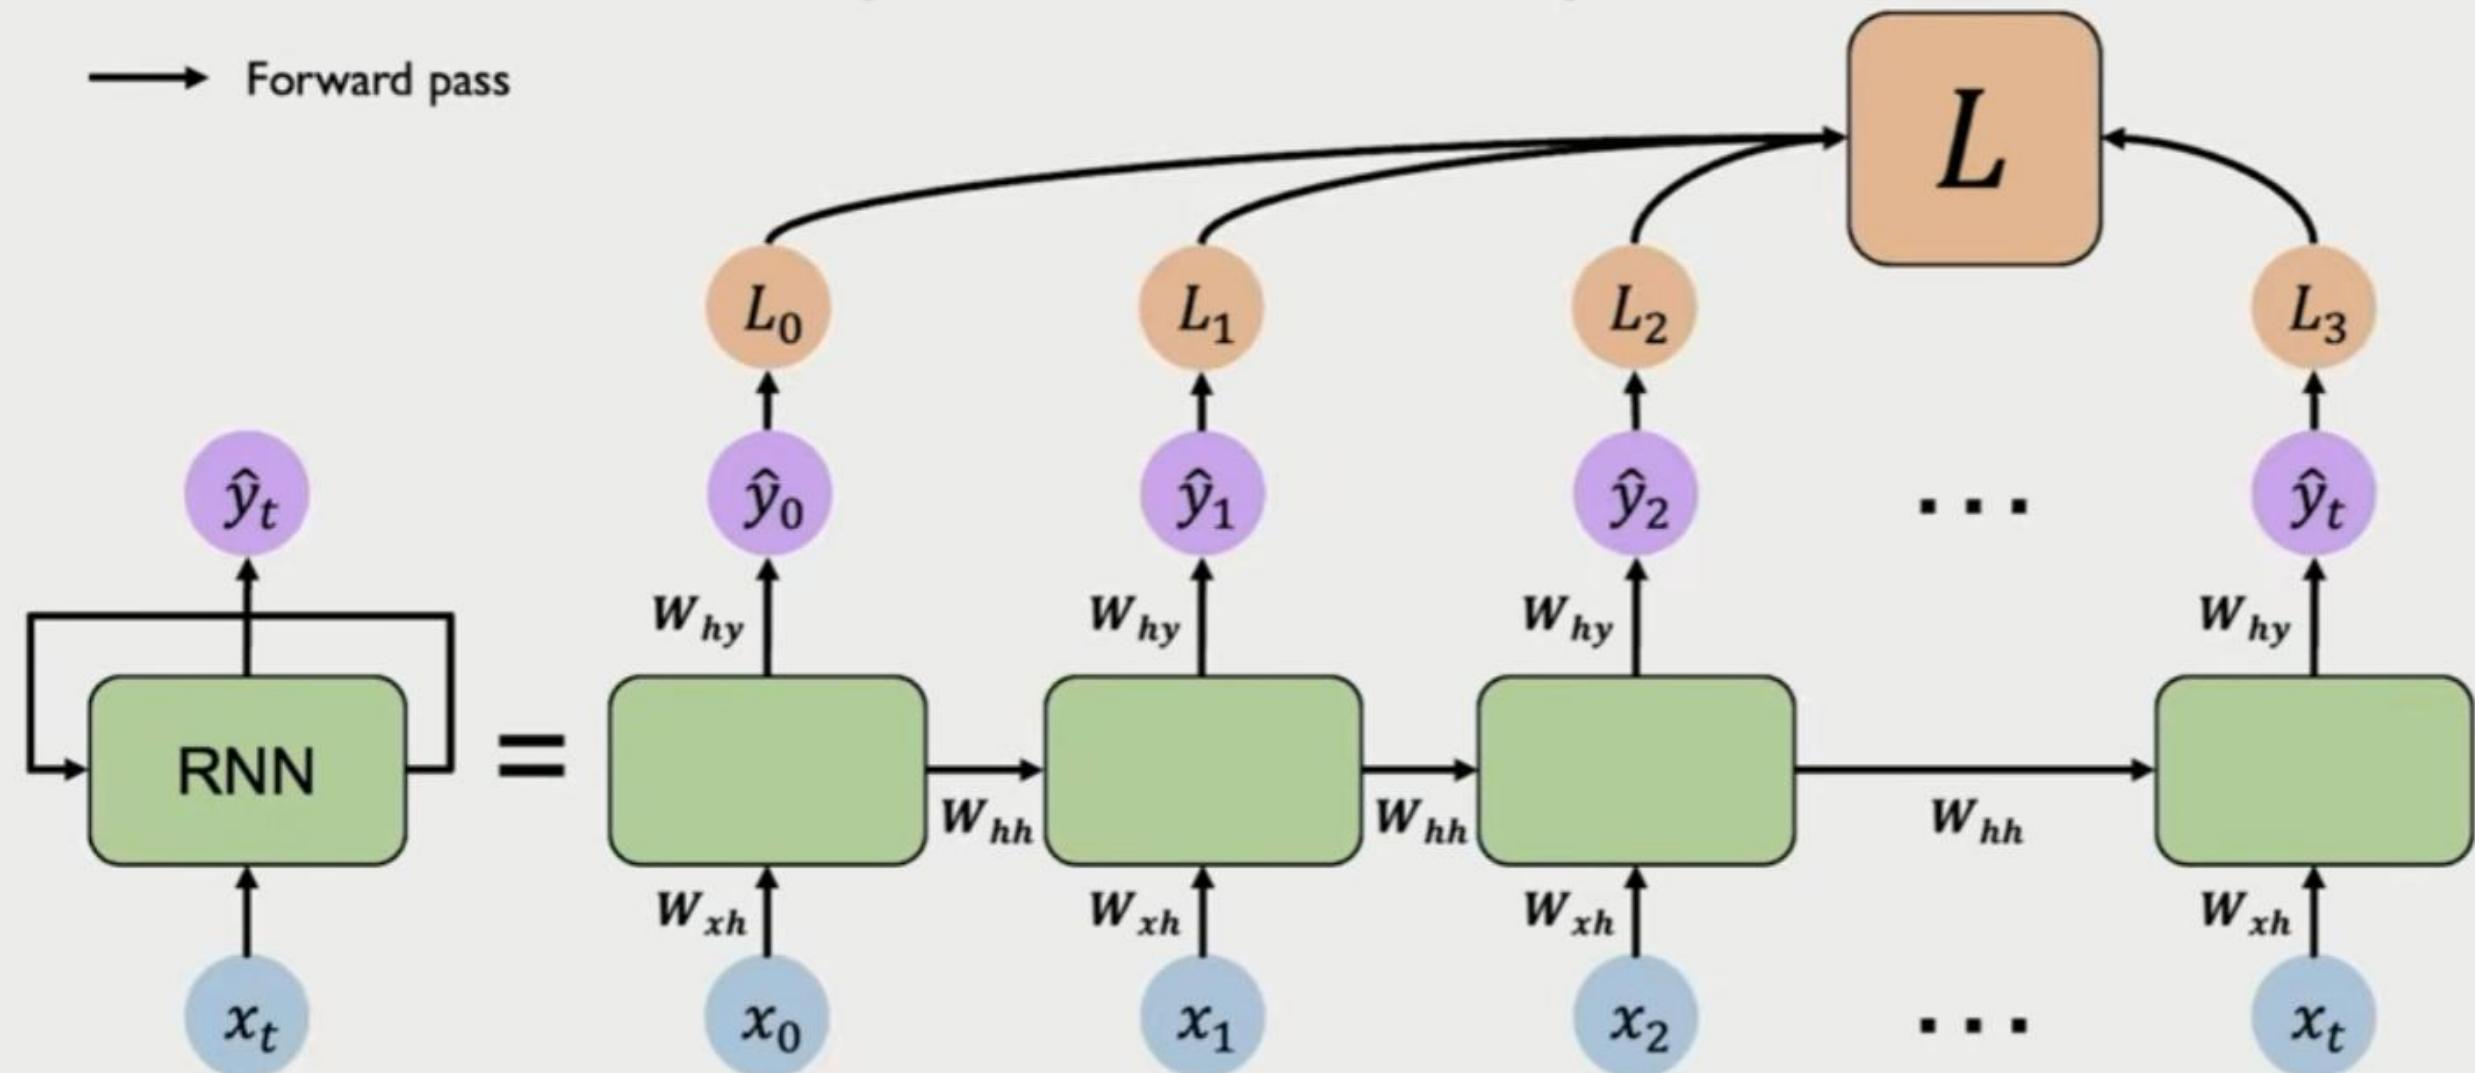
RNNs: Computational Graph Across Time

→ Forward pass



RNNs: Computational Graph Across Time

→ Forward pass



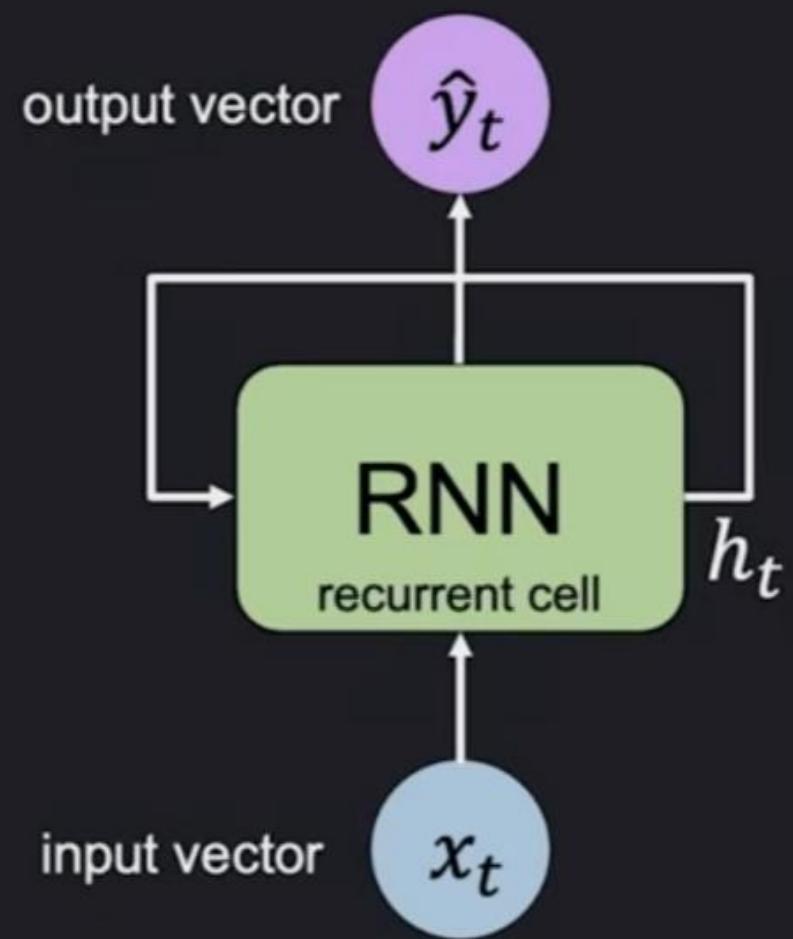


RNNs from Scratch in TensorFlow

```
class MyRNNCell(tf.keras.layers.Layer):
    def __init__(self, rnn_units, input_dim, output_dim):
        super(MyRNNCell, self).__init__()

        # Initialize weight matrices
        self.W_xh = self.add_weight([rnn_units, input_dim])
        self.W_hh = self.add_weight([rnn_units, rnn_units])
        self.W_hy = self.add_weight([output_dim, rnn_units])

        # Initialize hidden state to zeros
        self.h = tf.zeros([rnn_units, 1])
```





RNNs from Scratch in TensorFlow

```
class MyRNNCell(tf.keras.layers.Layer):
    def __init__(self, rnn_units, input_dim, output_dim):
        super(MyRNNCell, self).__init__()

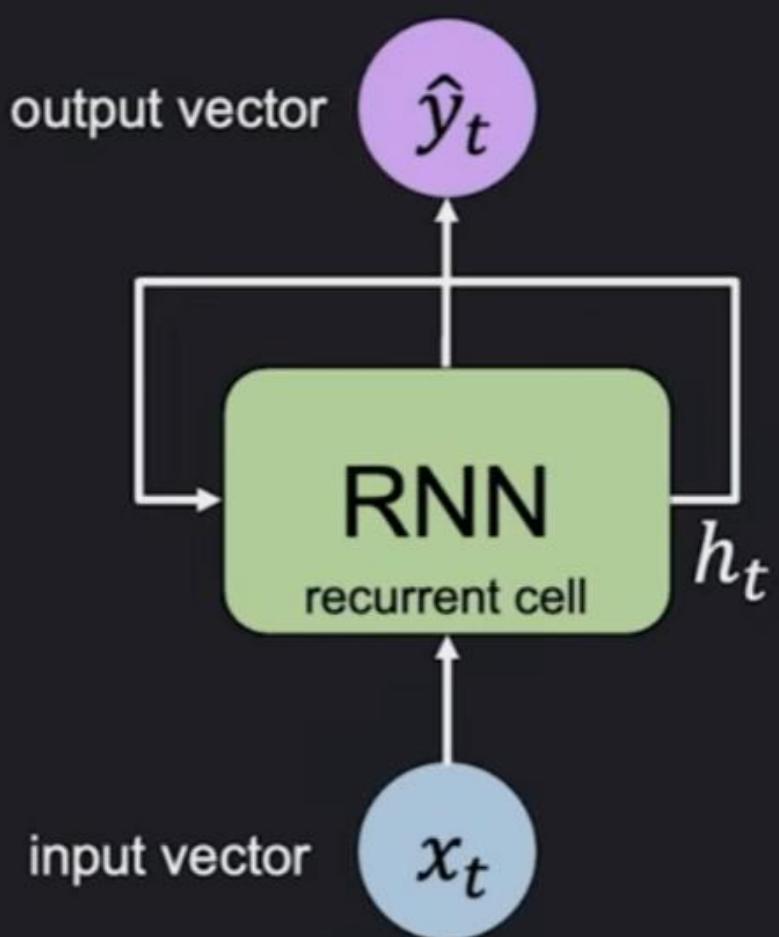
        # Initialize weight matrices
        self.W_xh = self.add_weight([rnn_units, input_dim])
        self.W_hh = self.add_weight([rnn_units, rnn_units])
        self.W_hy = self.add_weight([output_dim, rnn_units])

        # Initialize hidden state to zeros
        self.h = tf.zeros([rnn_units, 1])

    def call(self, x):
        # Update the hidden state
        self.h = tf.math.tanh( self.W_hh * self.h + self.W_xh * x )

        # Compute the output
        output = self.W_hy * self.h

        # Return the current output and hidden state
        return output, self.h
```



RNNs from Scratch in TensorFlow

```
class MyRNNCell(tf.keras.layers.Layer):
    def __init__(self, rnn_units, input_dim, output_dim):
        super(MyRNNCell, self).__init__()

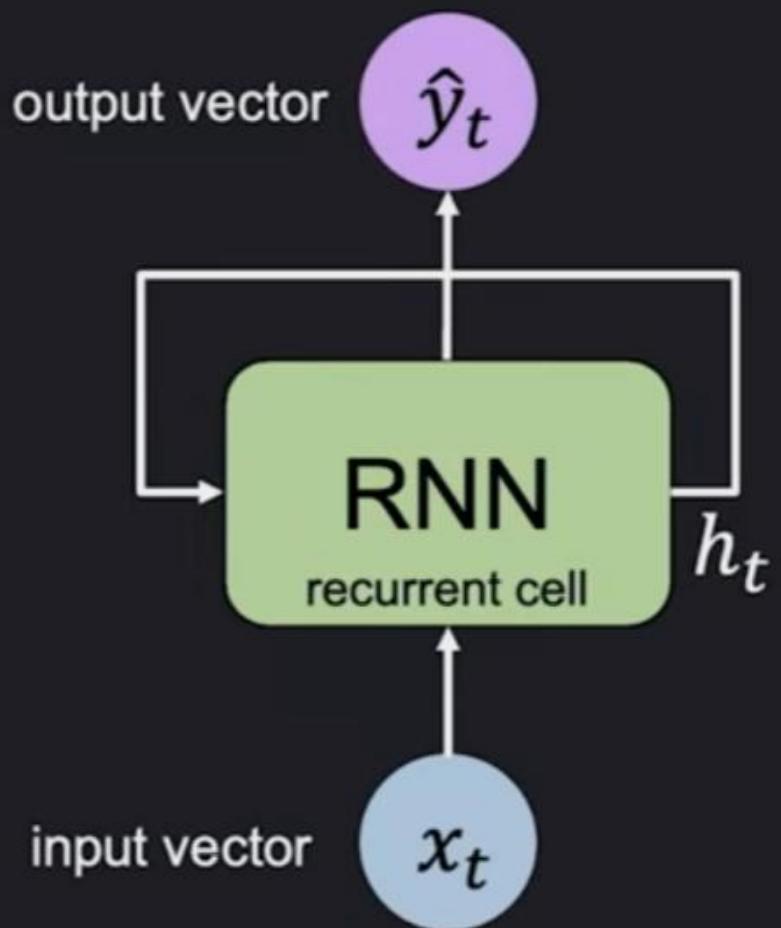
        # Initialize weight matrices
        self.W_xh = self.add_weight([rnn_units, input_dim])
        self.W_hh = self.add_weight([rnn_units, rnn_units])
        self.W_hy = self.add_weight([output_dim, rnn_units])

        # Initialize hidden state to zeros
        self.h = tf.zeros([rnn_units, 1])

    def call(self, x):
        # Update the hidden state
        self.h = tf.math.tanh( self.W_hh * self.h + self.W_xh * x )

        # Compute the output
        output = self.W_hy * self.h

        # Return the current output and hidden state
        return output, self.h
```





RNNs from Scratch in TensorFlow

```
class MyRNNCell(tf.keras.layers.Layer):
    def __init__(self, rnn_units, input_dim, output_dim):
        super(MyRNNCell, self).__init__()

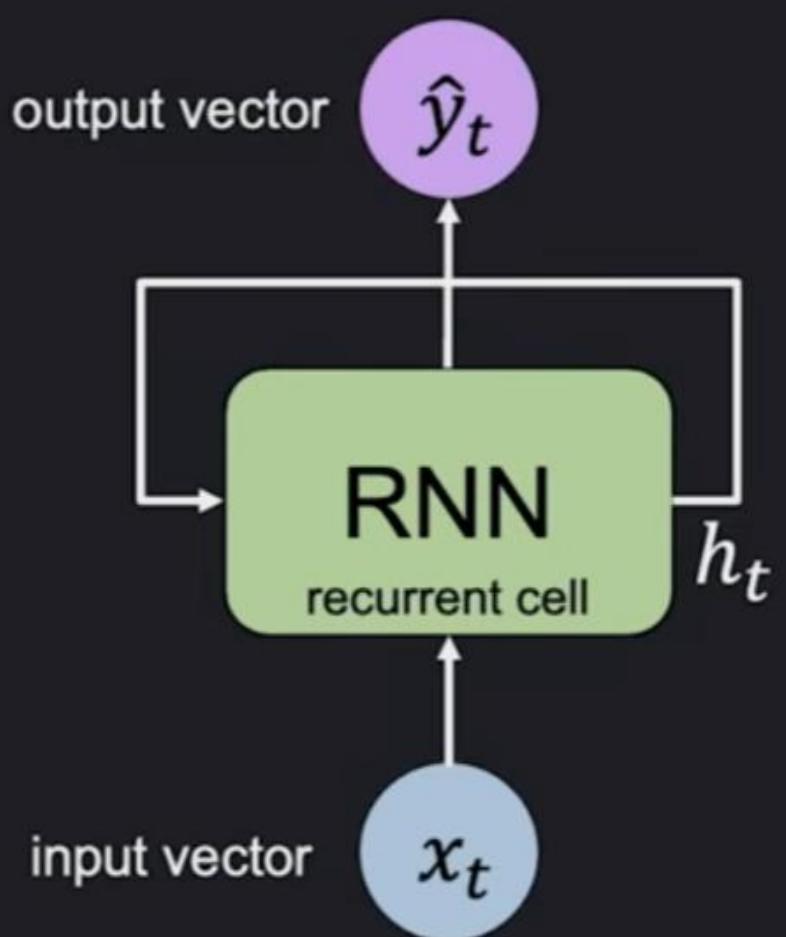
        # Initialize weight matrices
        self.W_xh = self.add_weight([rnn_units, input_dim])
        self.W_hh = self.add_weight([rnn_units, rnn_units])
        self.W_hy = self.add_weight([output_dim, rnn_units])

        # Initialize hidden state to zeros
        self.h = tf.zeros([rnn_units, 1])

    def call(self, x):
        # Update the hidden state
        self.h = tf.math.tanh( self.W_hh * self.h + self.W_xh * x )

        # Compute the output
        output = self.W_hy * self.h

    # Return the current output and hidden state
    return output, self.h
```





RNNs from Scratch in TensorFlow

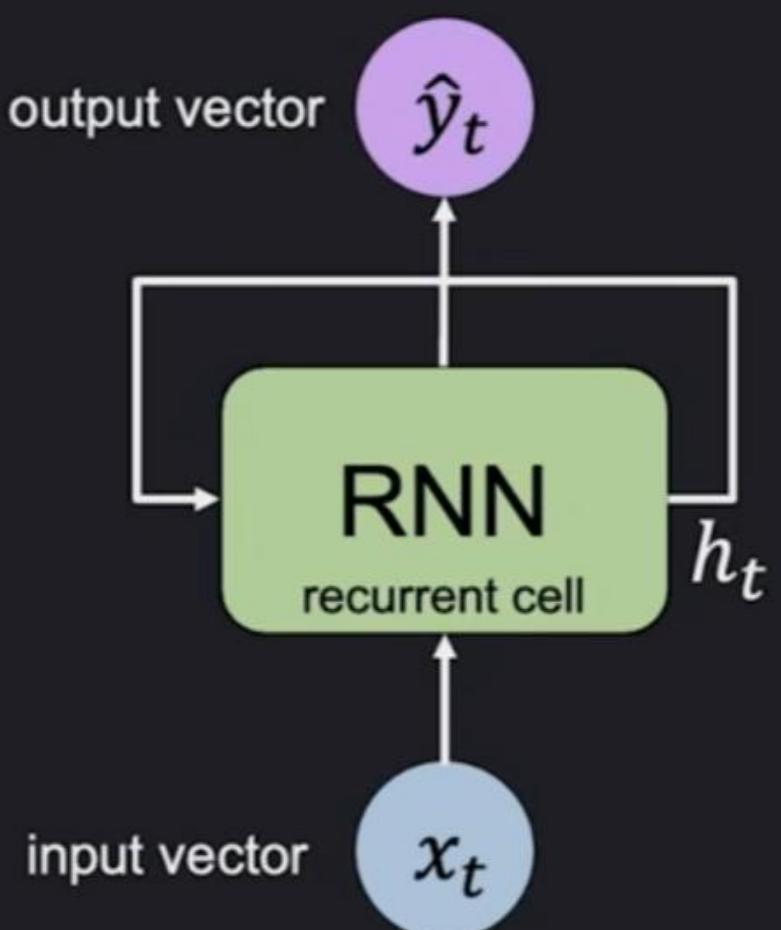
```
class MyRNNCell(tf.keras.layers.Layer):
    def __init__(self, rnn_units, input_dim, output_dim):
        super(MyRNNCell, self).__init__()
        # Initialize weight matrices
        self.W_xh = self.add_weight([rnn_units, input_dim])
        self.W_hh = self.add_weight([rnn_units, rnn_units])
        self.W_hy = self.add_weight([output_dim, rnn_units])

        # Initialize hidden state to zeros
        self.h = tf.zeros([rnn_units, 1])

    def call(self, x):
        # Update the hidden state
        self.h = tf.math.tanh( self.W_hh * self.h + self.W_xh * x )

        # Compute the output
        output = self.W_hy * self.h

        # Return the current output and hidden state
        return output, self.h
```

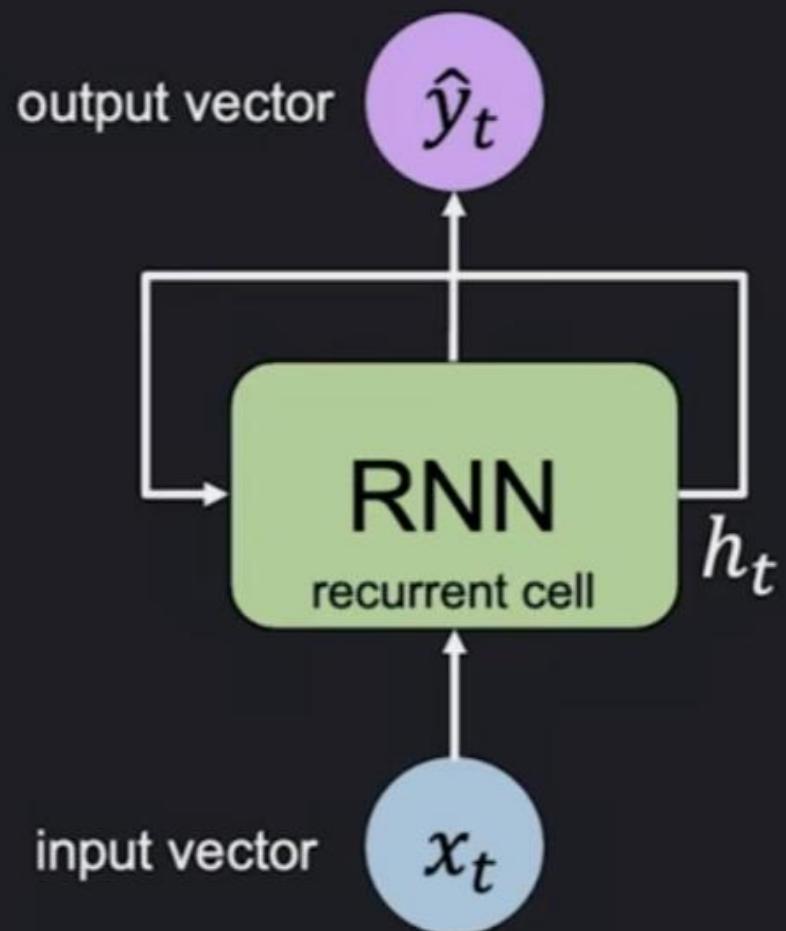


RNN Implementation: TensorFlow & PyTorch

```
from tf.keras.layers import SimpleRNN  
model = SimpleRNN(rnn_units)
```



```
from torch.nn import RNN  
model = RNN(input_size, rnn_units)
```

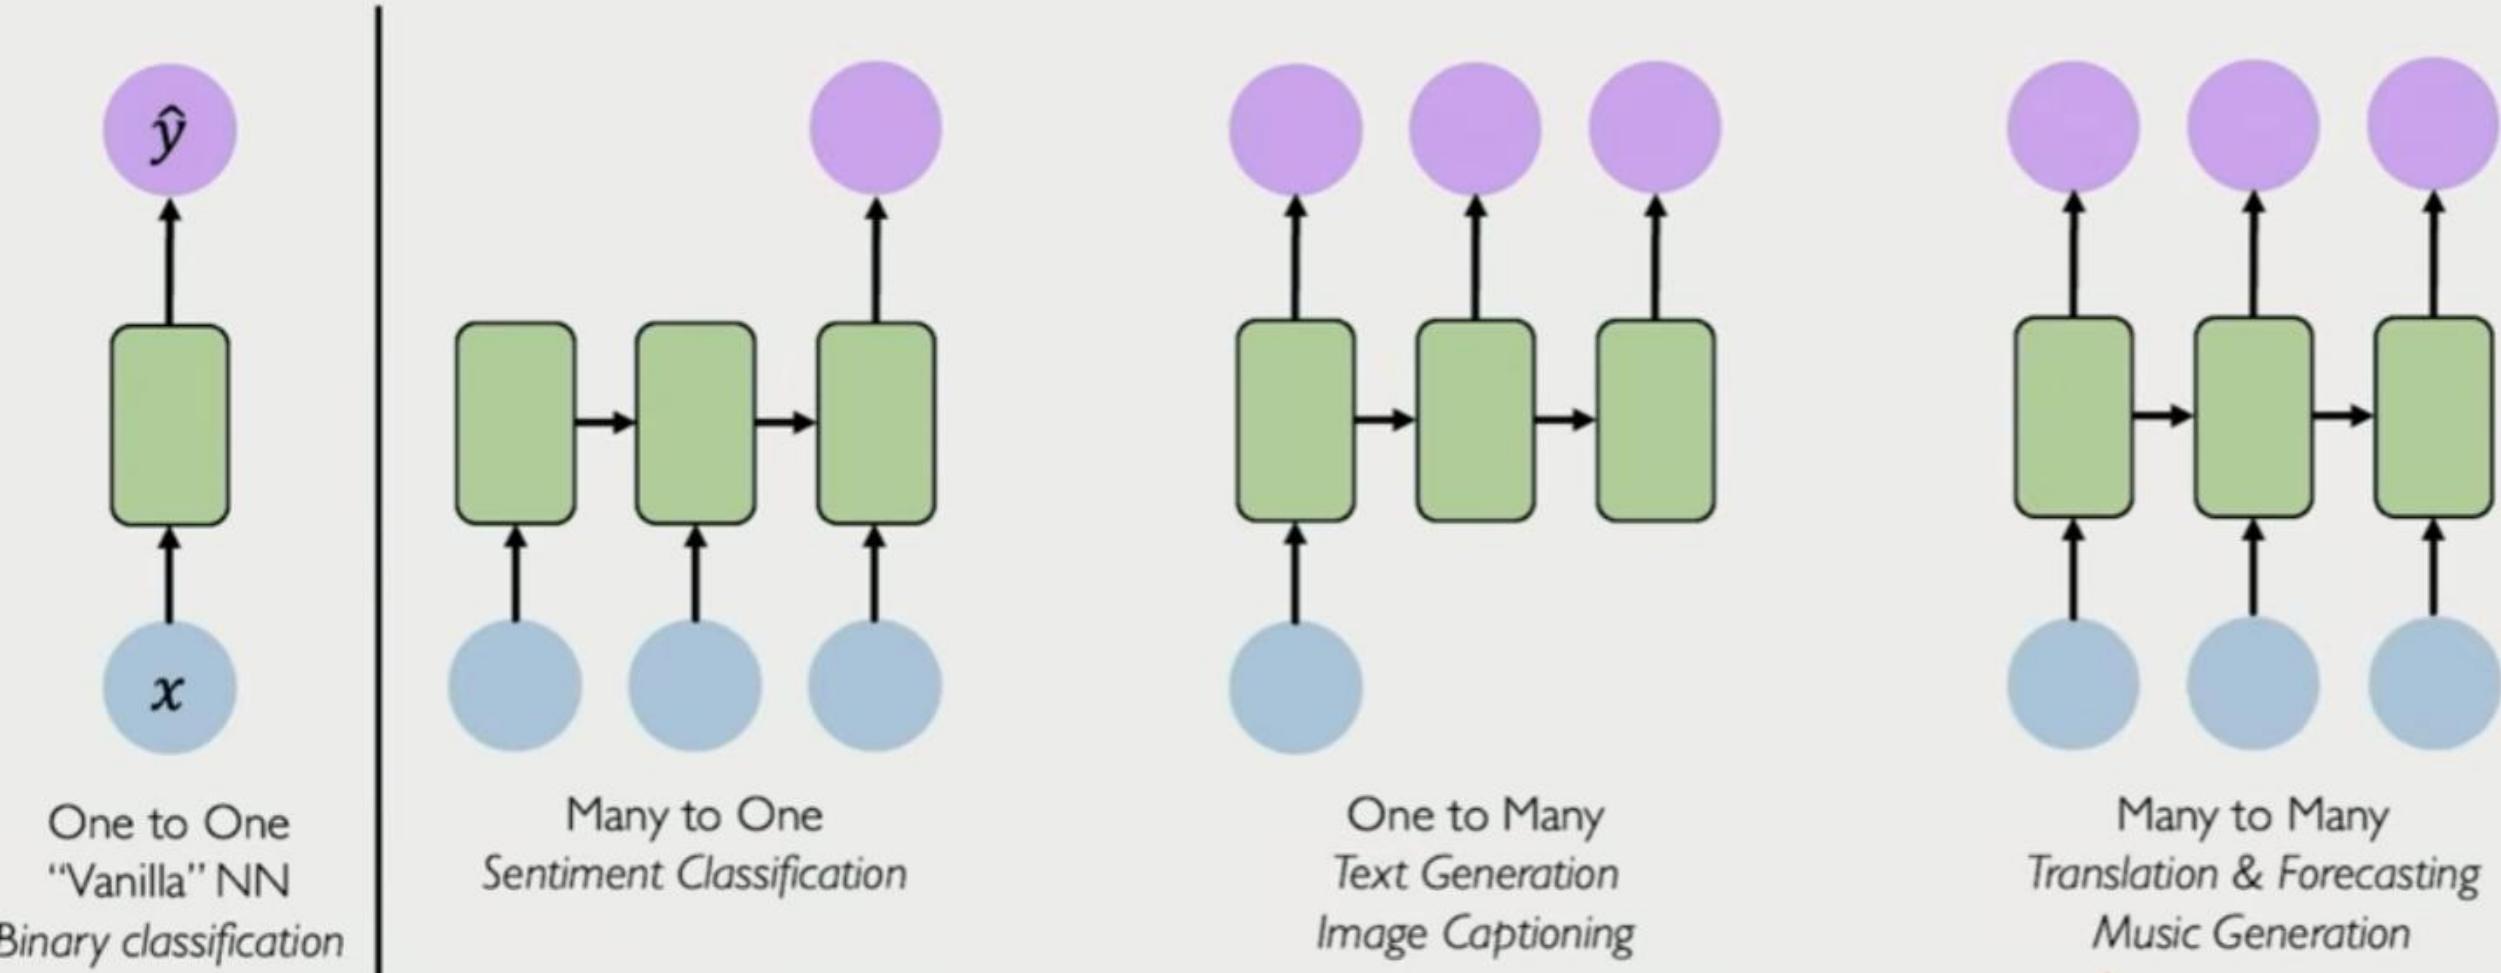


RNNs for Sequence Modeling



One to One
"Vanilla" NN
Binary classification

RNNs for Sequence Modeling



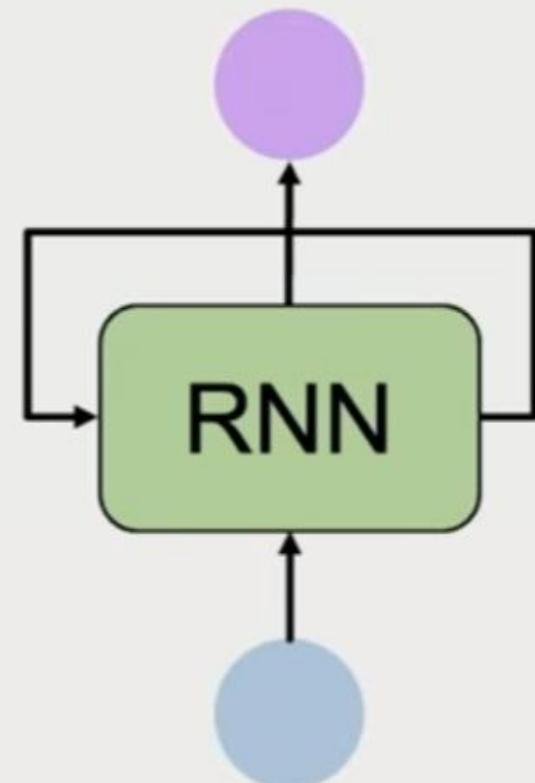
Sequence Modeling: Design Criteria

To model sequences, we need to:

Sequence Modeling: Design Criteria

To model sequences, we need to:

1. Handle **variable-length** sequences
2. Track **long-term** dependencies
3. Maintain information about **order**
4. **Share parameters** across the sequence



Recurrent Neural Networks (RNNs) meet
these sequence modeling design criteria

A Sequence Modeling Problem: Predict the Next Word

A Sequence Modeling Problem: Predict the Next Word

“This morning I took my cat for a walk.”

A Sequence Modeling Problem: Predict the Next Word

“This morning I took my cat for a walk.”

given these words

predict the
next word

A Sequence Modeling Problem: Predict the Next Word

“This morning I took my cat for a walk.”

given these words

predict the
next word

Representing Language to a Neural Network

A Sequence Modeling Problem: Predict the Next Word

“This morning I took my cat for a walk.”

given these words

predict the
next word

Representing Language to a Neural Network



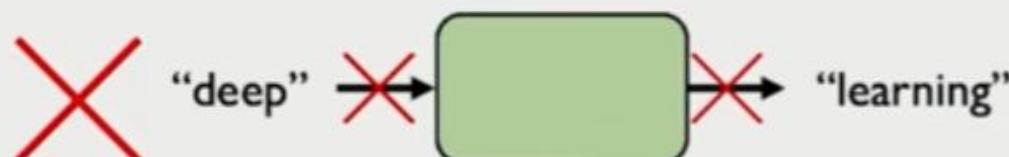
A Sequence Modeling Problem: Predict the Next Word

“This morning I took my cat for a walk.”

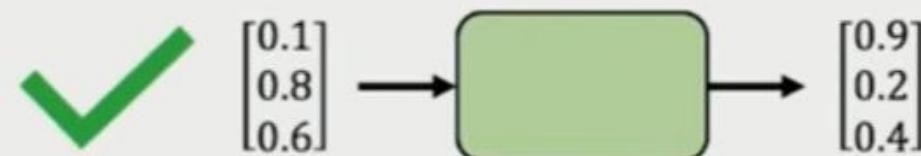
given these words

predict the
next word

Representing Language to a Neural Network

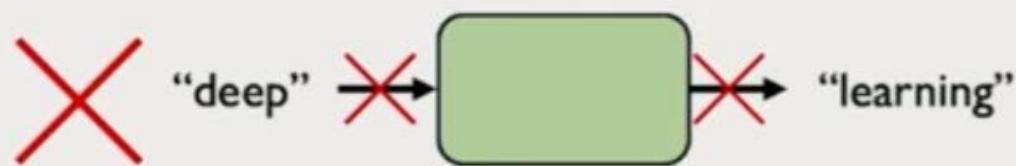


Neural networks cannot interpret words

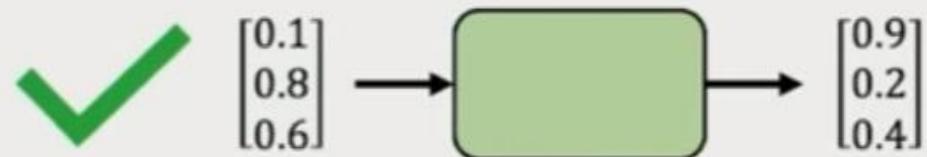


Neural networks require numerical inputs

Encoding Language for a Neural Network

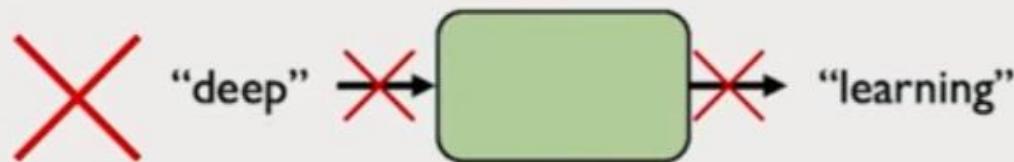


Neural networks cannot interpret words

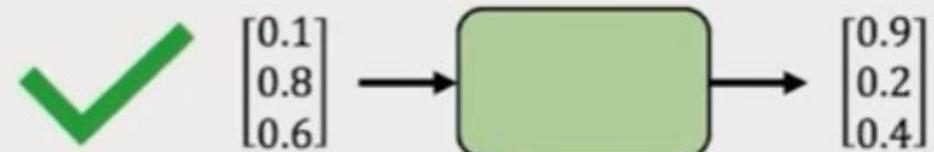


Neural networks require numerical inputs

Encoding Language for a Neural Network



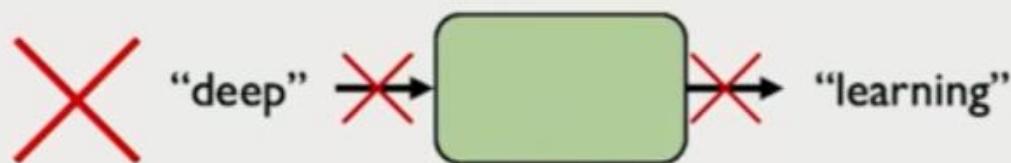
Neural networks cannot interpret words



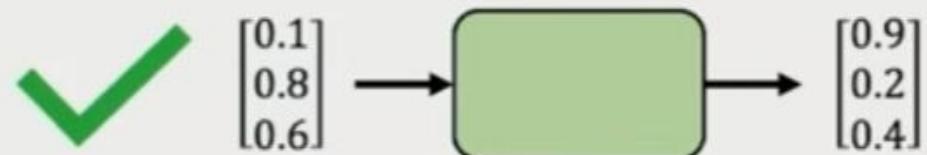
Neural networks require numerical inputs

Embedding: transform indexes into a vector of fixed size.

Encoding Language for a Neural Network

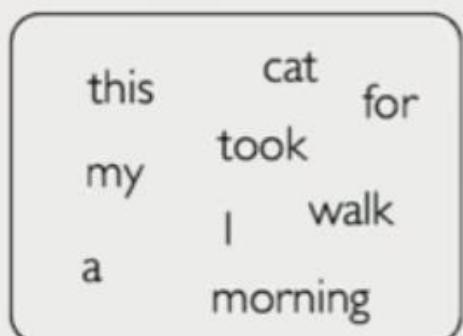


Neural networks cannot interpret words



Neural networks require numerical inputs

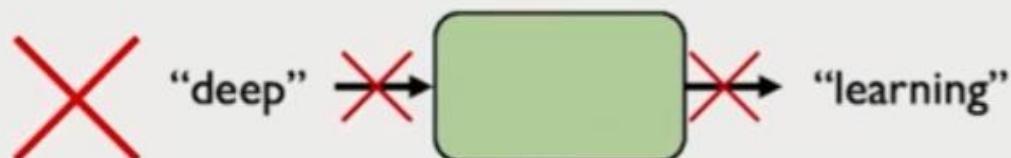
Embedding: transform indexes into a vector of fixed size.



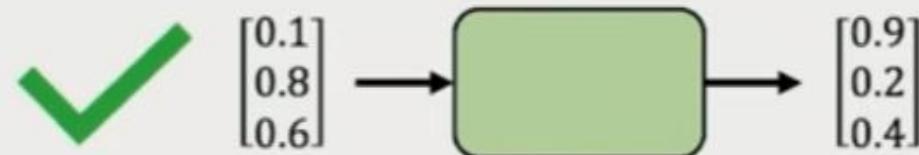
I. Vocabulary:

Corpus of words

Encoding Language for a Neural Network



Neural networks cannot interpret words



Neural networks require numerical inputs

Embedding: transform indexes into a vector of fixed size.

this	cat	for
my	took	I
a		walk
		morning

1. Vocabulary:

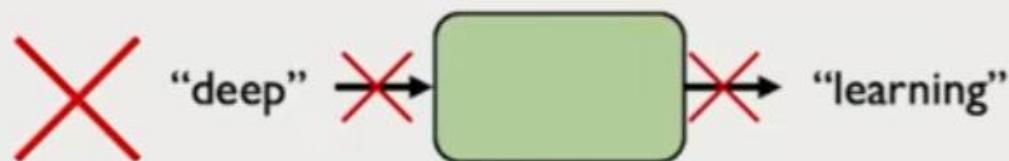
Corpus of words

a	→	1
cat	→	2
...
walk	→	N

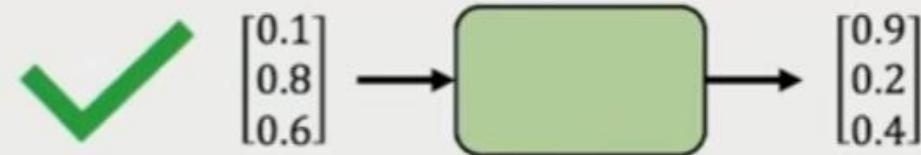
2. Indexing:

Word to index

Encoding Language for a Neural Network



Neural networks cannot interpret words



Neural networks require numerical inputs

Embedding: transform indexes into a vector of fixed size.

this	cat	for
my	took	I
a	walk	morning

I. Vocabulary:
Corpus of words

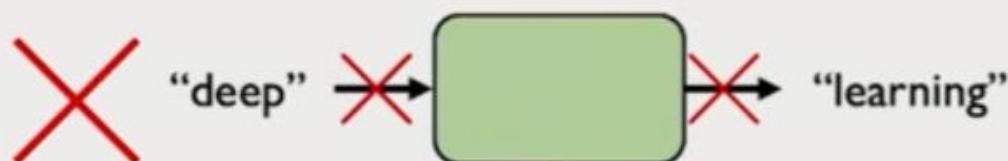
a	→	1
cat	→	2
...
walk	→	N

2. Indexing:
Word to index

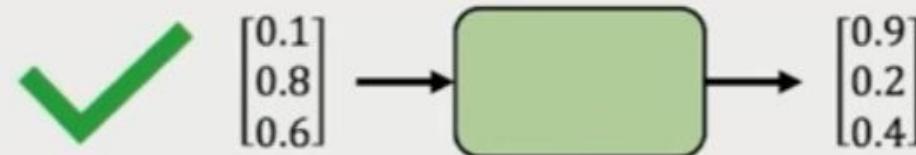
One-hot embedding
"cat" = [0, 1, 0, 0, 0, 0]
i-th index

3. Embedding:
Index to fixed-sized vector

Encoding Language for a Neural Network



Neural networks cannot interpret words



Neural networks require numerical inputs

Embedding: transform indexes into a vector of fixed size.

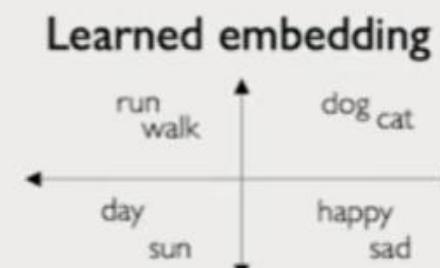
this	cat	for
my	took	I
a	walk	
	morning	

1. Vocabulary:
Corpus of words

a	→	1
cat	→	2
...	...	
walk	→	N

2. Indexing:
Word to index

One-hot embedding
“cat” = $[0, 1, 0, 0, 0, 0]$
i-th index



3. Embedding:
Index to fixed-sized vector

Handle Variable Sequence Lengths

The food was great

Handle Variable Sequence Lengths

The food was great

vs.

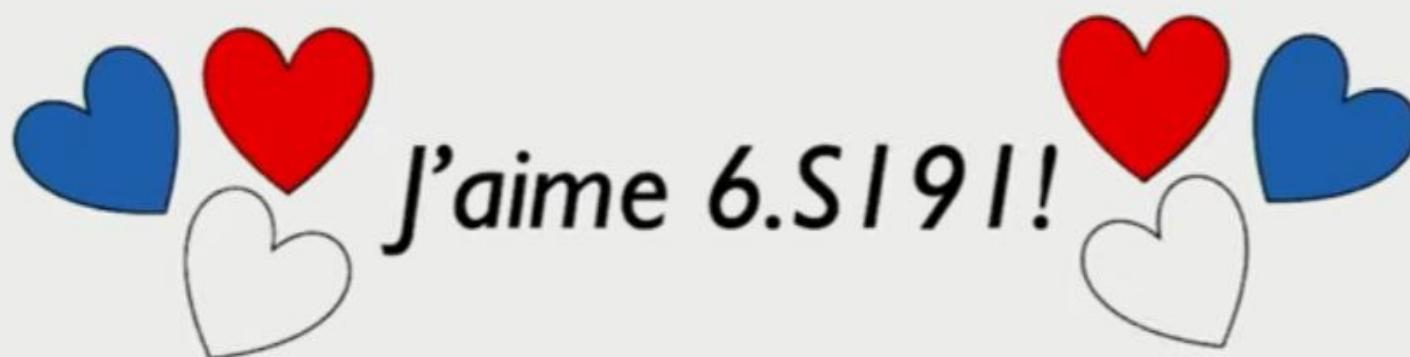
We visited a restaurant for lunch

vs.

We were hungry but cleaned the house before eating

Model Long-Term Dependencies

“France is where I grew up, but I now live in Boston. I speak fluent ____.”



We need information from **the distant past** to accurately predict the correct word.

Capture Differences in Sequence Order



The food was good, not bad at all.

vs.

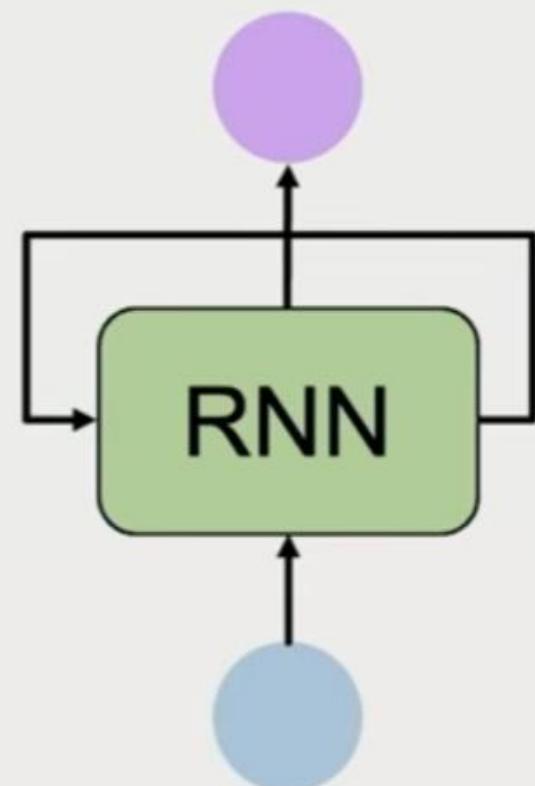
The food was bad, not good at all.



Sequence Modeling: Design Criteria

To model sequences, we need to:

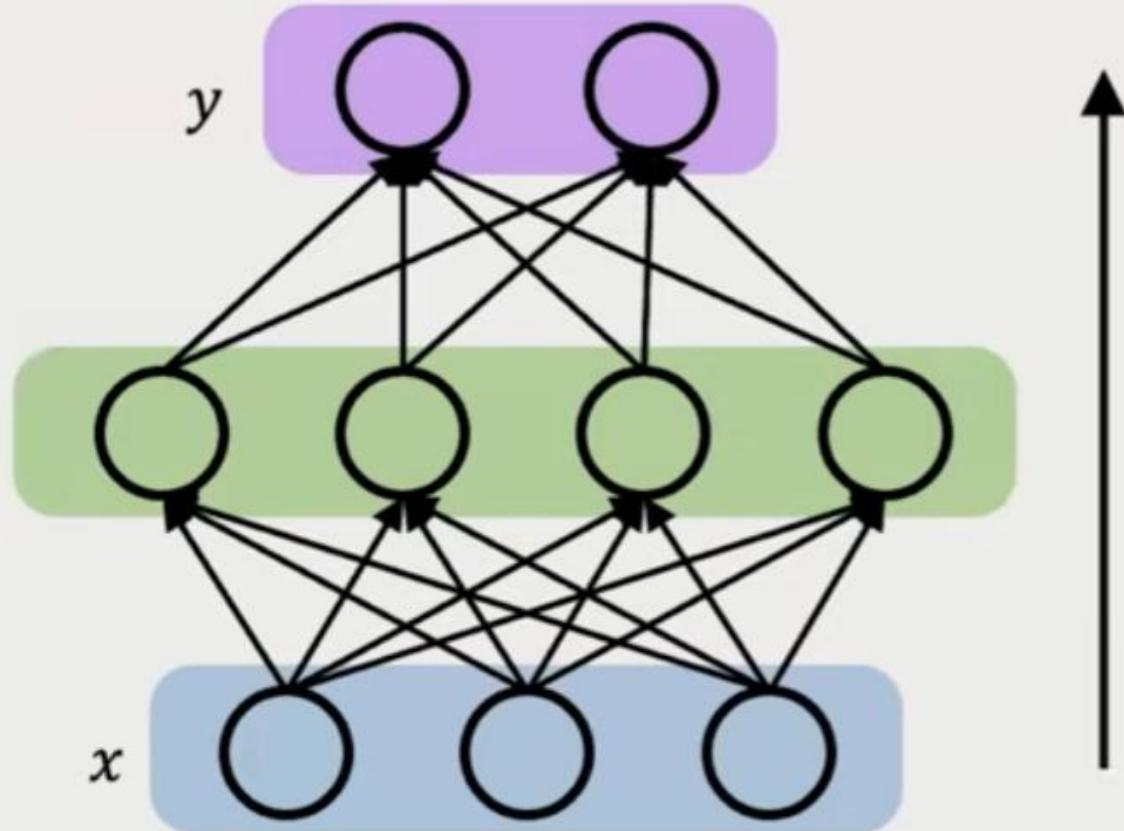
1. Handle **variable-length** sequences
2. Track **long-term** dependencies
3. Maintain information about **order**
4. **Share parameters** across the sequence



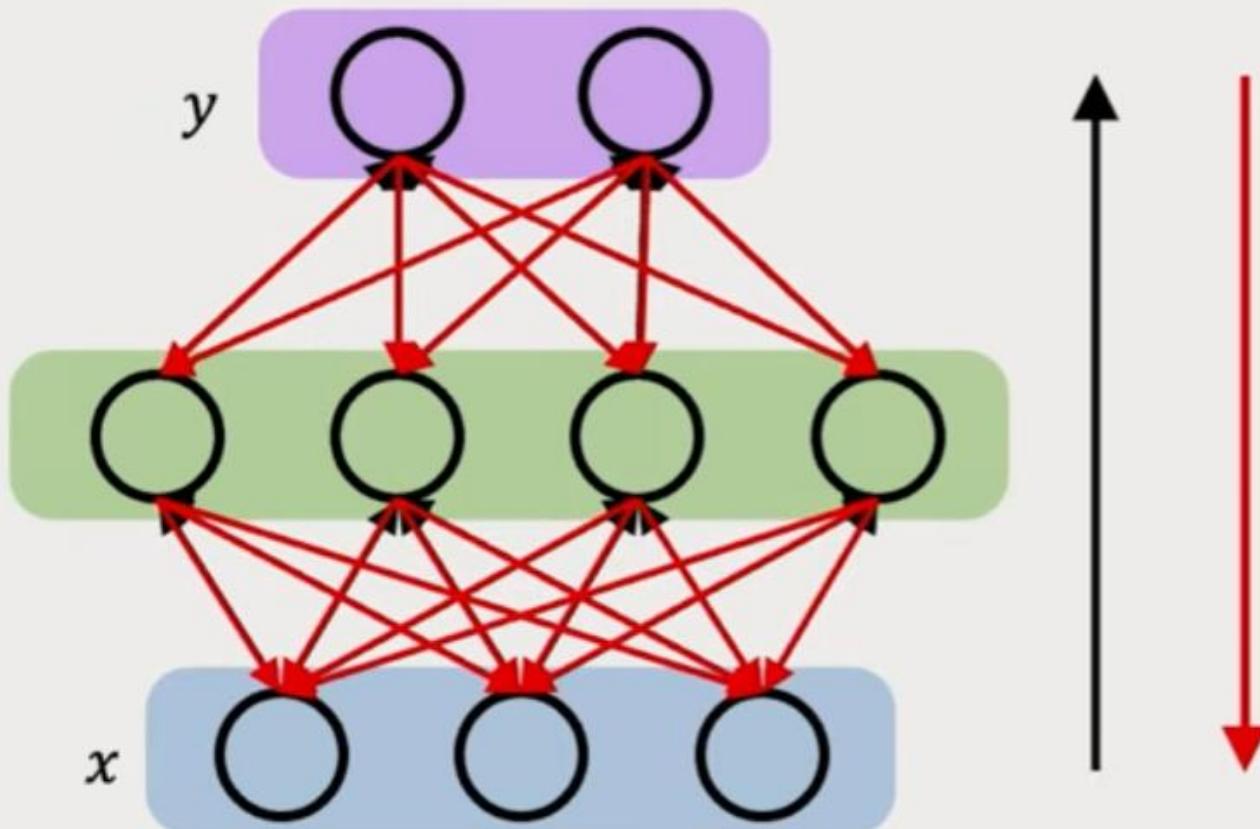
Recurrent Neural Networks (RNNs) meet
these sequence modeling design criteria

Backpropagation Through Time (BPTT)

Recall: Backpropagation in Feed Forward Models



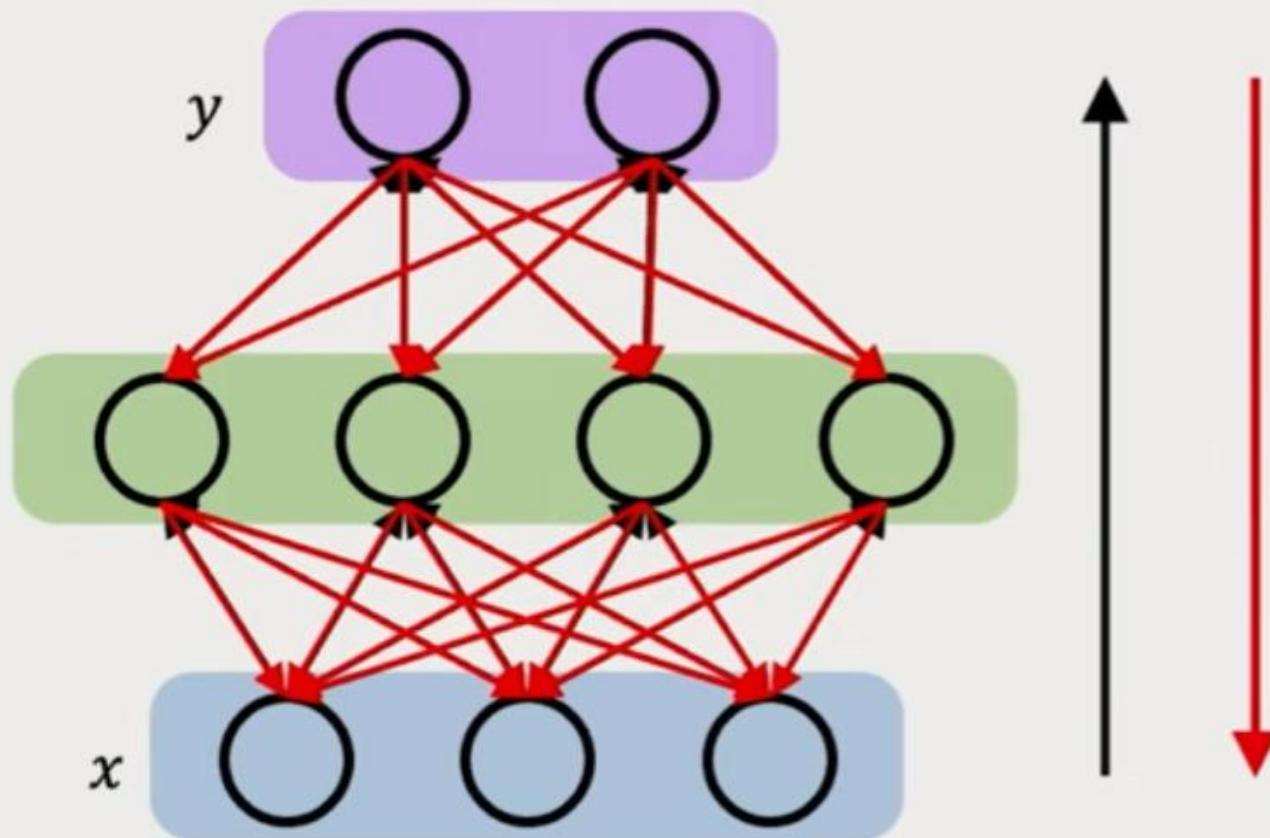
Recall: Backpropagation in Feed Forward Models



Backpropagation algorithm:

- I. Take the derivative (gradient) of the loss with respect to each parameter

Recall: Backpropagation in Feed Forward Models

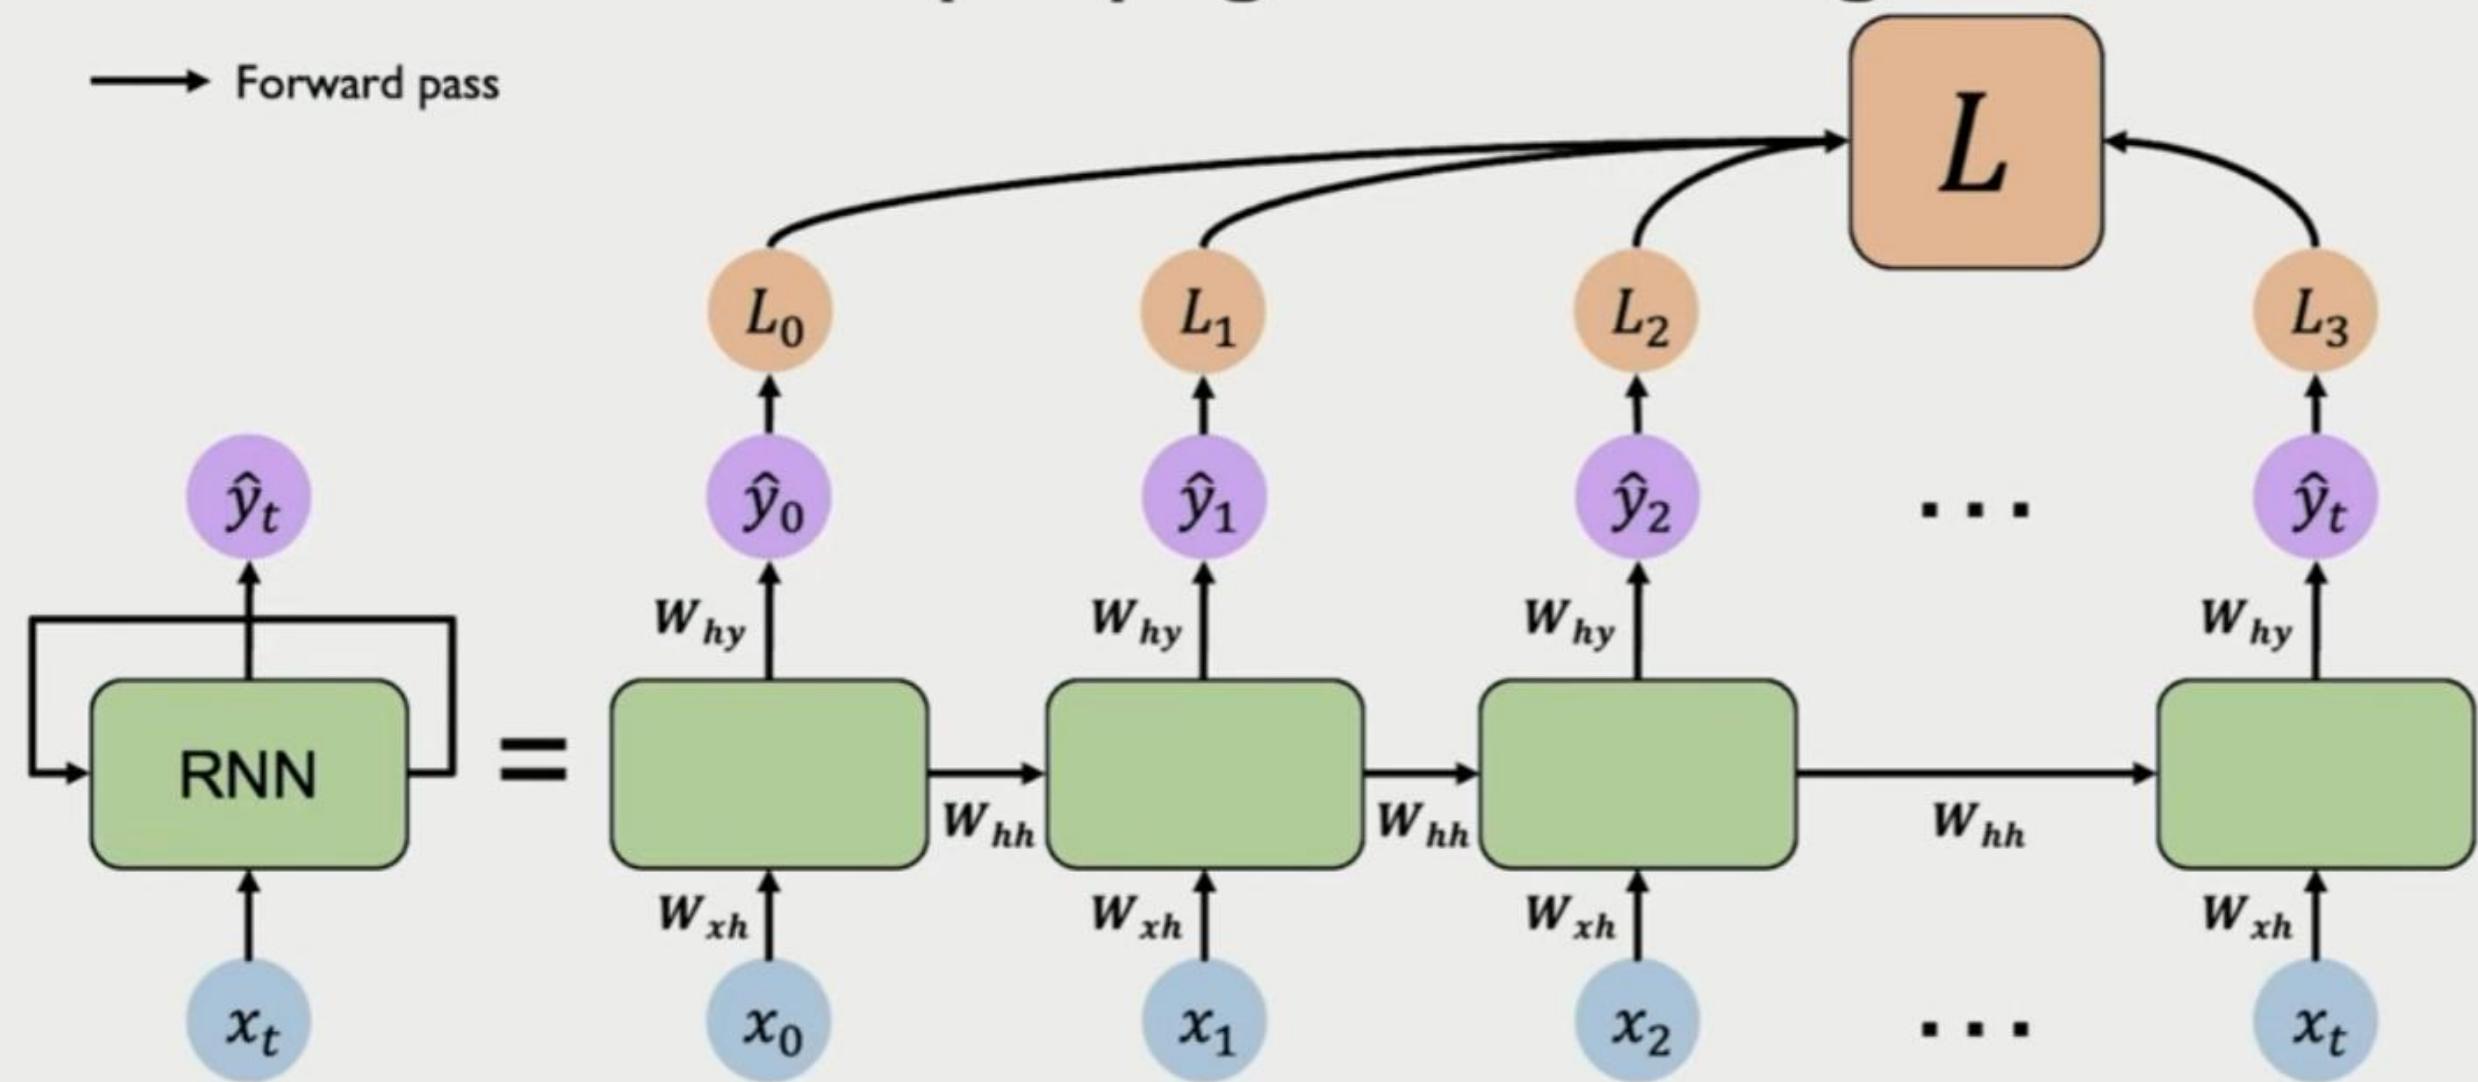


Backpropagation algorithm:

1. Take the derivative (gradient) of the loss with respect to each parameter
2. Shift parameters in order to minimize loss

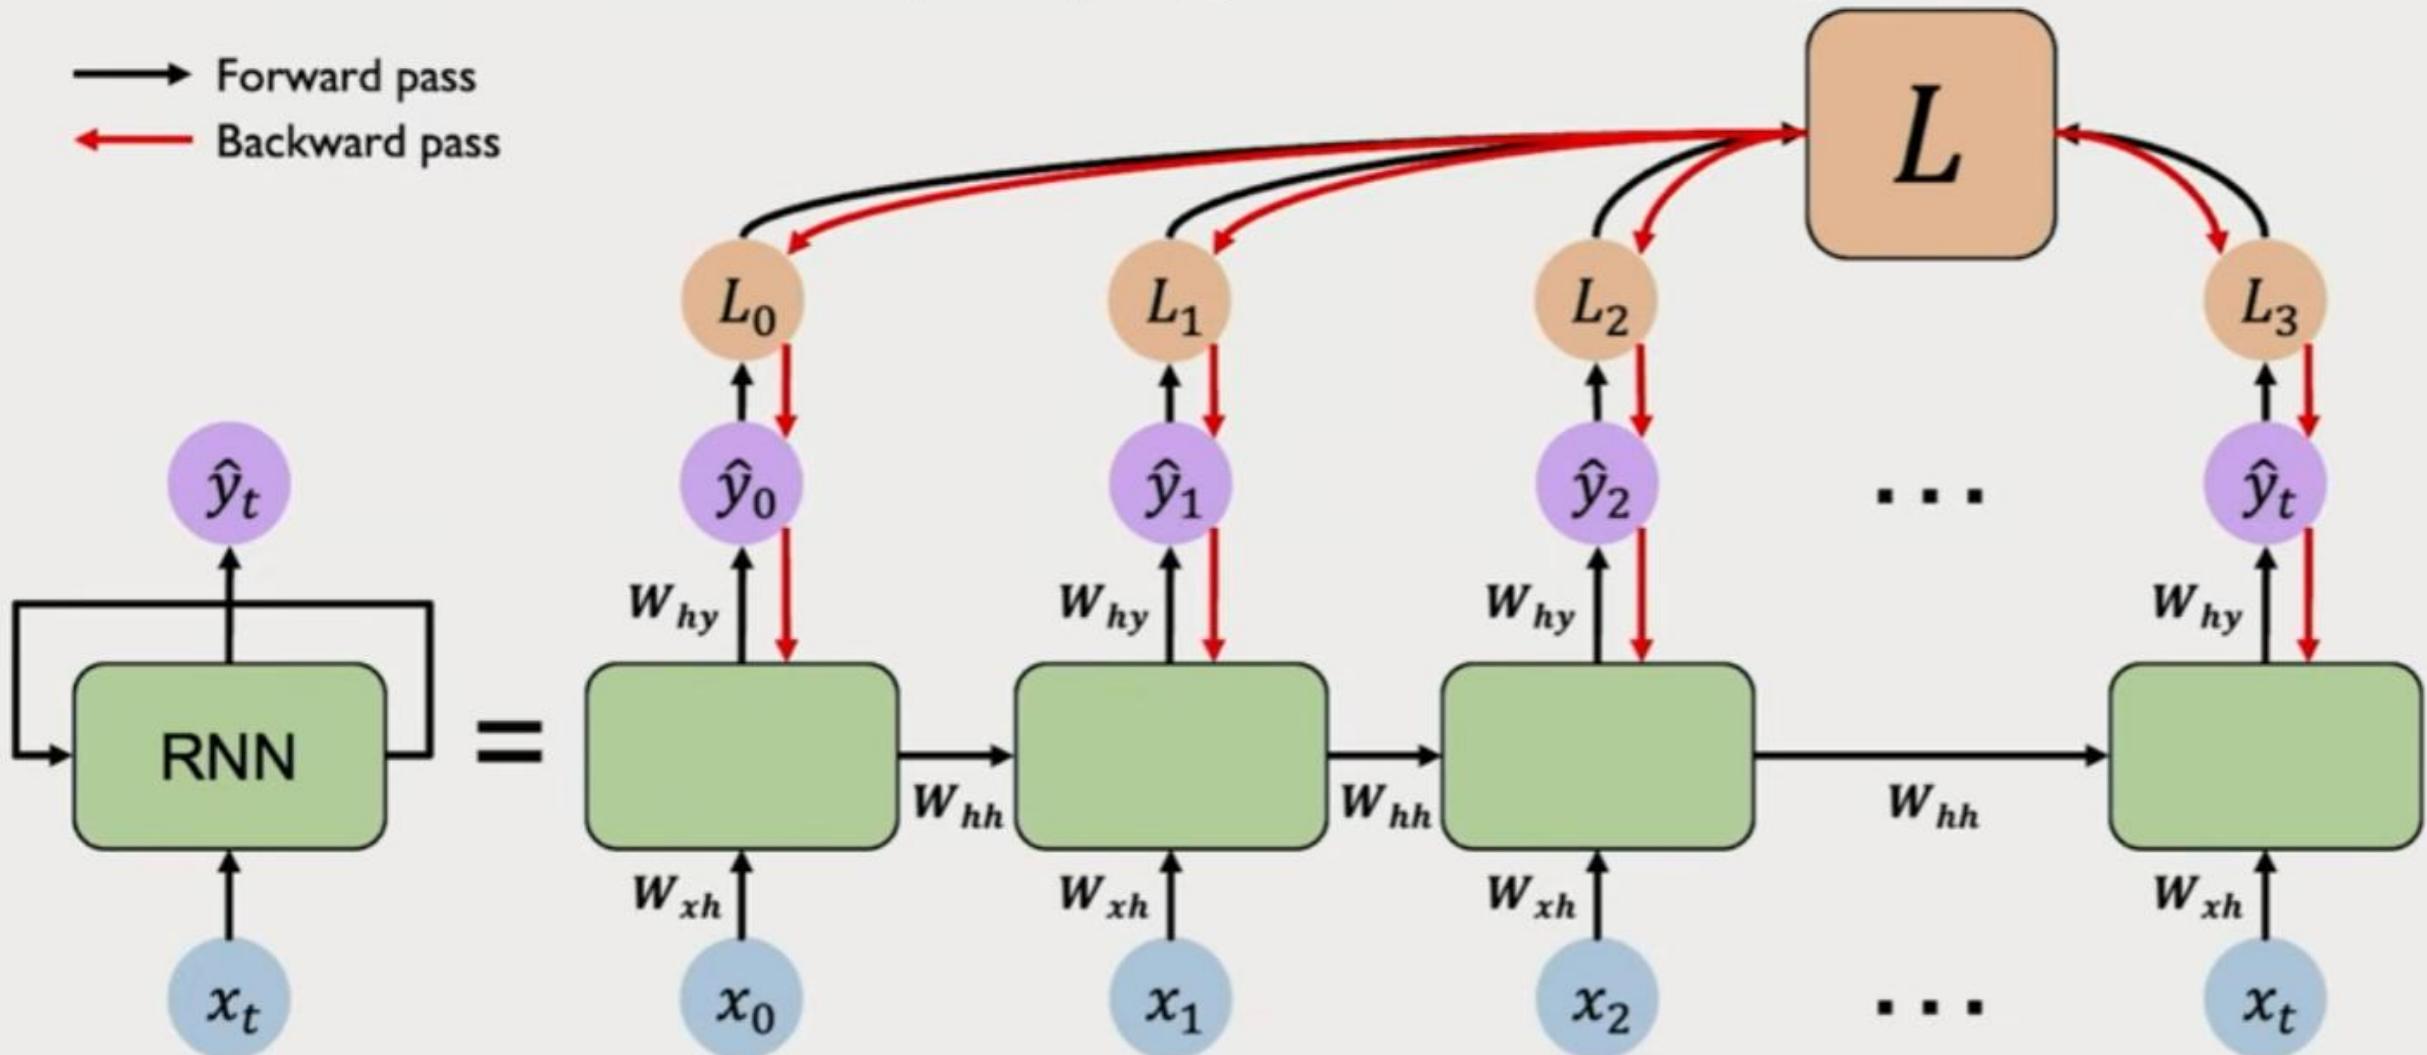
RNNs: Backpropagation Through Time

→ Forward pass



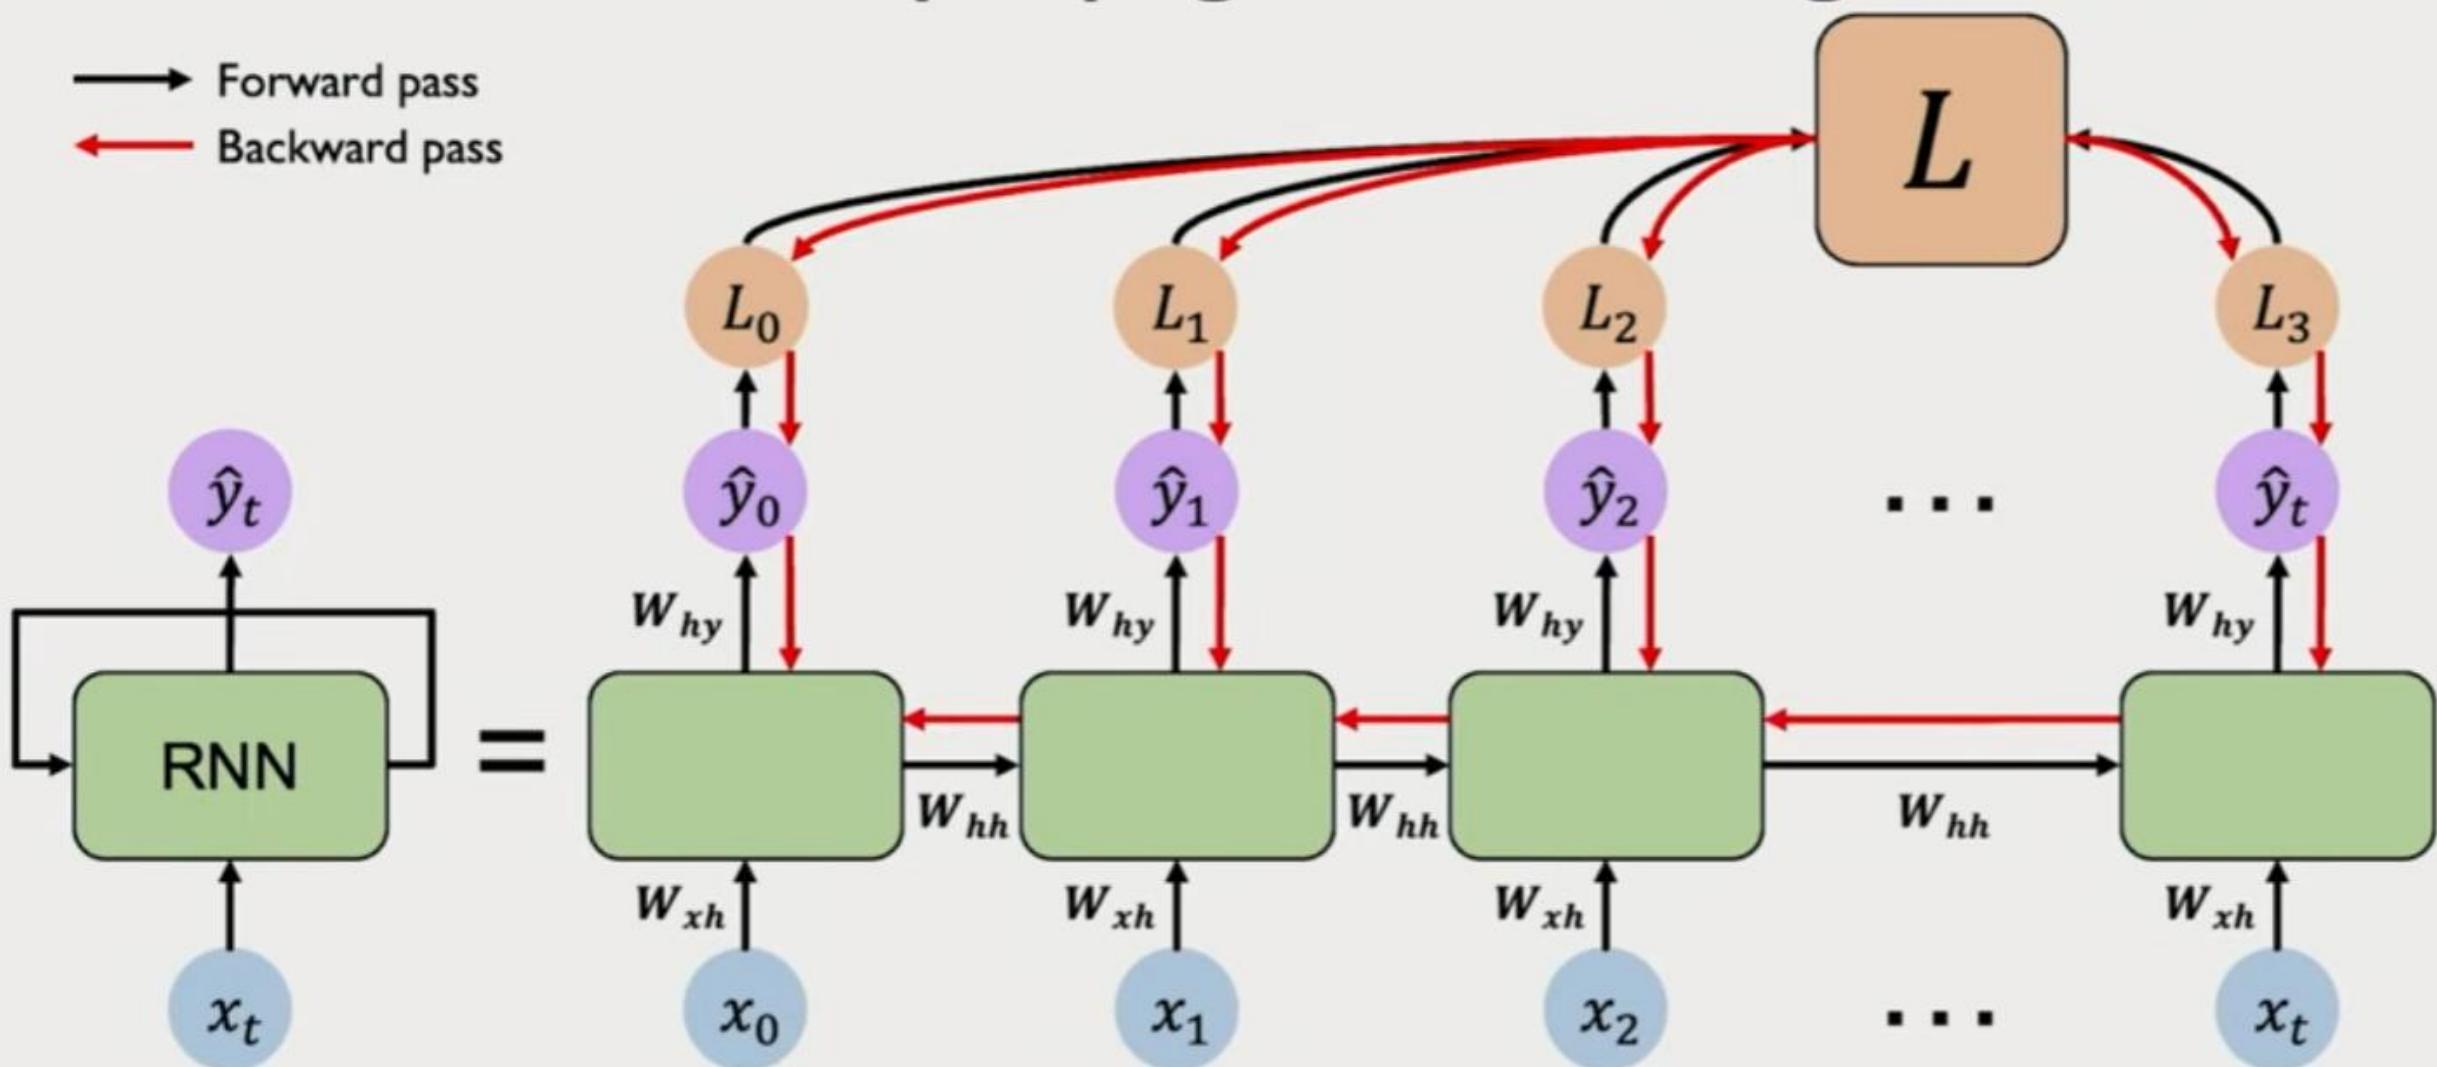
RNNs: Backpropagation Through Time

→ Forward pass
← Backward pass



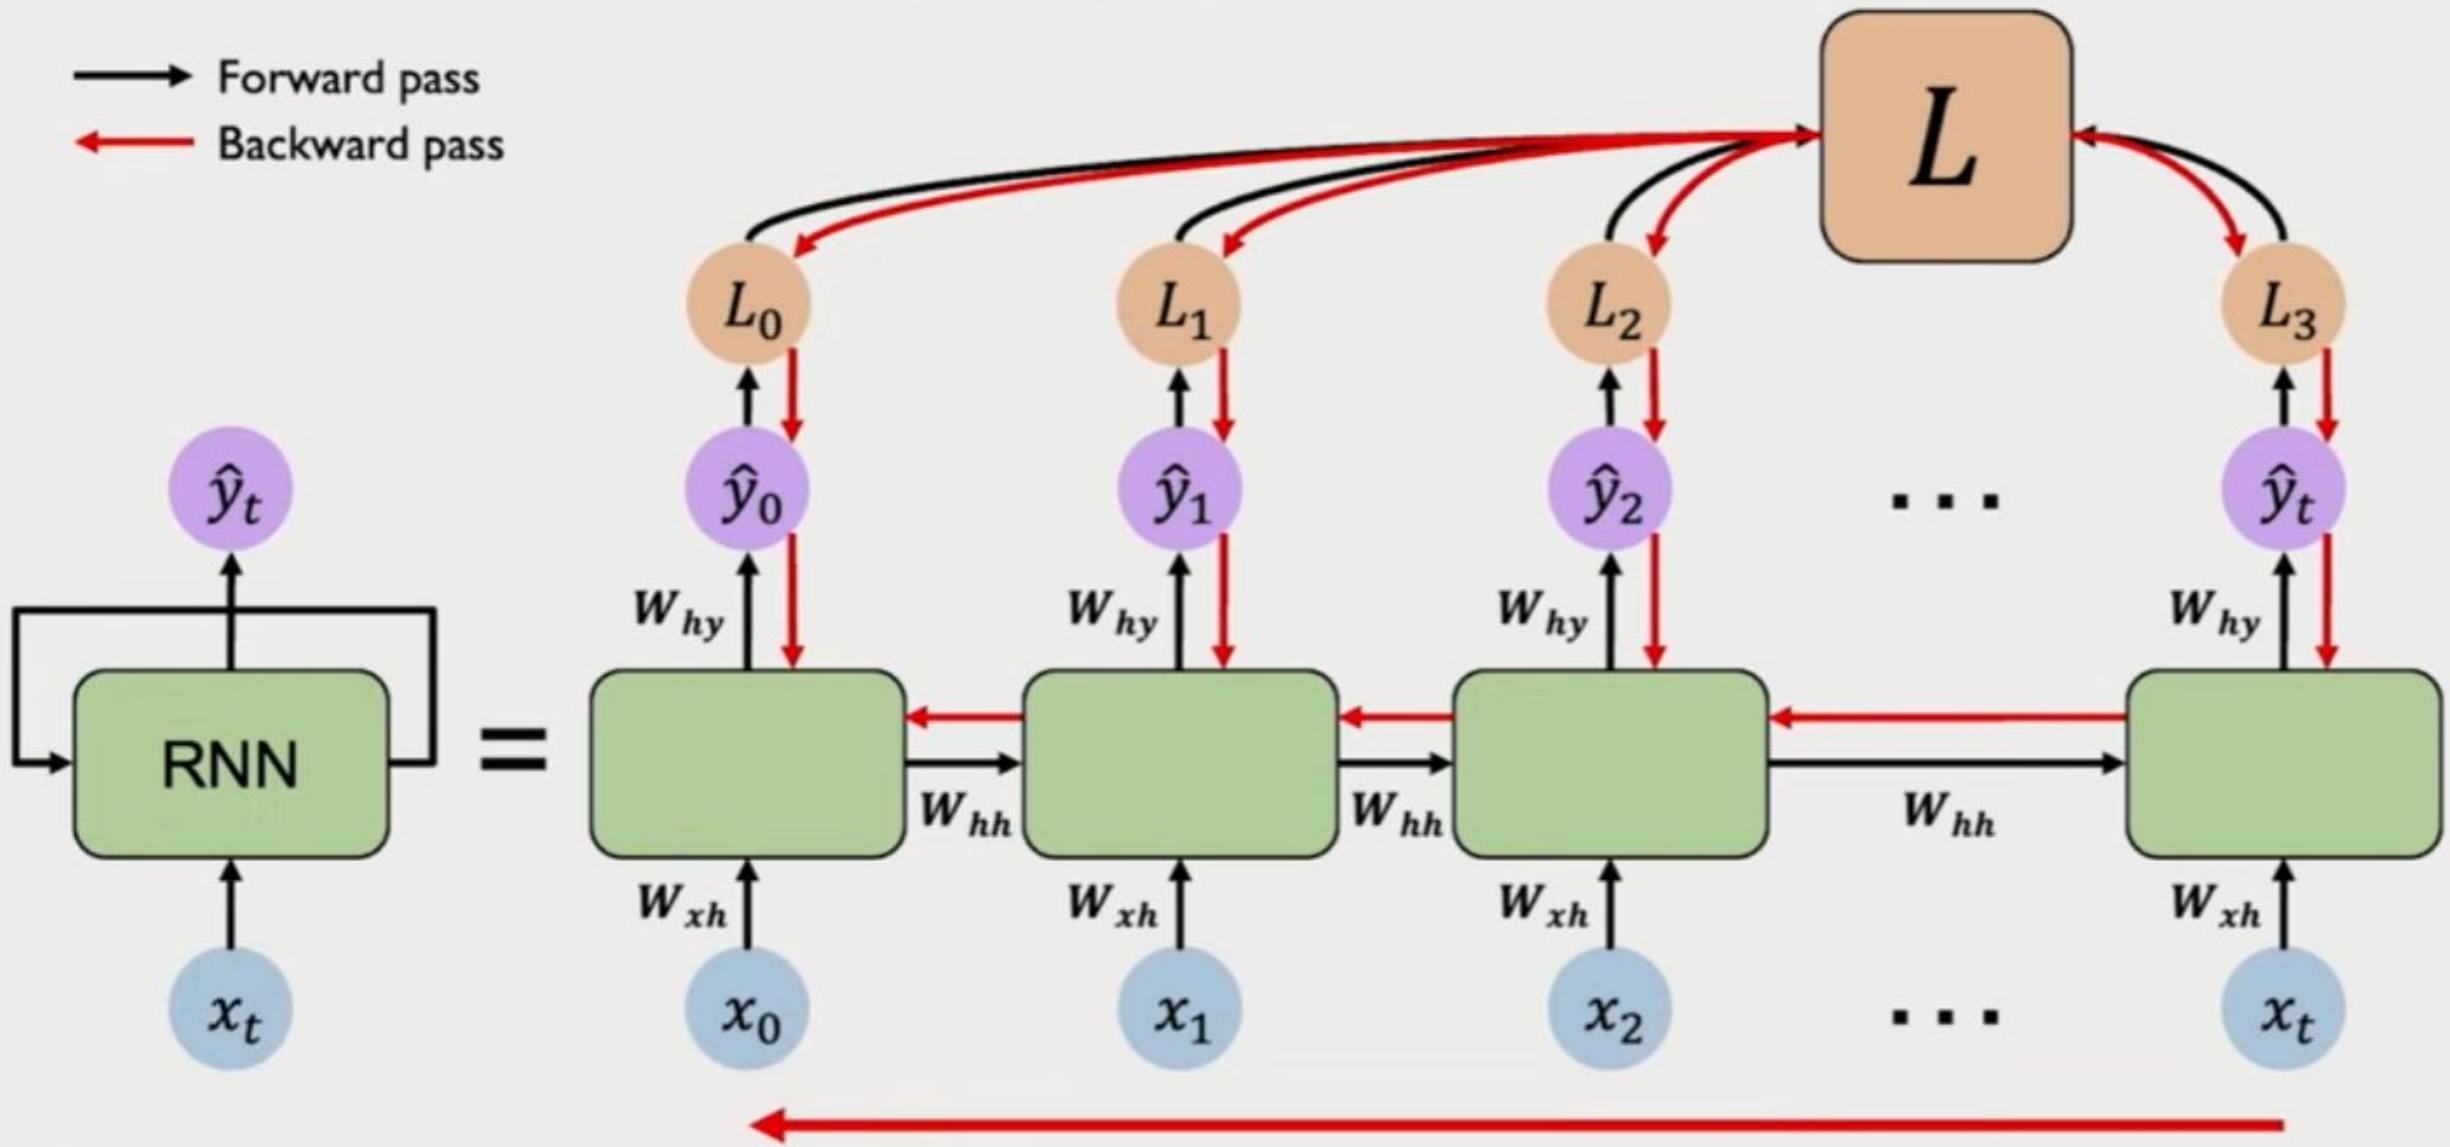
RNNs: Backpropagation Through Time

→ Forward pass
← Backward pass

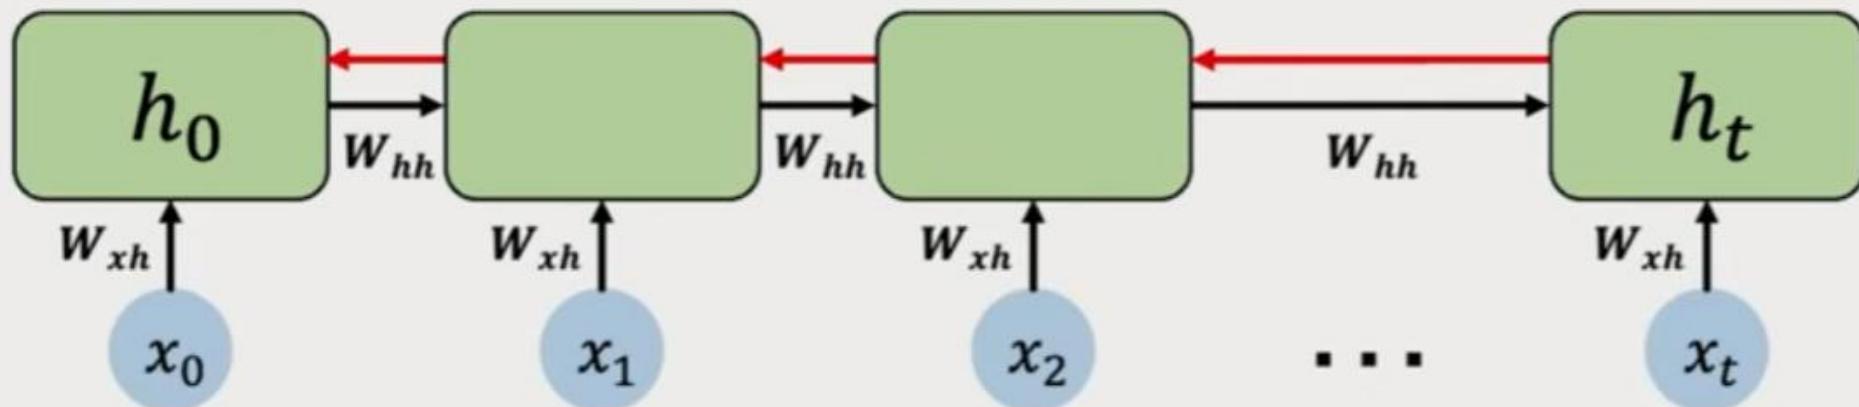


RNNs: Backpropagation Through Time

→ Forward pass
← Backward pass

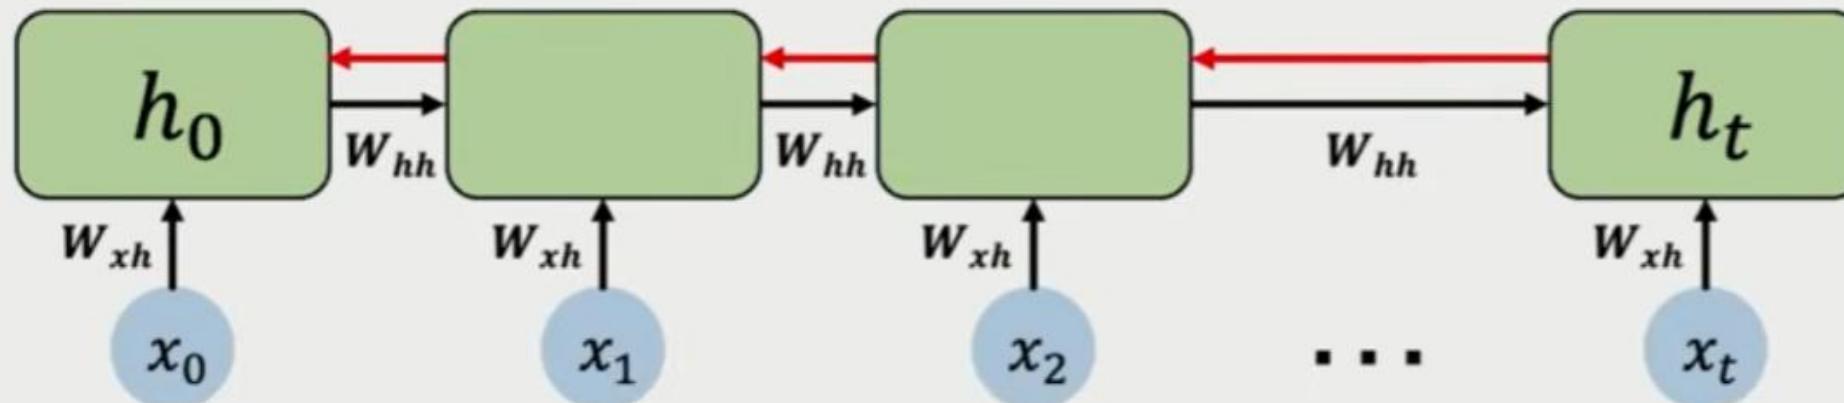


Standard RNN Gradient Flow



Computing the gradient wrt h_0 involves many factors of W_{hh} + repeated gradient computation!

Standard RNN Gradient Flow: Exploding Gradients

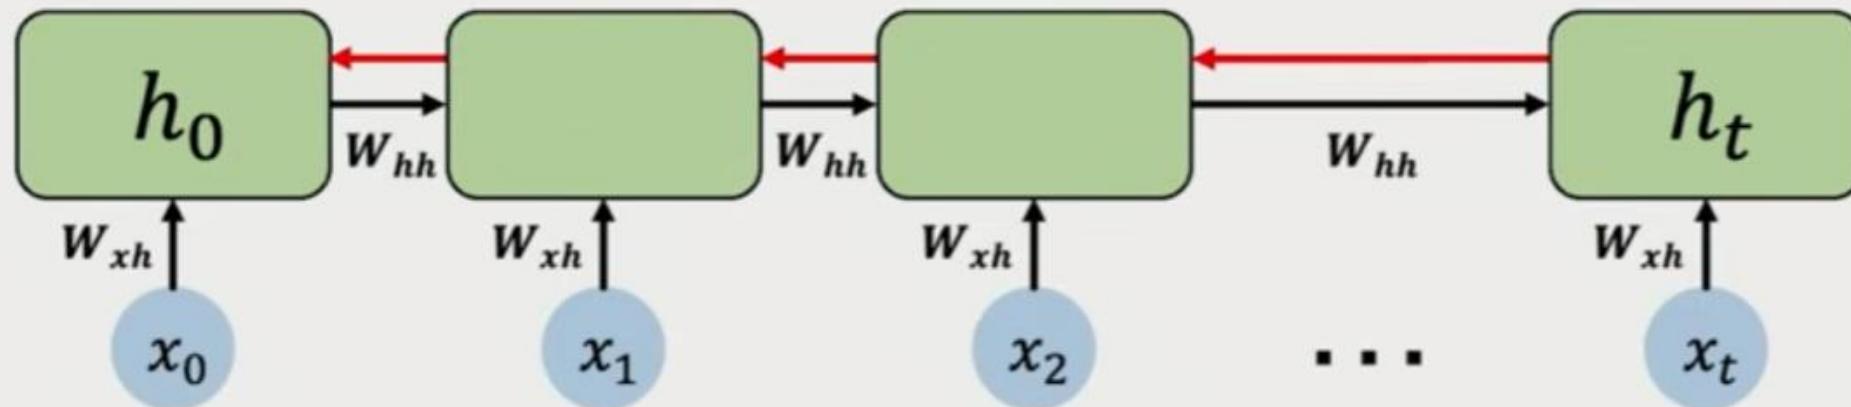


Computing the gradient wrt h_0 involves **many factors of W_{hh}** + repeated gradient computation!

Many values > 1:
exploding gradients

Gradient clipping to
scale big gradients

Standard RNN Gradient Flow: Vanishing Gradients



Computing the gradient wrt h_0 involves many factors of W_{hh} + repeated gradient computation!

Many values > 1 :
exploding gradients

Gradient clipping to
scale big gradients

Many values < 1 :
vanishing gradients

1. Activation function
2. Weight initialization
3. Network architecture

The Problem of Long-Term Dependencies

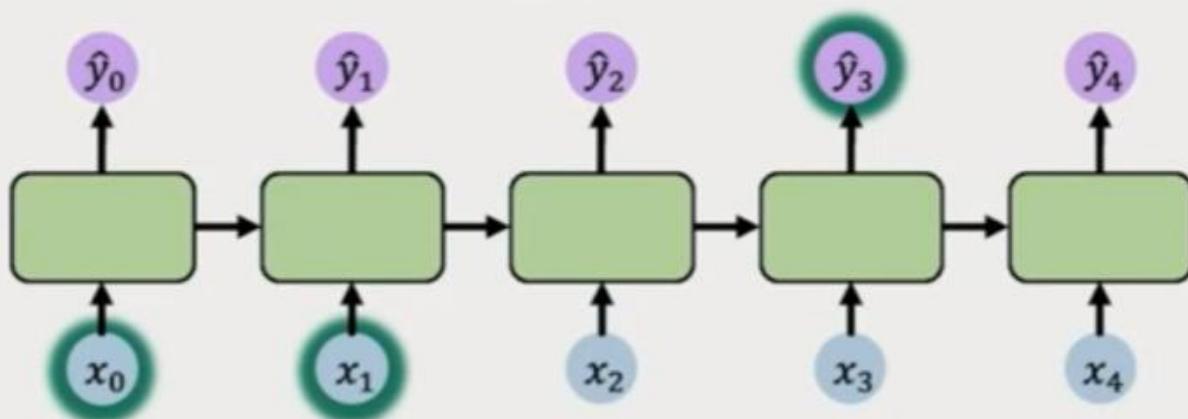
Why are vanishing gradients a problem?

Multiply many **small numbers** together

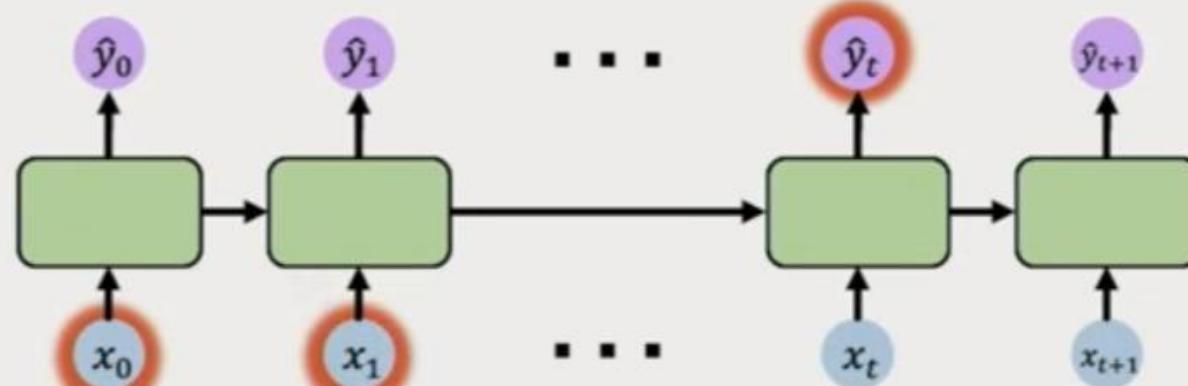
↓
Errors due to further back time steps
have smaller and smaller gradients

↓
Bias parameters to capture short-term
dependencies

"The clouds are in the ___"



"I grew up in France, ... and I speak fluent ___ "

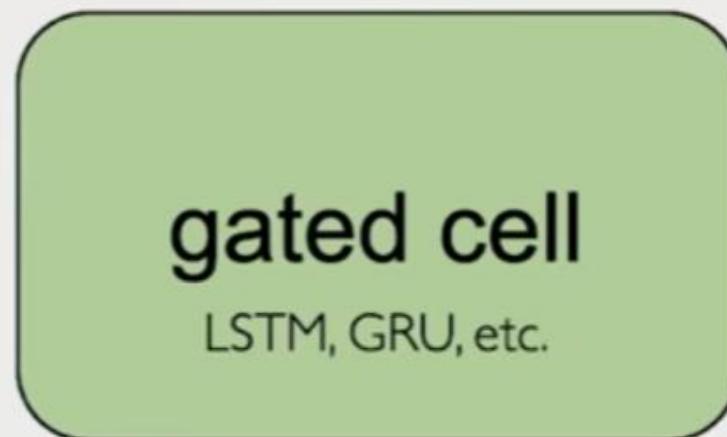
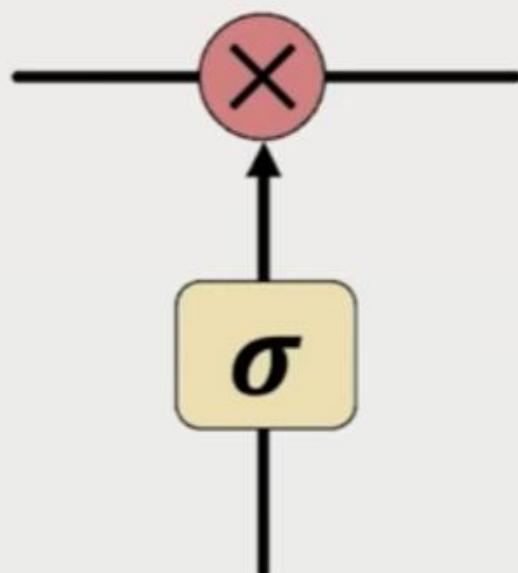


Gating Mechanisms in Neurons

Idea: use **gates** to selectively **add** or **remove** information within **each recurrent unit** with

Pointwise multiplication

Sigmoid neural net layer



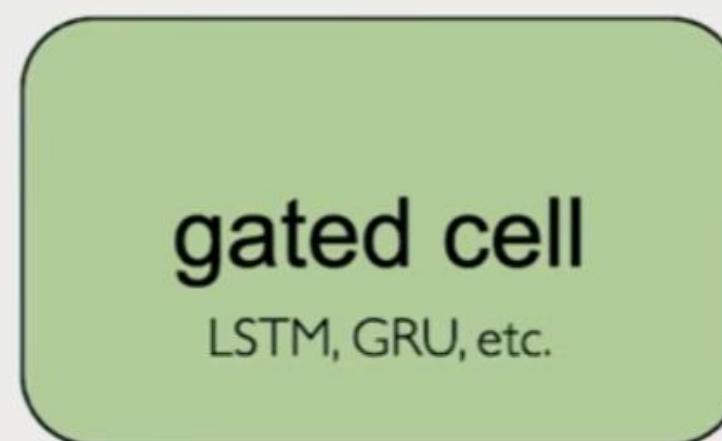
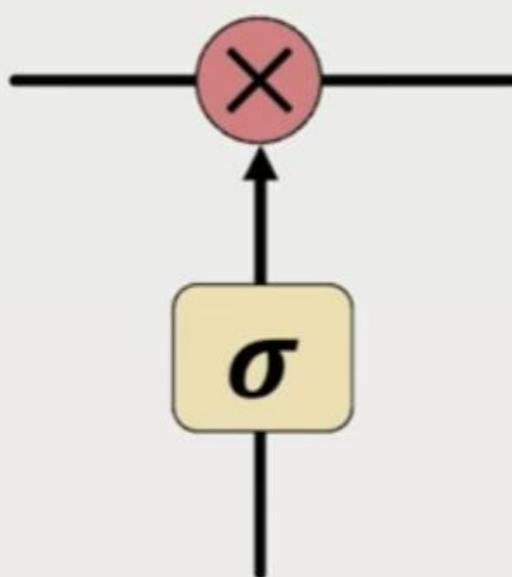
Gates optionally let information through the cell

Gating Mechanisms in Neurons

Idea: use **gates** to selectively **add** or **remove** information within **each recurrent unit with**

Pointwise multiplication

Sigmoid neural net layer



Gates optionally let information through the cell

Long Short Term Memory (LSTMs) networks rely on a gated cell to track information throughout many time steps.